

CS561 Fall 2018

Foundation of Artificial Intelligence

Intelligent Systems

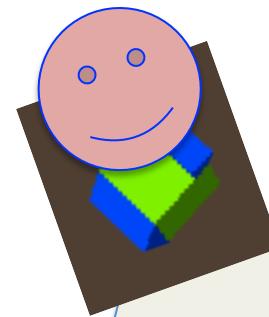
Tuesday Lectures, Session 24-25

6:40-7:55 & 8:05-9:20

Professor Wei-Min Shen

What You Have Learned This Semester

Learn AI and ML
Learn how to learn by yourself



EM Algorithm

Machine Learning

Probability Reasoning

Knowledge Representation

General AI



Self-Reconfigurable Robots, Digital Hormones, and Surprise-Based Learning

Wei-Min (“wingman”) Shen

Polymorphic Robotics Laboratory

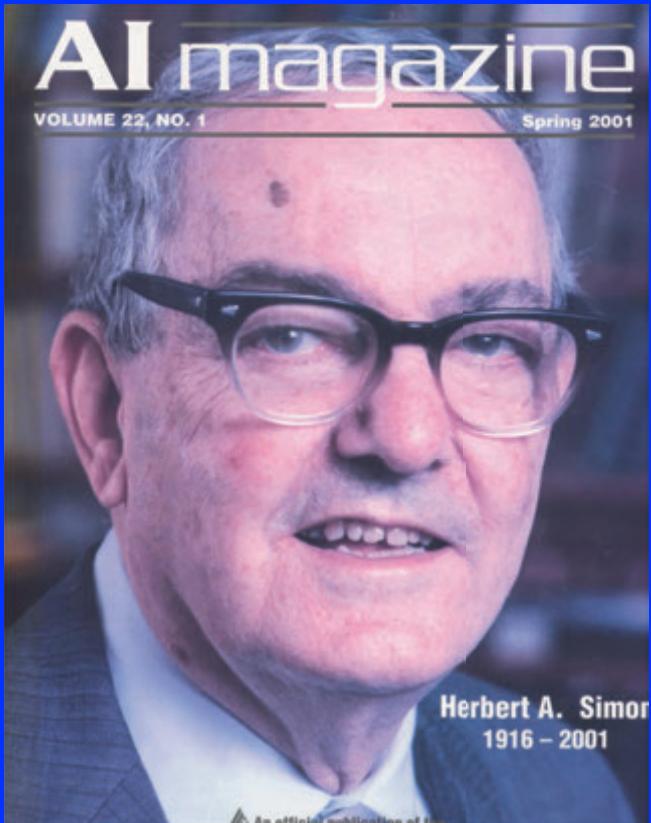
Information Sciences Institute (ISI)

Computer Science Department (CS)

Center for Robotics and Embedded Systems (CRES)

University of Southern California

<http://www.isi.edu/robots>



Forecasting the Future or Shaping it?
October 19, 2000

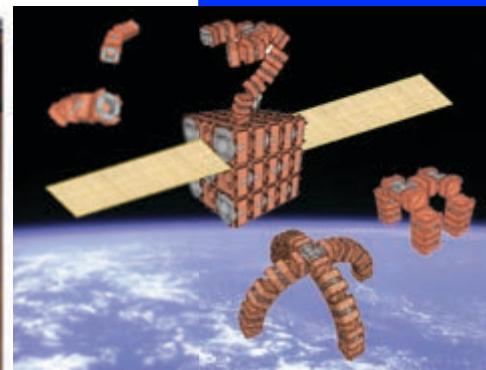
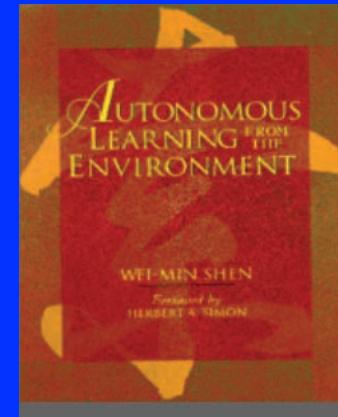
Our task is not to *predict* the future; our task is to *design* a future for a sustainable and acceptable world, and then to devote our efforts to bringing that future about.

Professor Herbert A. Simon
Nobel Price Laureate
A Founder of Artificial Intelligence

Prof Wei-Min Shen <http://www.isi.edu/robots>

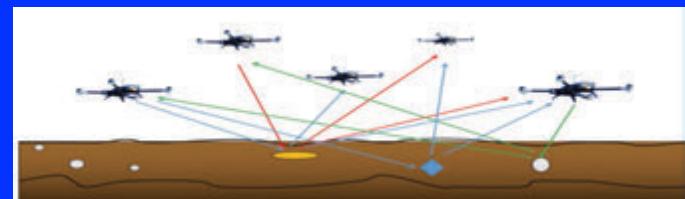
Self-Reconfigurable Modular Robots (System of Systems)

Surprise-Based Learning



Projects

We conduct research in adaptive, self-reconfigurable, autonomous robots and systems, including StarCell, modular, multifunctional and self-reconfigurable SuperBot, Hormone-based Control (HBC), Game Playing, etc.



Polymorphic Robotics Lab

<http://www.isi.edu/robots>

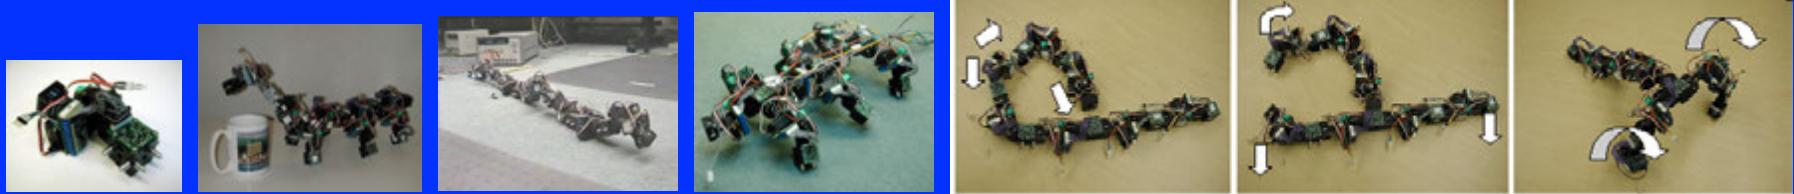
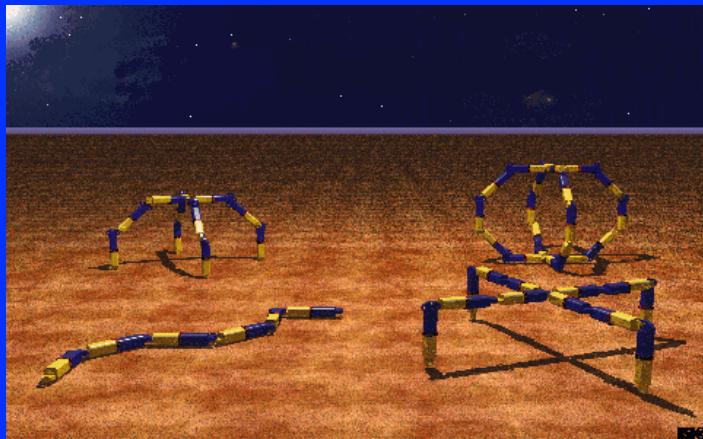
- Our Research
 - ***Self-Reconfigurable Systems***, including robots, AUVs, UAVs, agents, structures, networks, bio-models, ...
- Funding Agencies
 - Government: DARPA, NASA, NSF, AFOSR, ARO, NavLab
 - Industry: Chevron, Motorola, LA, ...
- Awards
 - Navigation Robots: AAAI Robotic Competition (1996)
 - World-championship Soccer-playing robots (1997)
 - Best Paper at SAB Conference (2002)
 - USC Research Awards (2003)
 - Robotic Contiguous Challenge Award, ICRA (2008)
 - Best Paper at 26th Army Science Conference (2008)



World Champion Soccer Robots (Nagoya, Japan, 1997)



Self-Reconfigurable Robots



Introduction Movie

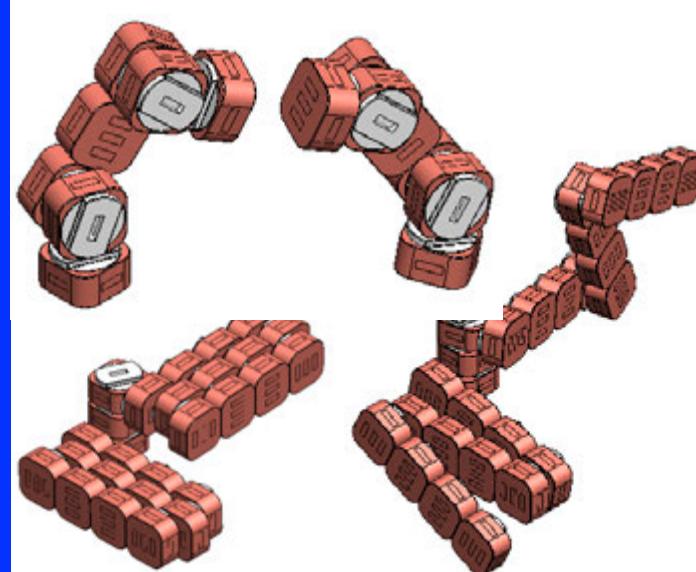
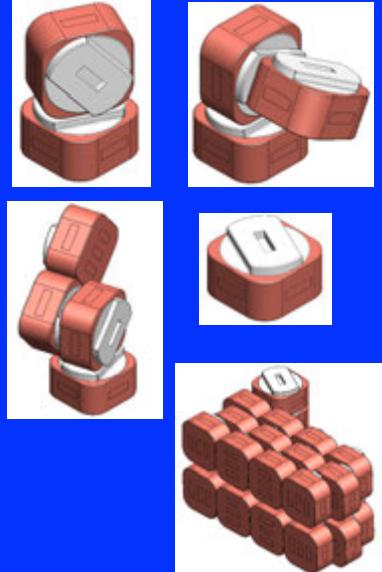
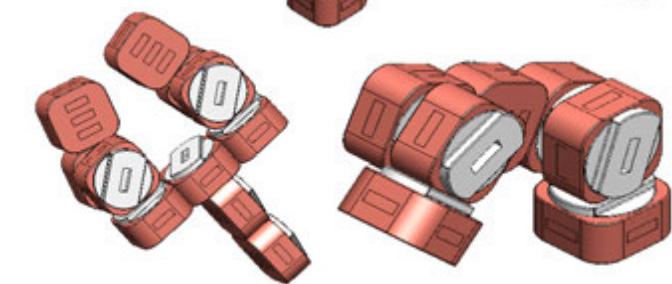
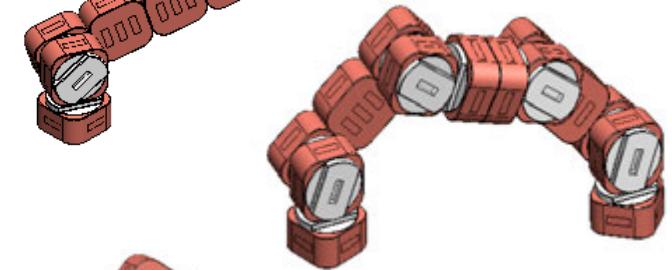
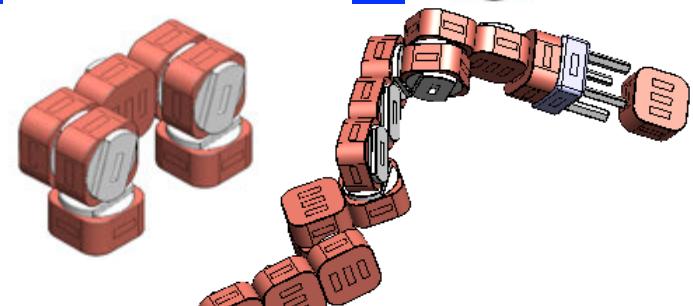
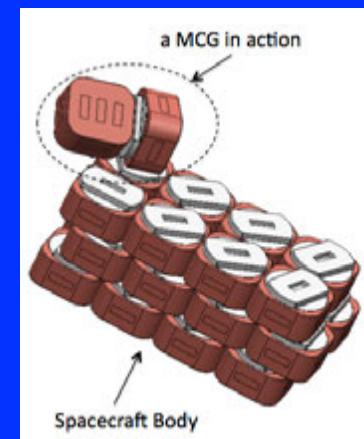
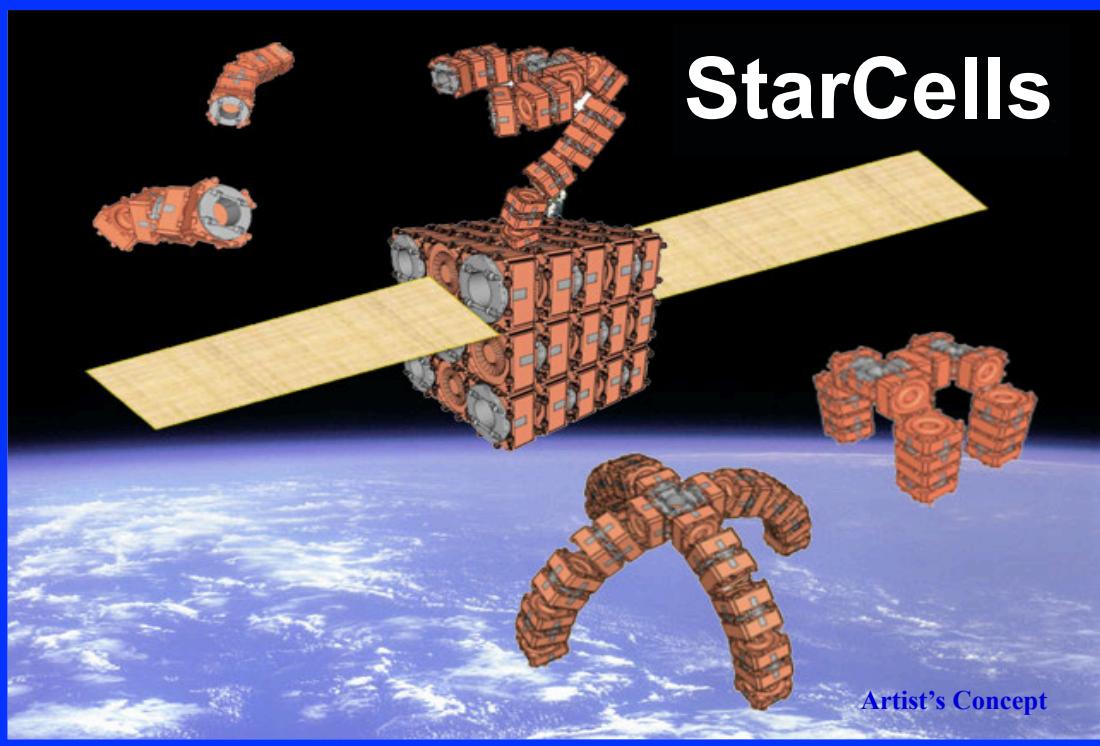


SuperBot Vision for Space

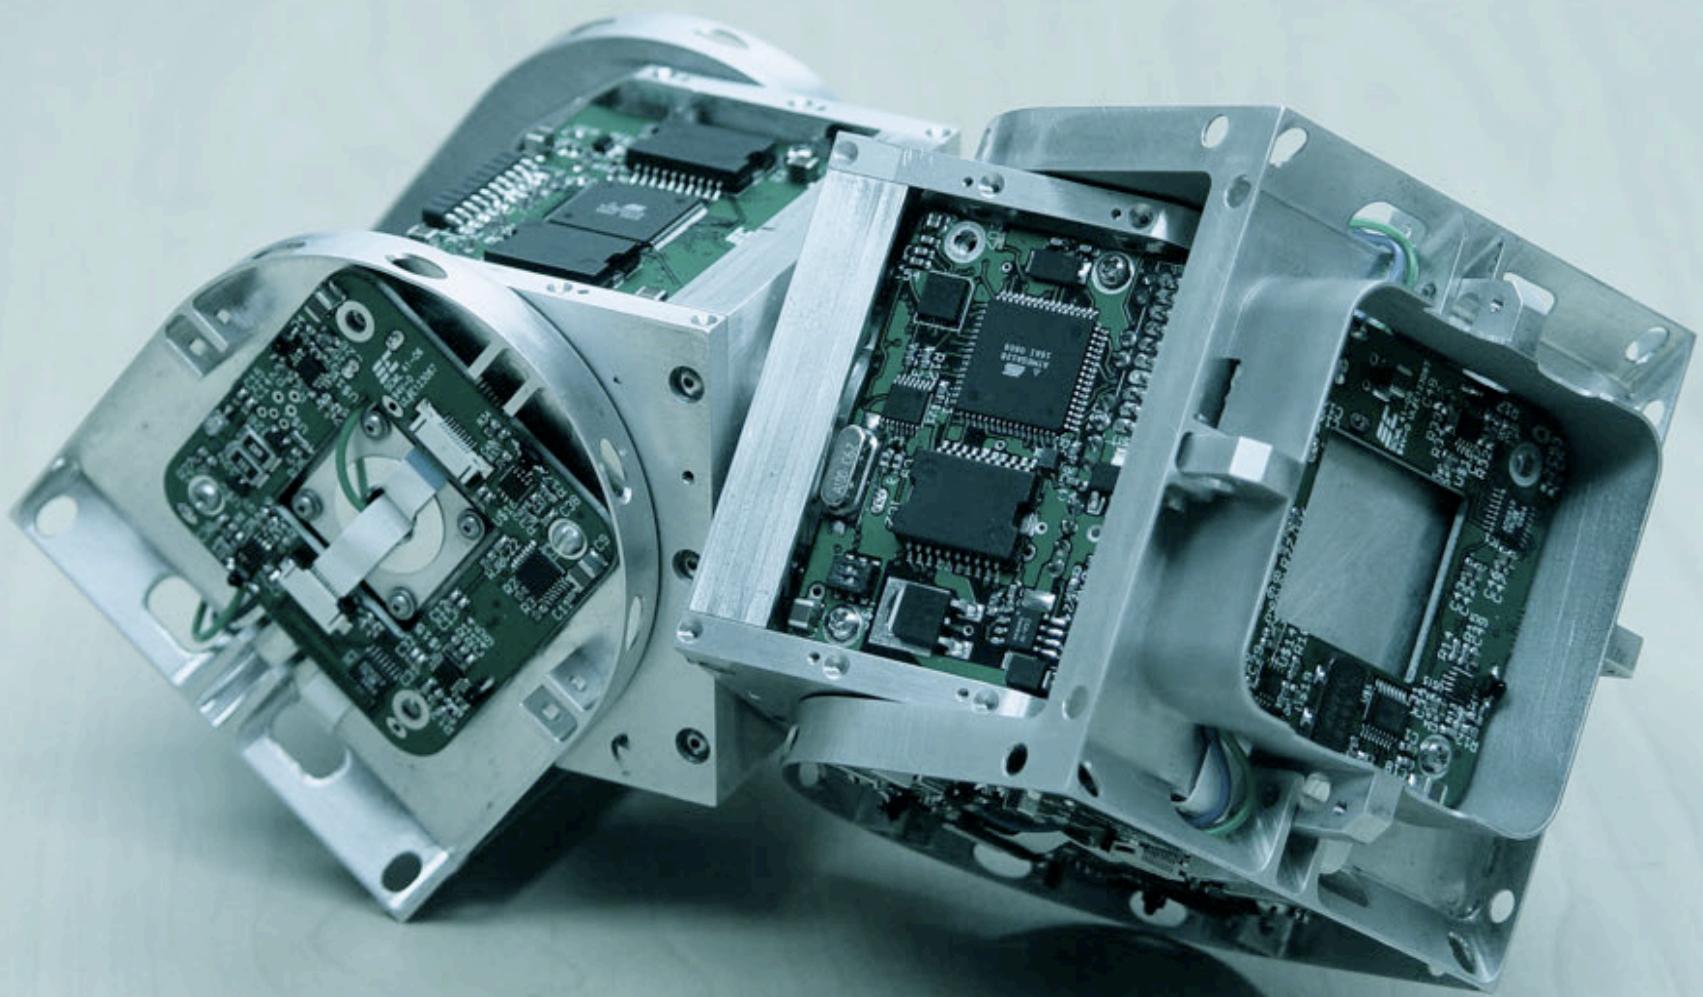
Modular, Multifunction, Self-Reconfiguration



StarCells



Approved for Public Release, Distribution Unlimited



SuperBot: Key Concepts

Invention

- Development of connectable modules with onboard power, connectivity, motion, sensing, and intelligence

Key Features

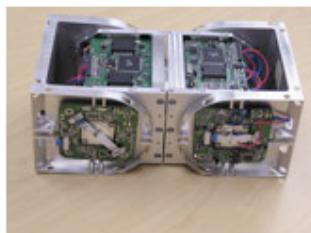
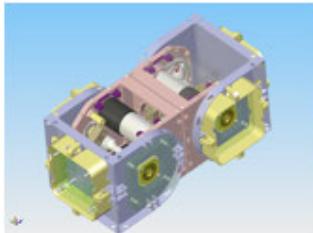
- Capable of operating with other similar modules to morph into multi-shapes and multifunctional self-reconfigurable robots

Applications

- High-risk environments, such as space, underwater, nuclear sites, military, civilian, and reconfigurable/flexible manufacturing

“SuperBot” module specification

- 3-DOF rotational mechanism
- 6 connectors for reconfiguration
- Onboard batteries (7.4v 1.6Ah)
- 2 micro-computers (Atmage128)
- Communication by 4x6 infrared
- 3D gravity sensors
- 4x6 Proximity sensors
- Radio remote control
- Real-time operating system
- C programming environment
- Distributed control
- Topology-triggered behaviors

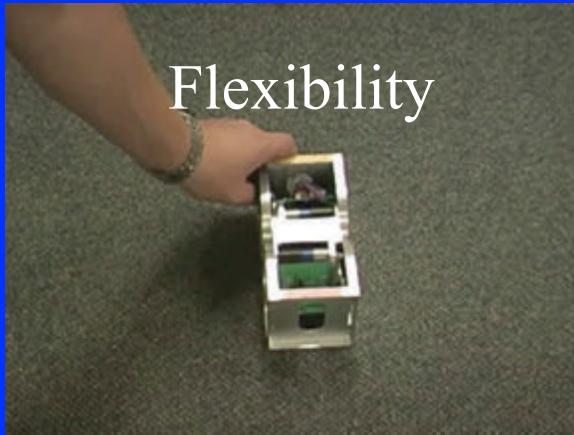


Key Inventions & Patents

- Mechanical mechanisms
 - 3D reconfigurable modules
 - Self-healing connectors (US Patent)
 - Multiplex thrusts (US Patent)
- Electronic Components
 - Power Sharing, local and global communication
 - Multifunctional Mechatronics Architectures
- Intelligent Control Software
 - Digital hormones (US Patent)
 - Surprise-based Learning and Adaptation
 - Self-reconfiguration on-line planning
 - Self-assembly and self-healing

Single Module Autonomous Robots

Flexibility



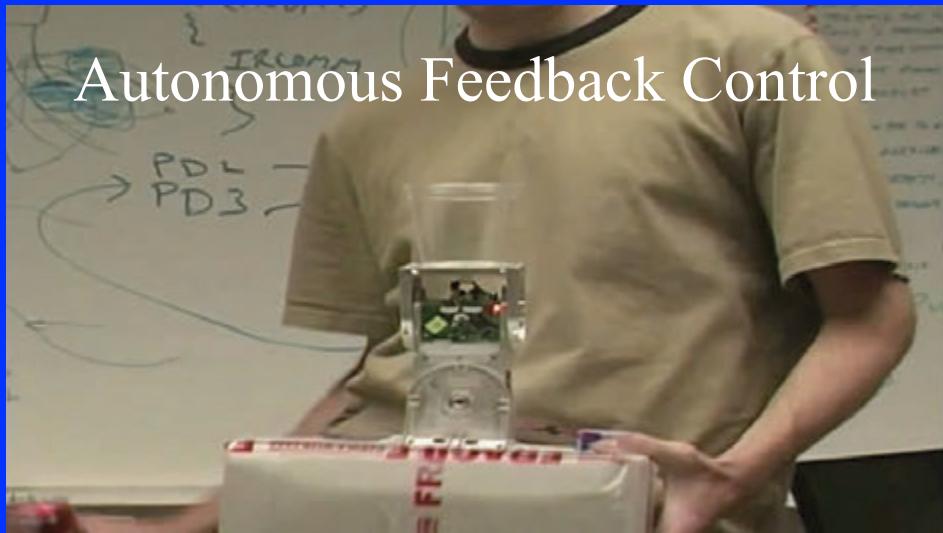
Diversity



Durability



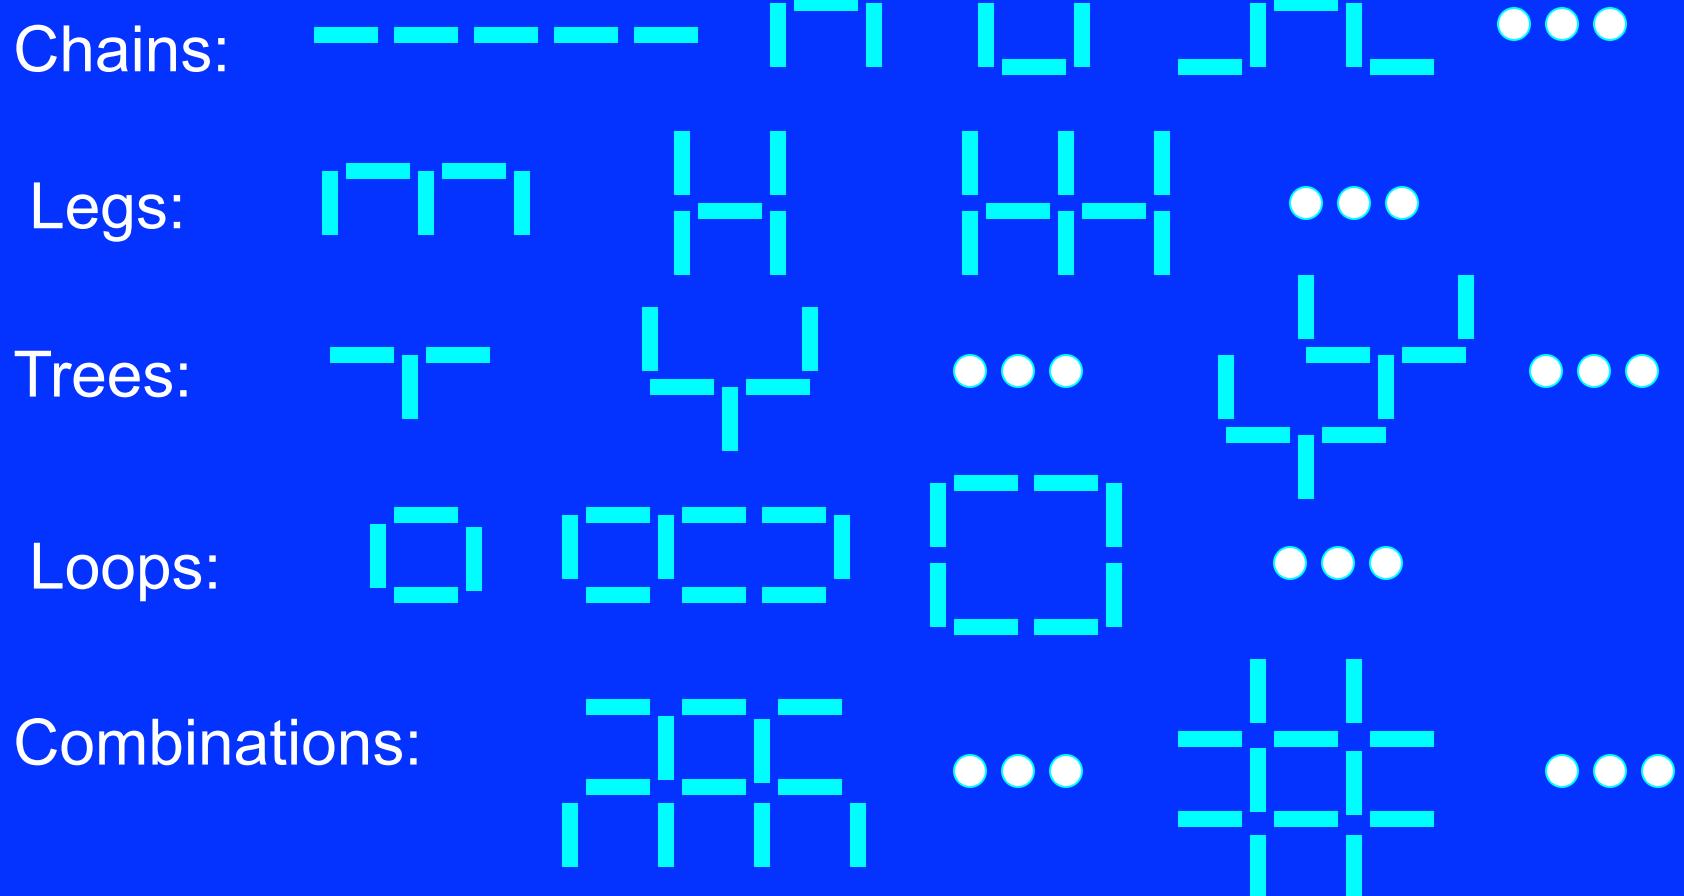
Autonomous Feedback Control



Synchronization



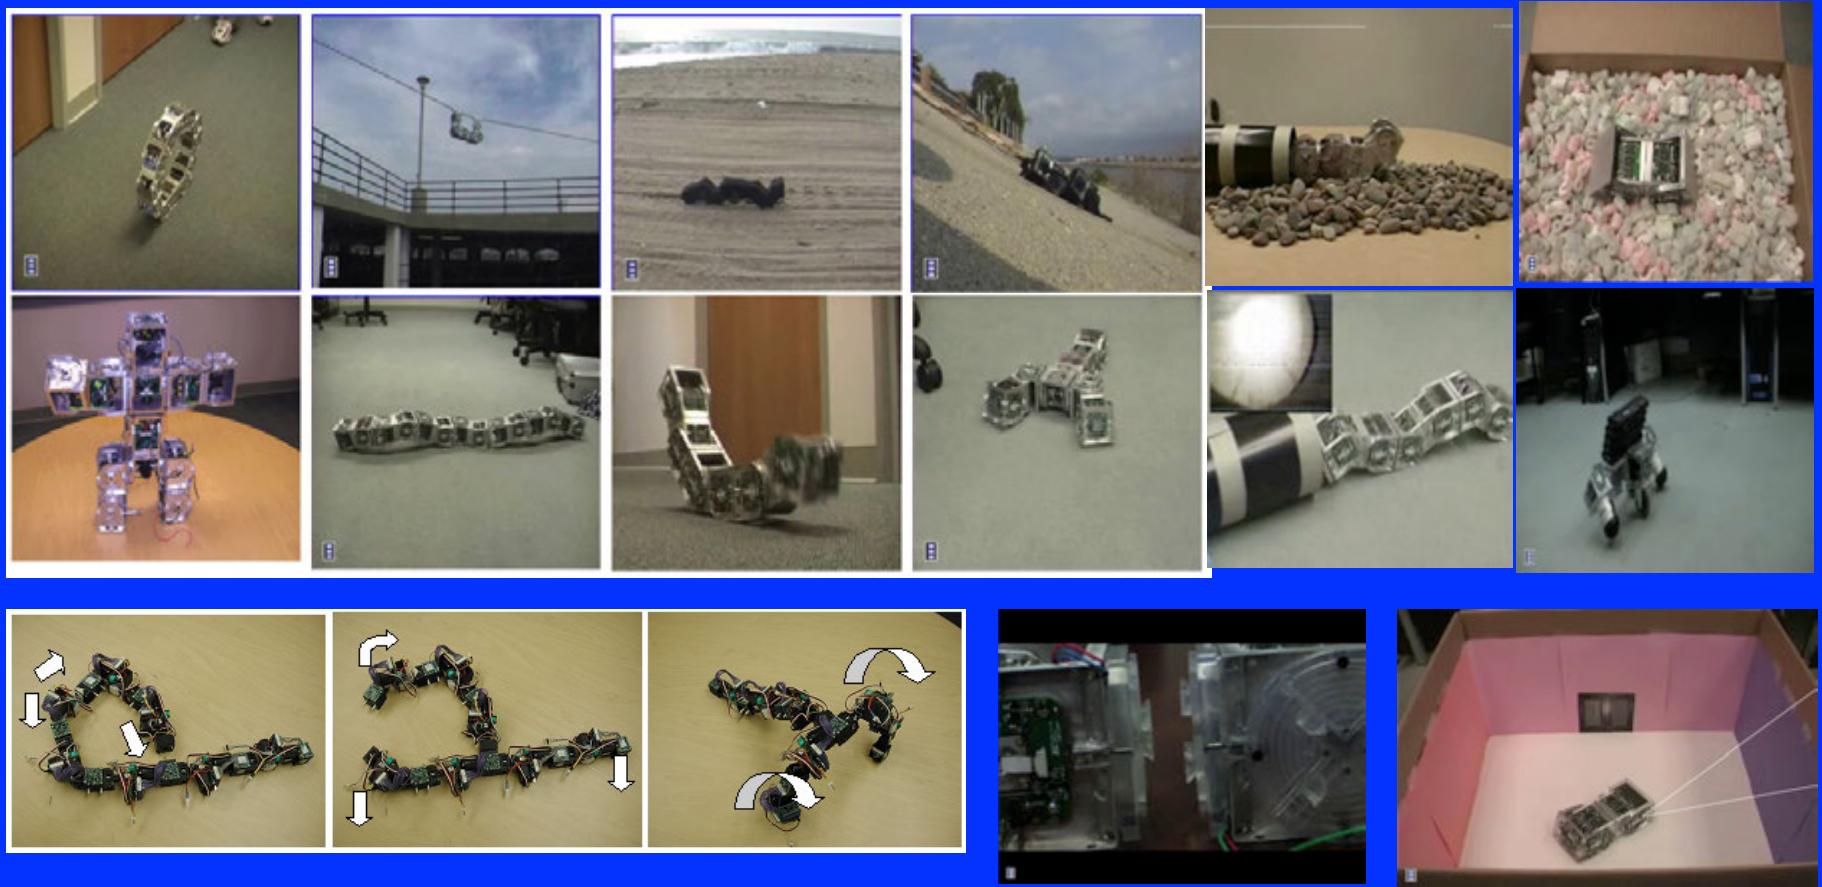
Possible Configurations



Modules in Configurations



Multifunction, Reconfiguration, Distributed Control, & Learning



Caterpillars on Beach



Locomotion Forms



Sidewinder

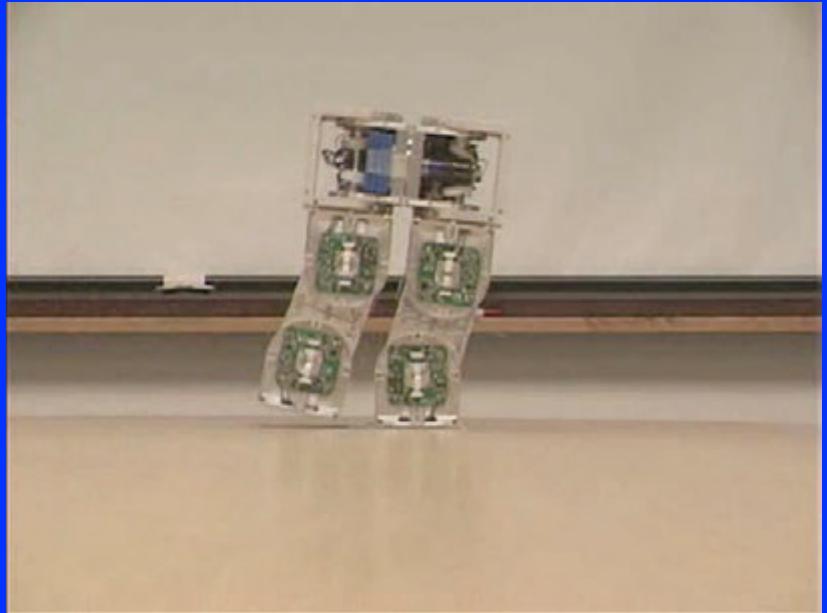
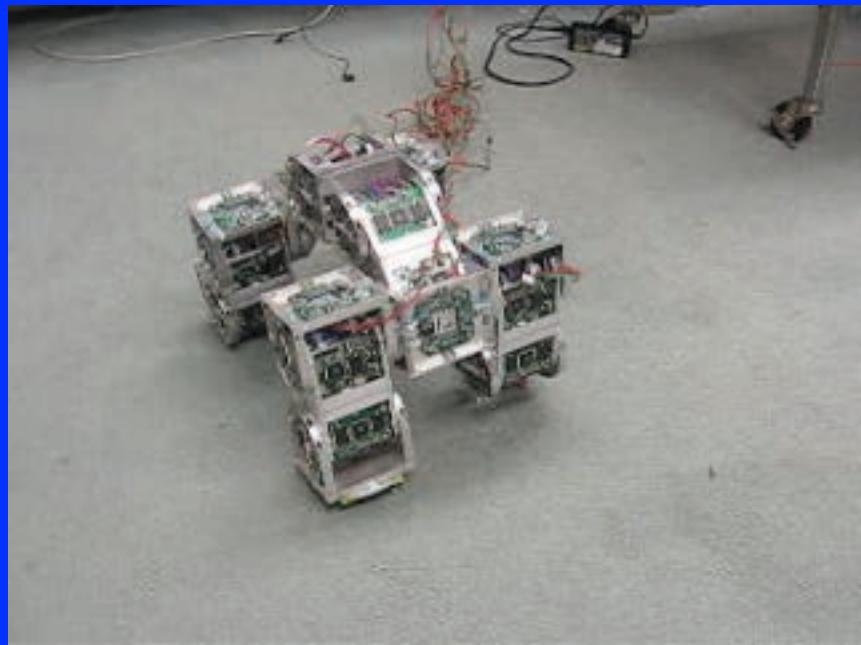


Butterfly



Scorpion

Walking with Legs



Digging and Burying



Going through Pipes



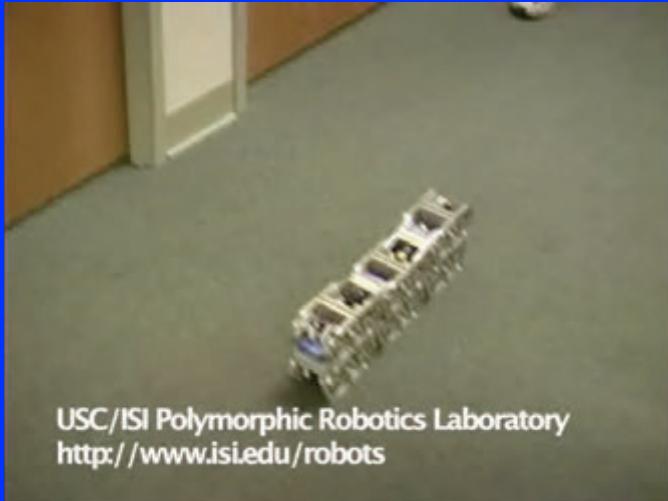
Crawl up Steep Slopes



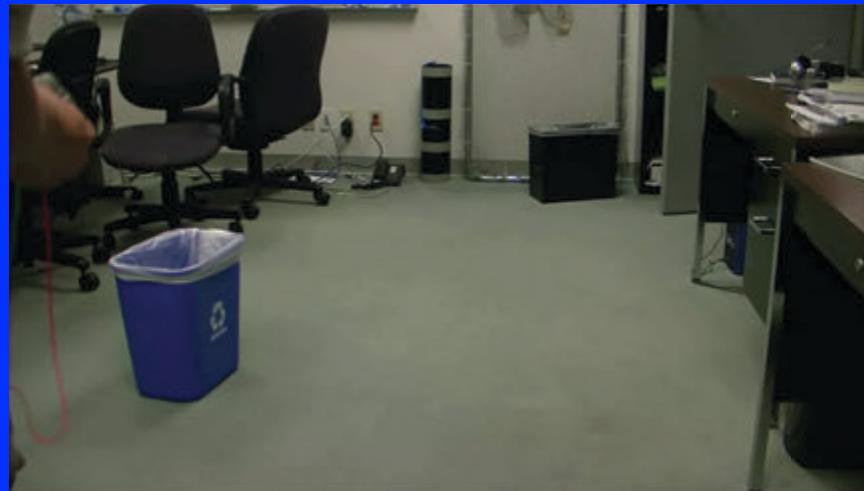
Sand Dune Experiments



Rolling, Recovering, & Turning



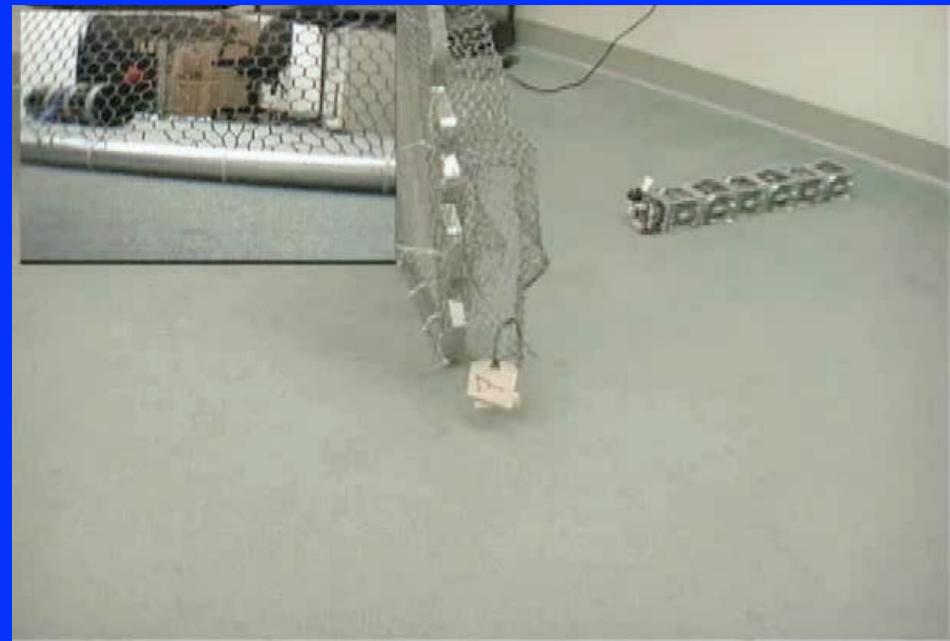
USC/ISI Polymorphic Robotics Laboratory
<http://www.isi.edu/robots>



Climbing Ropes



Tele Operations



Roller Skates and Payload



Rolling on Disks



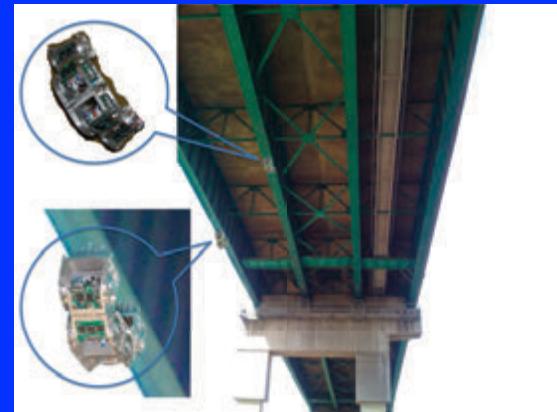
Self Reconfiguration



Resilient to Shape Changes

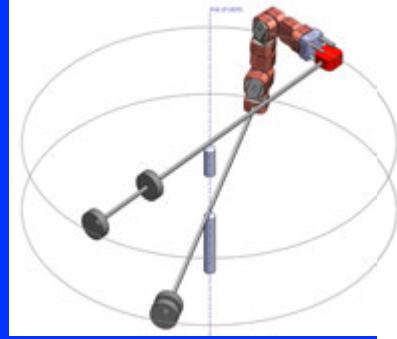
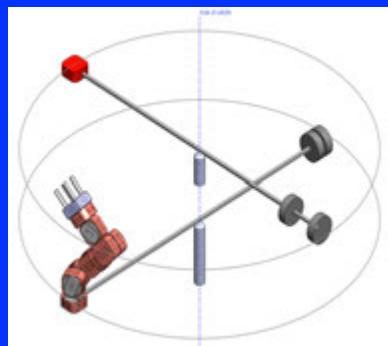
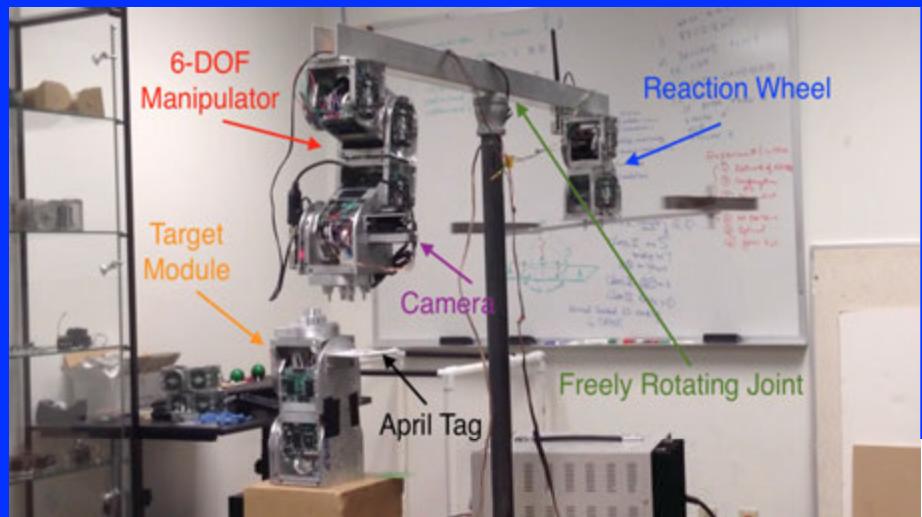
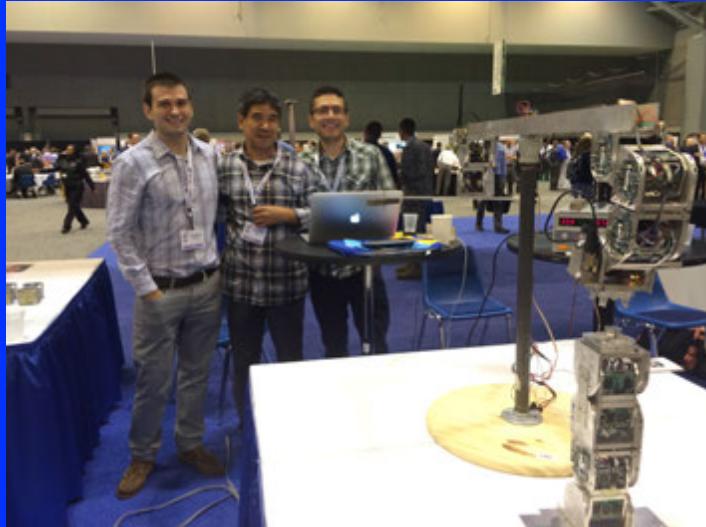


Infrastructure Inspections

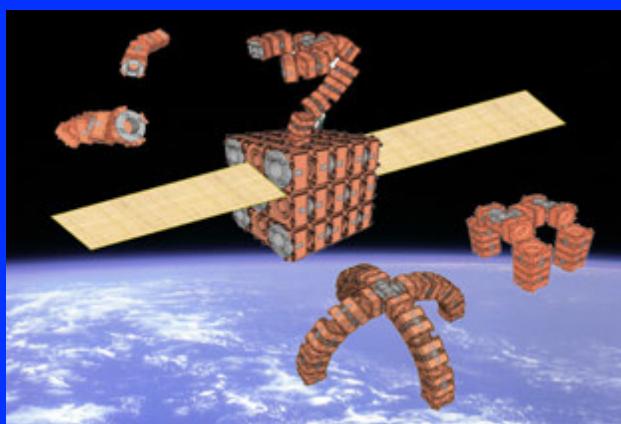
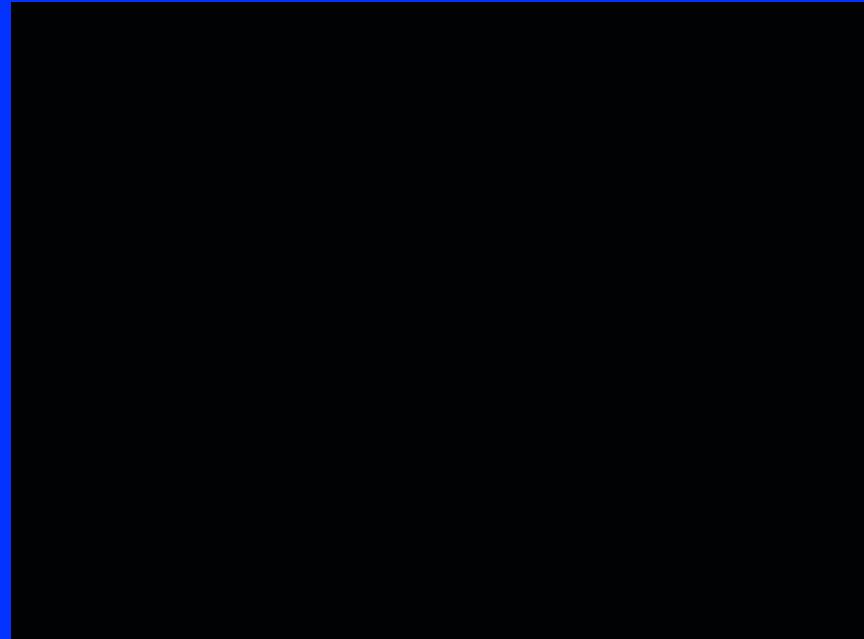
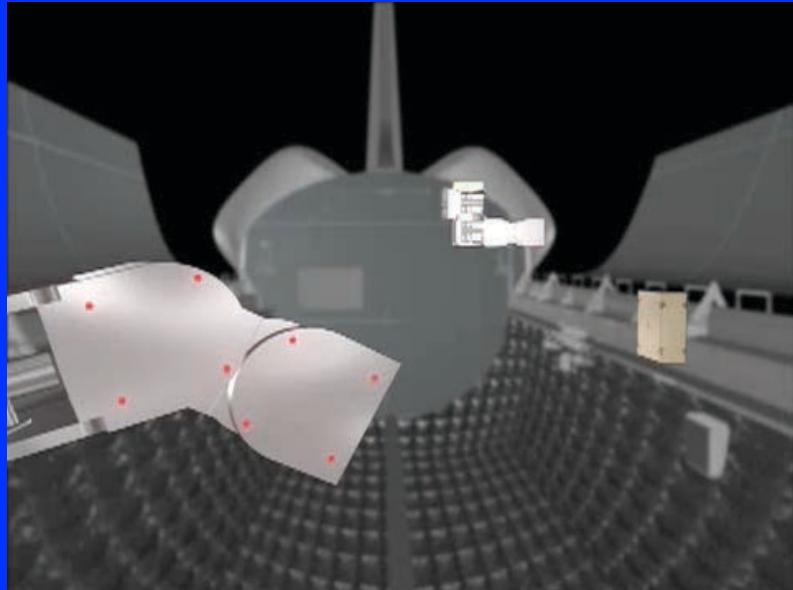


6D Autonomous Docking (in ~0G)

Featured at DARPA Wait-What Forum 2015



Future Operations in Space



Research Challenges

- How to make modules?
 - Modularity, nano, fluid, ...
- How to control a configuration?
 - One structure but many uses
- How to adapt configurations?
 - Many structures, adaptive



Advantages

Flexibility, Reusability,
Repair-ability, Low-cost, ...

Challenges in Systems

- **Challenges in Mechanics**

- Modularity
- Homogeneous and/or heterogeneous modules
- Flexible, connectable, strong, small/large, ...
- Robust, durable, replaceable, ...

- **Challenges in Connectors**

- Genderless
- Single side connectable/releasable
- Thin but strong profile (no back-drive)
- Efficient when connected and disconnected
- Support multi-orientation engagements
- Tolerant to misalignment
- Communication
- *Power sharing*
- *Dirt tolerant*

- **Challenges in Electronics**

- Sensors, actuators, communication, control, interferences, collaboration, ...
- Real-time operating system
- Programming environment
- Debugging facilities
- Power sharing
- Integration of all these!

- **Challenges in Intelligent Control**

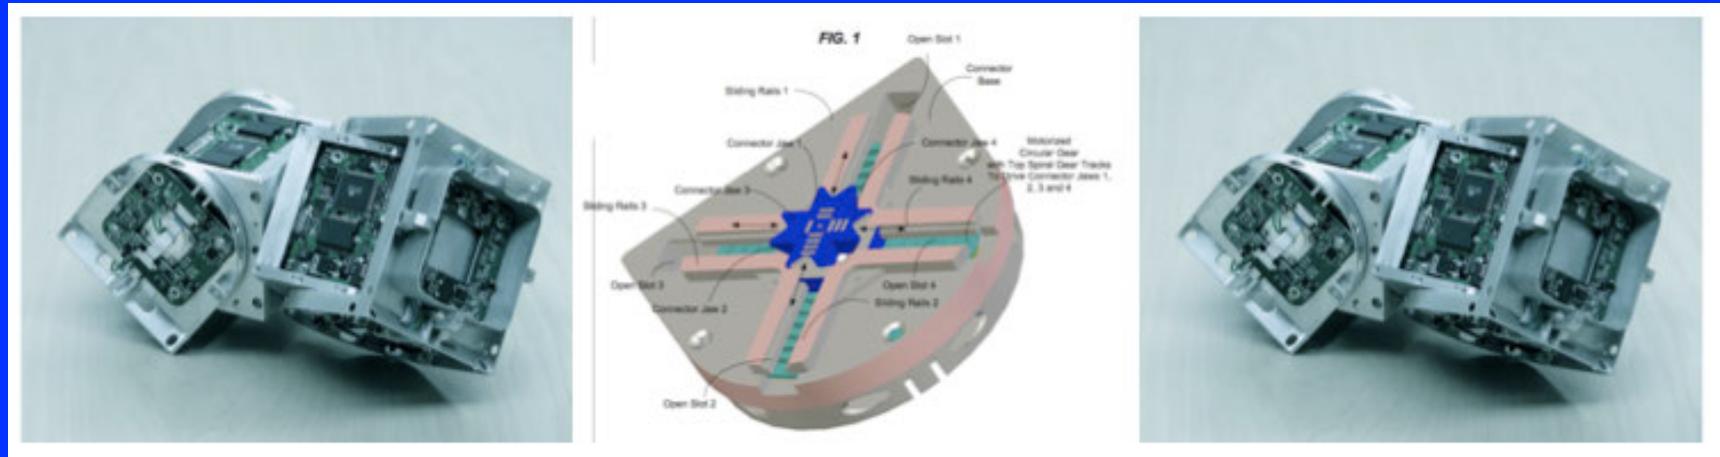
- Intelligent behaviors
 - Mobility adaptive to tasks/environment
 - Maneuverability adaptive to tasks/environment
- Adaptive shapes and structures
 - Higher level challenges for AI & Robotics
 - Self-organization
 - Self-healing
 - Reconfiguration deciding, planning, execution
- New theory for self-reconfigurable systems

Challenges in Connectors

- Genderless
- Strong mechanical torque and endurance
- Data communication
- Power sharing
- Thin profile <15mm
- Dock-able between variable sizes
- Single side operational and releasable
- No back drive
- Efficient power consumption for connect/disconnect
- Multi-orientation engagements
- Tolerant to misalignment during docking process

SINGO Connectors

(US Patent 8,234,950)



A movie of
connecting and
disconnecting

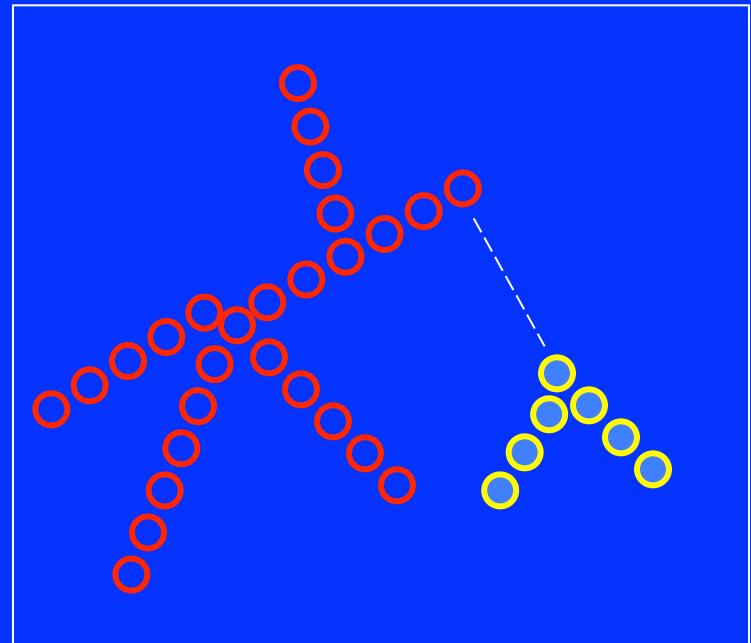
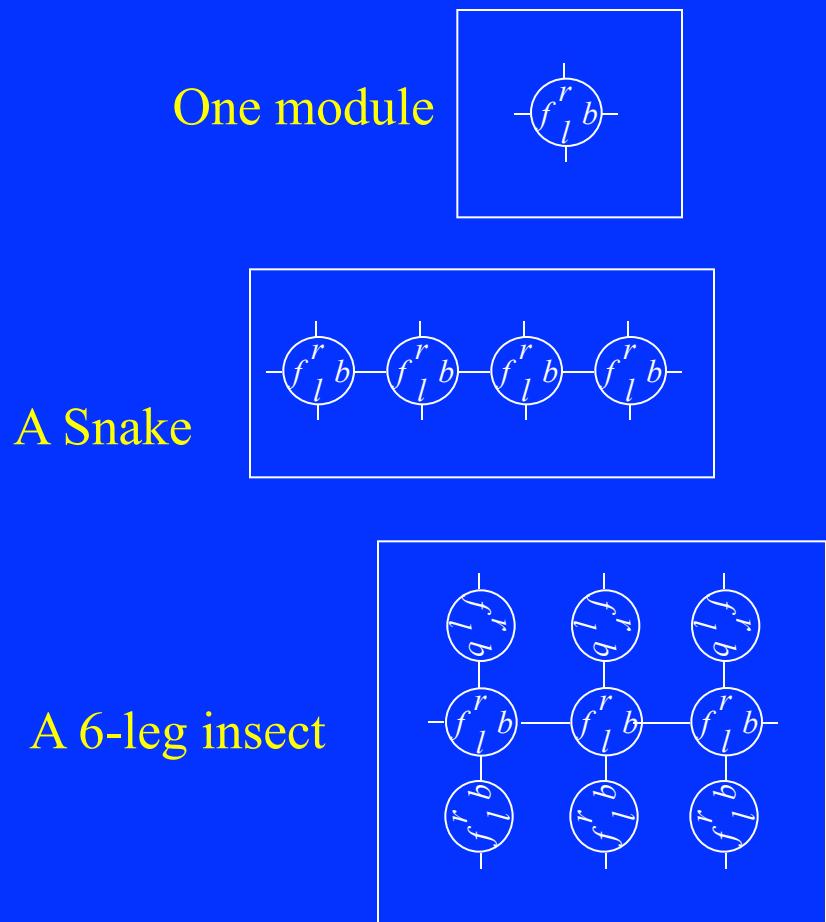
Challenges in Distributed Control

- Distributed among all modules
 - No “fixed brains” or single-point-failures
 - Free from any unique global names or identifiers
 - Discover topology and negotiate tasks
 - Change behaviors by topological locations
- Dynamic network topology
 - Endure unexpected changes in the network
- Asynchronous
 - Collaborate without any synchronized or global clocks
- Scalable
 - Function with any size and shape
- Adaptive
 - Perform complex tasks in different environments
 - Recover from component failures and self-repair
 - Autonomous learning for unexpected situations and tasks

“Digital Hormones” (US Patent 6,636,781)

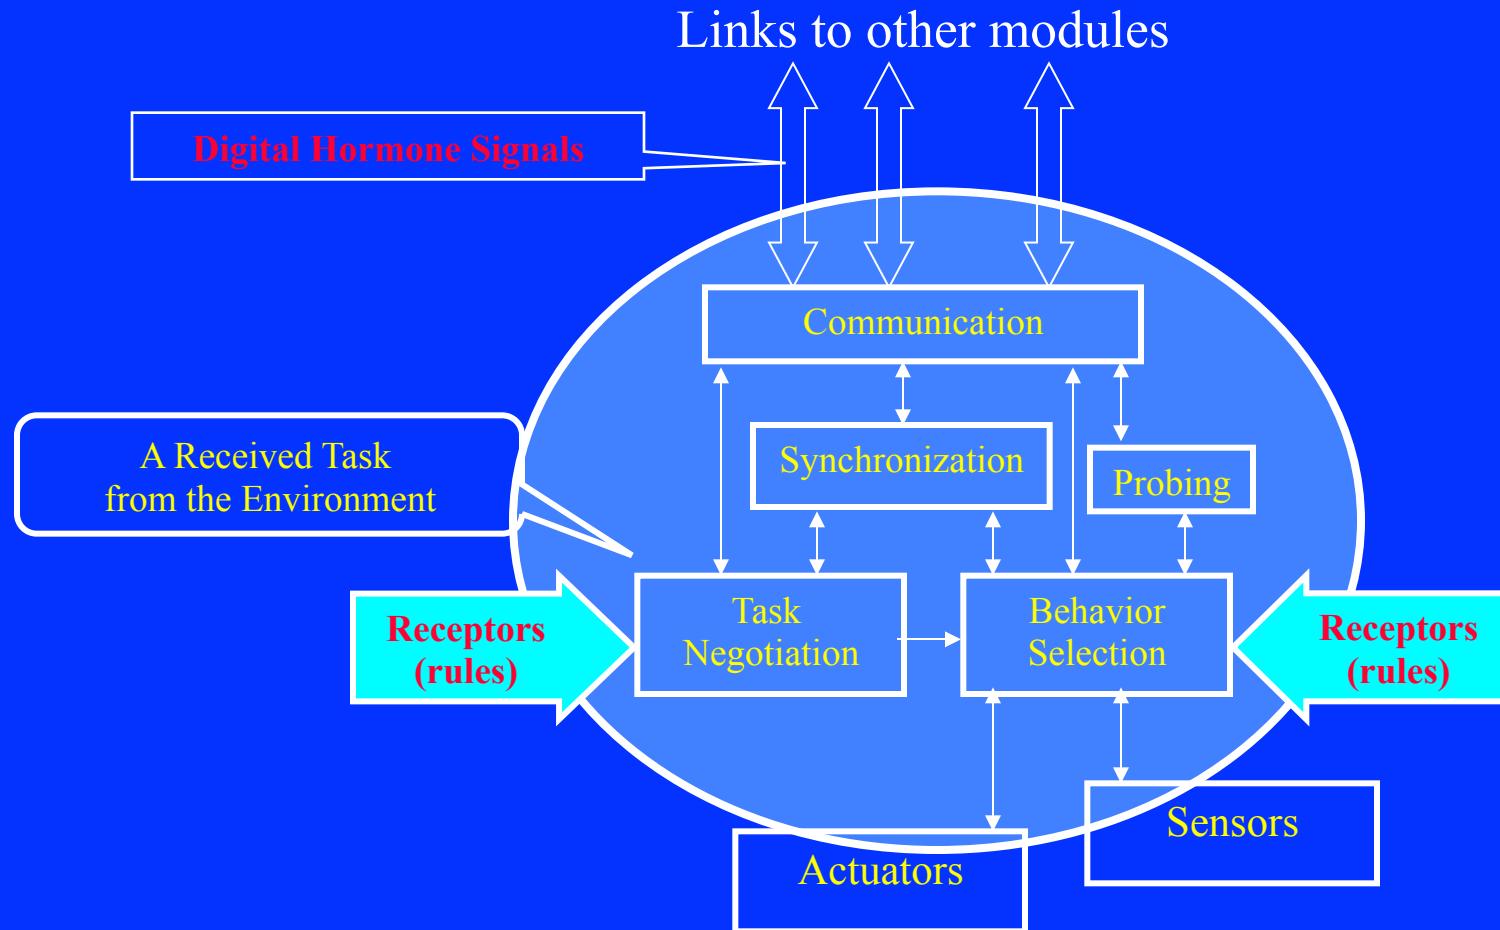
- Content-based messages
 - No addresses nor identifiers
 - Have finite life time
 - Trigger different actions at different sites
- Floating in a global medium
 - Propagated, not broadcast
 - Different from *pheromones*
 - Internal circulation, not external deposit
- Preserve local autonomy for individual sites
- Could modify module behaviors (as RNA)

Configuration Representation



Communication between
two separate creatures

Internal Control Architecture

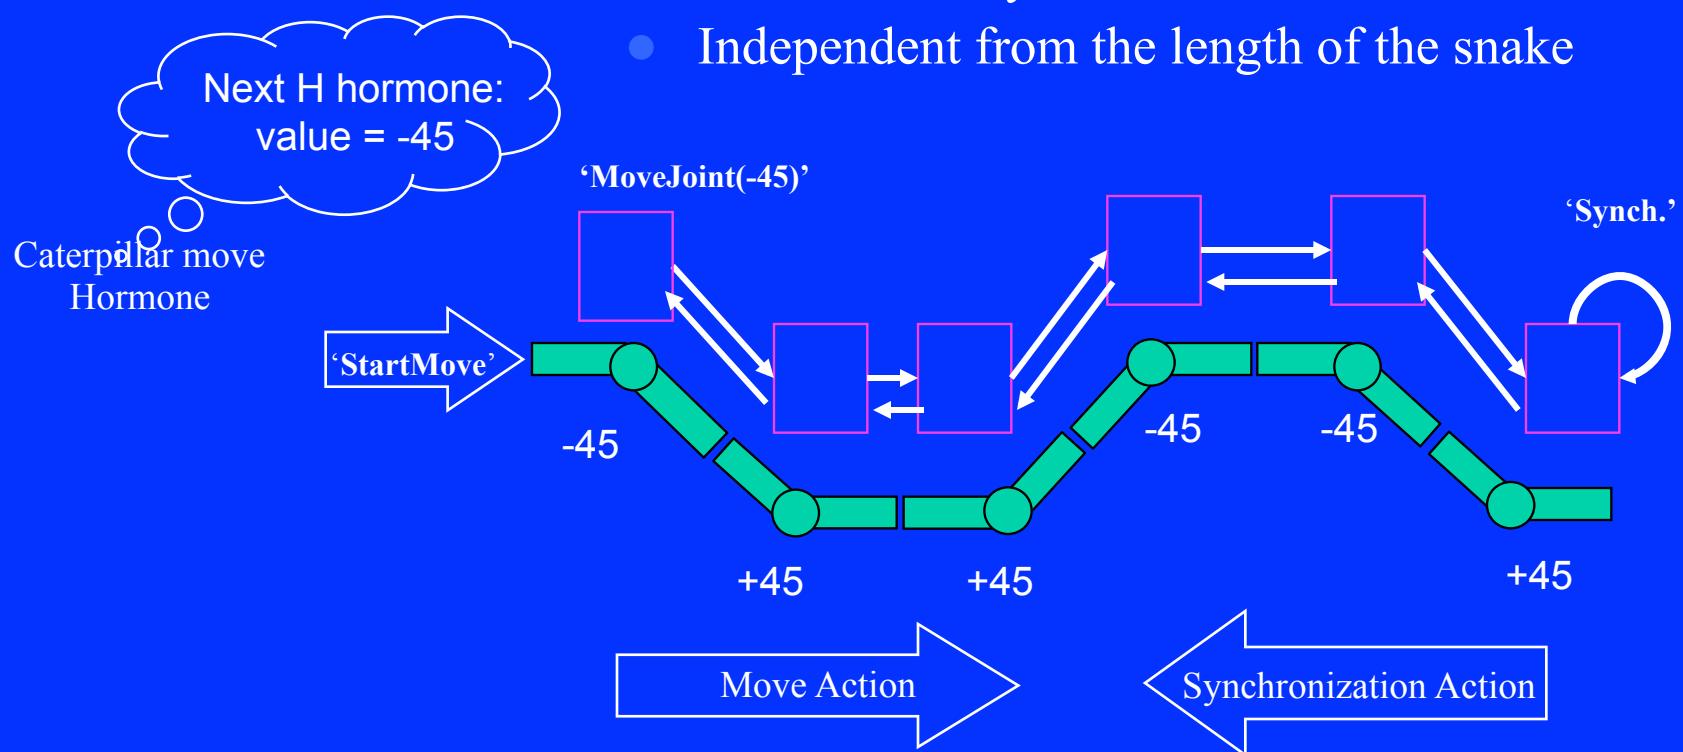


The Uses of Digital Hormones

- Communication in dynamic network
- Cooperation of autonomous robotic modules
 - Locomotion
 - Reconfiguration
 - Synchronization
 - Global effects by weak local actions
 - Conflict resolution (multi hormone management)
 - Navigation
- Configuration adaptation (vision, telescope, ...)
- Self-healing systems

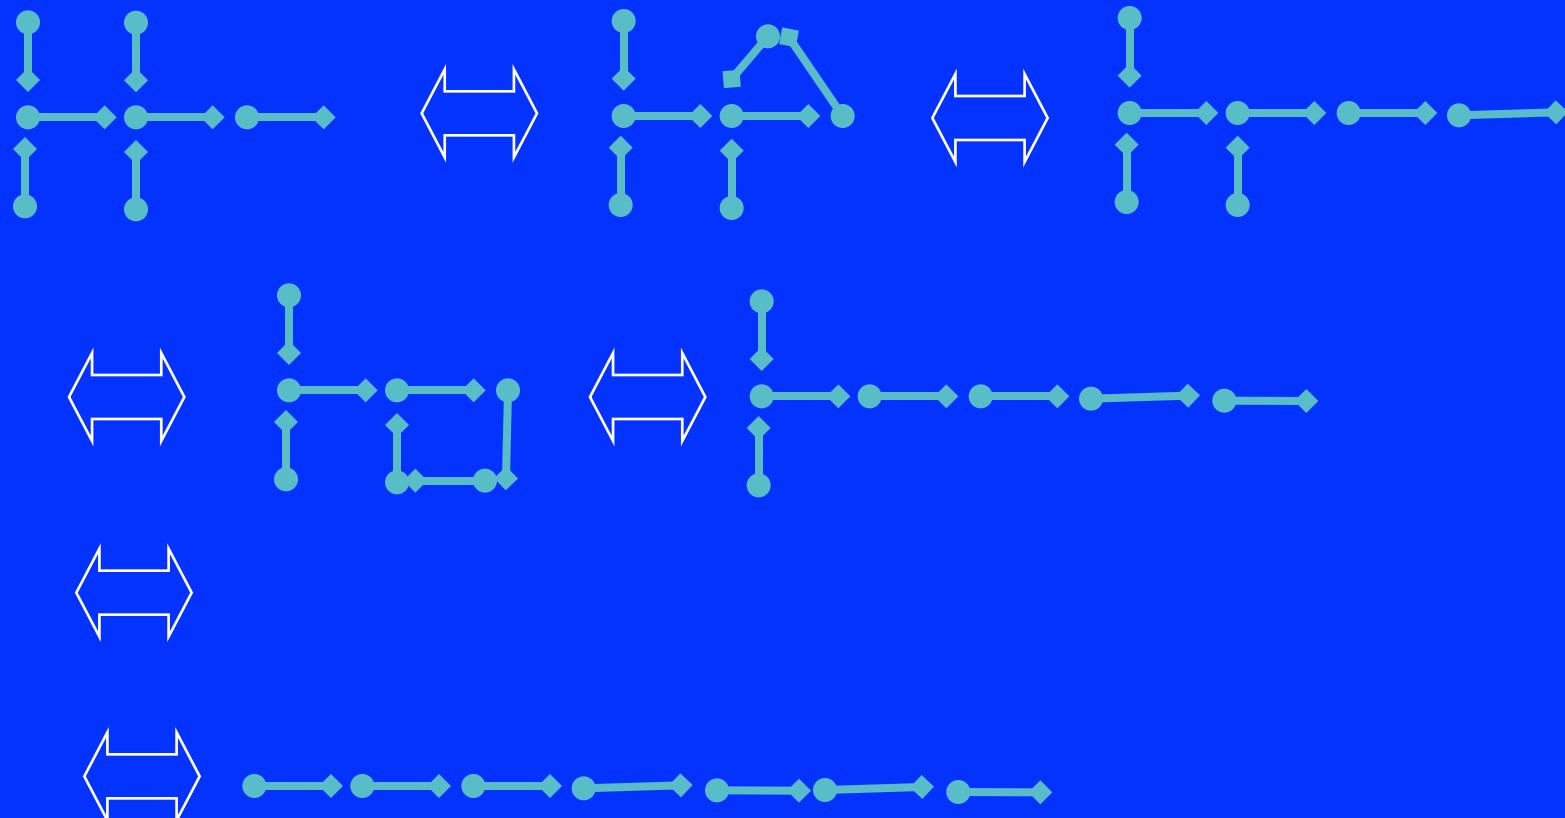
Hormones for Caterpillar Move

- A simple one-pass hormone from head to tail
- Controls and synchronizes all motor actions
- Independent from the length of the snake

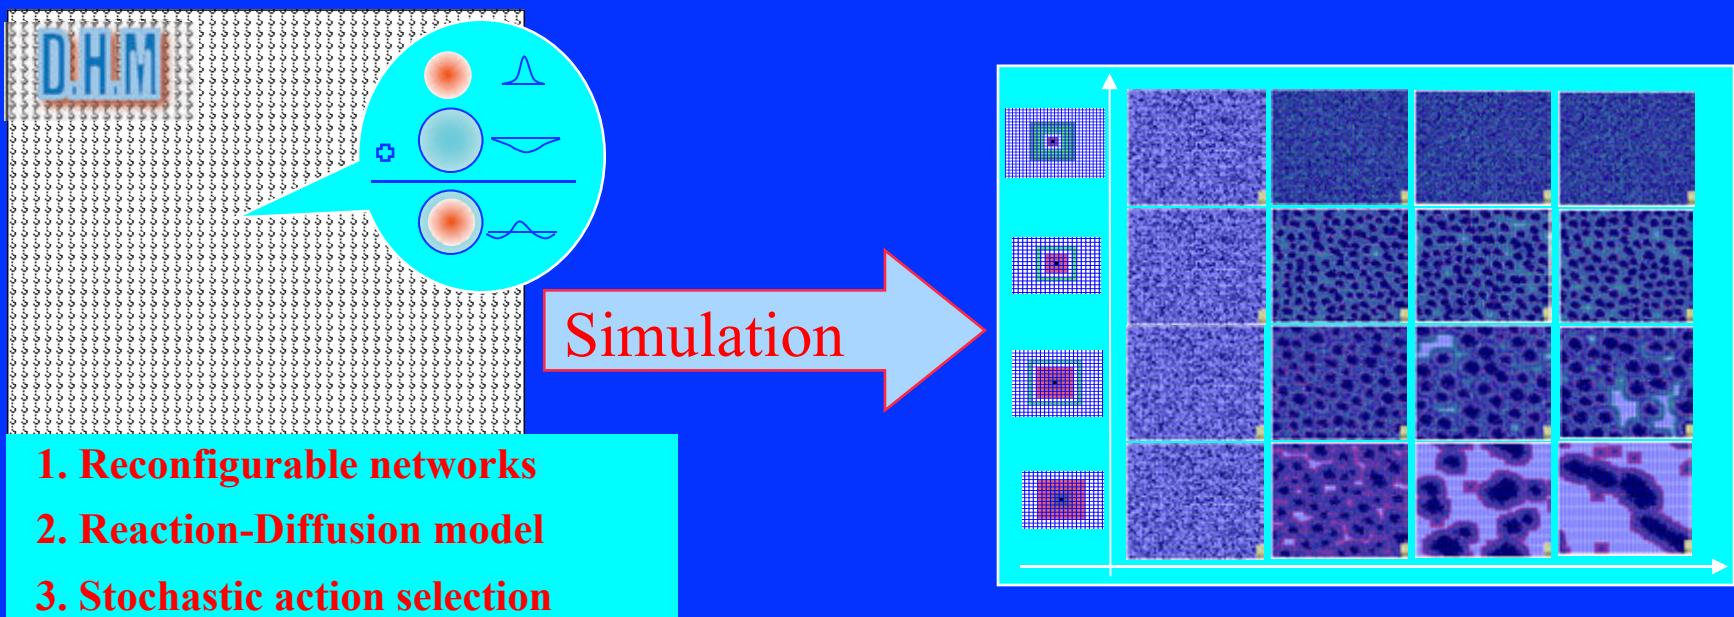
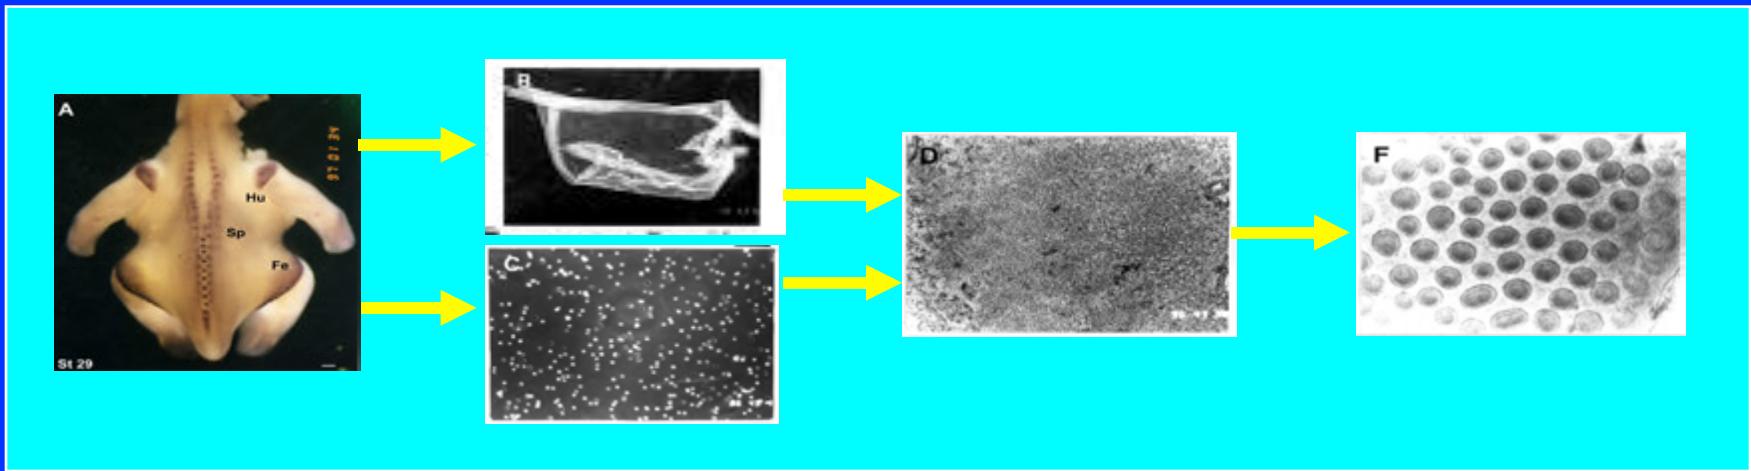


Hormones for Reconfiguration

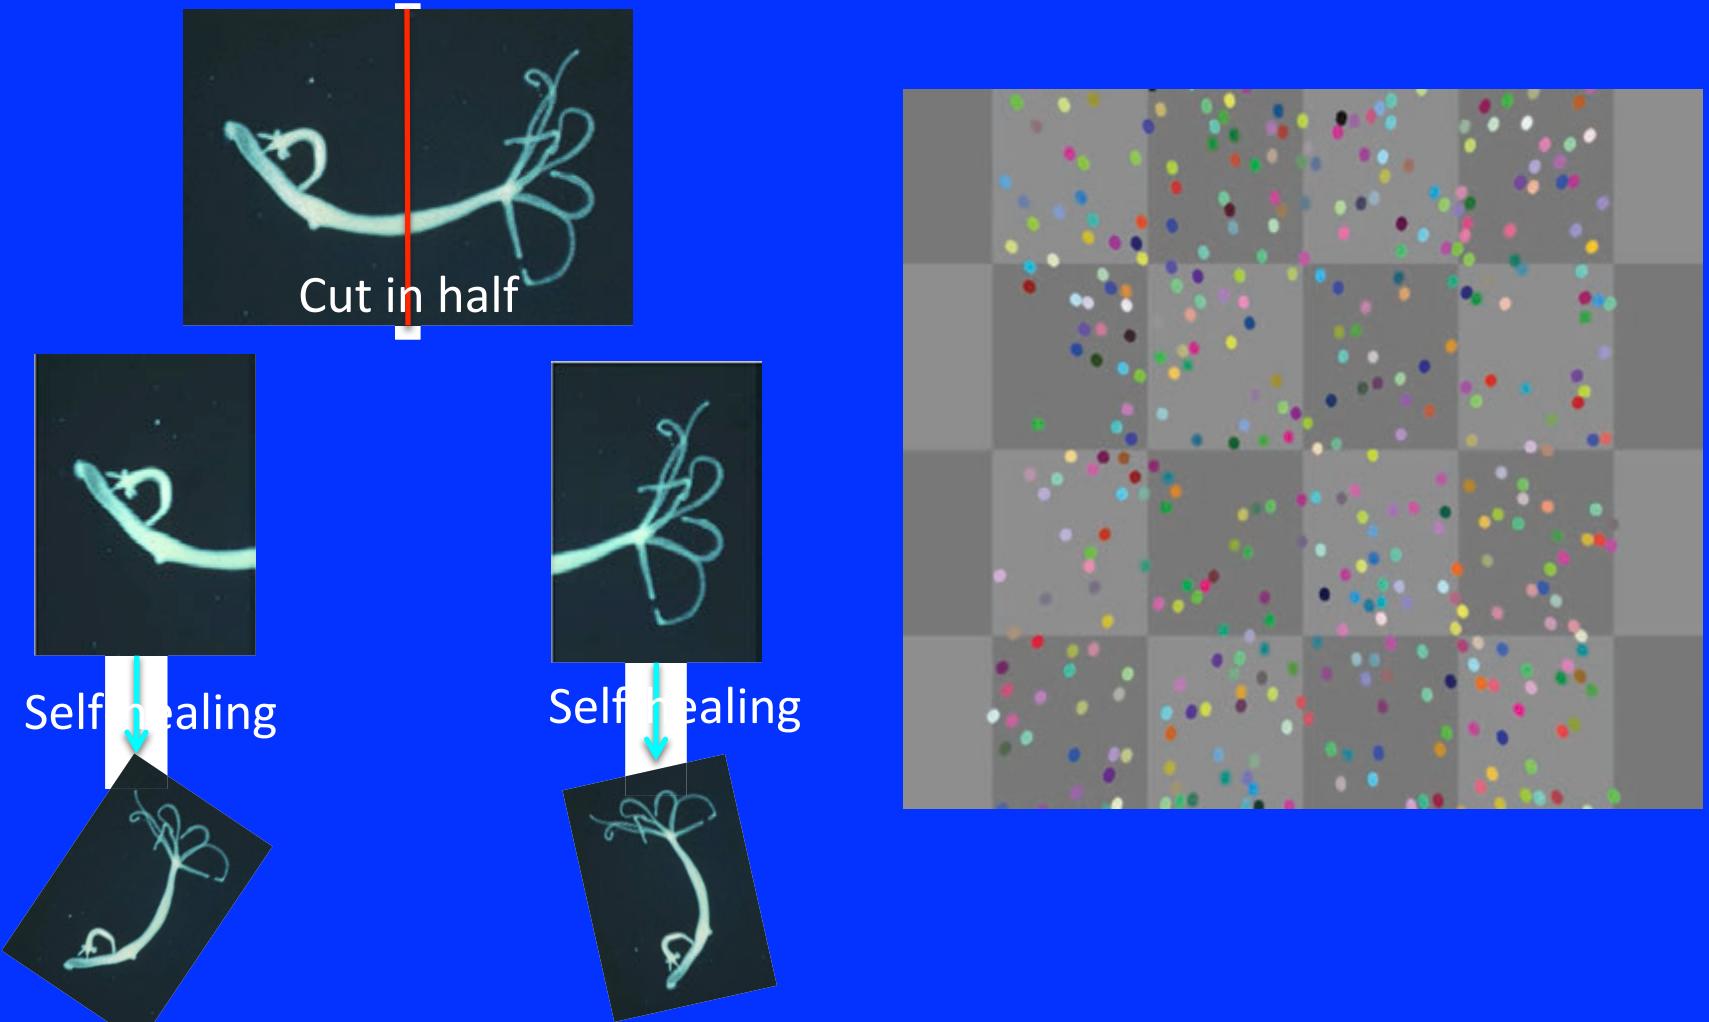
Insect \leftrightarrow Snake



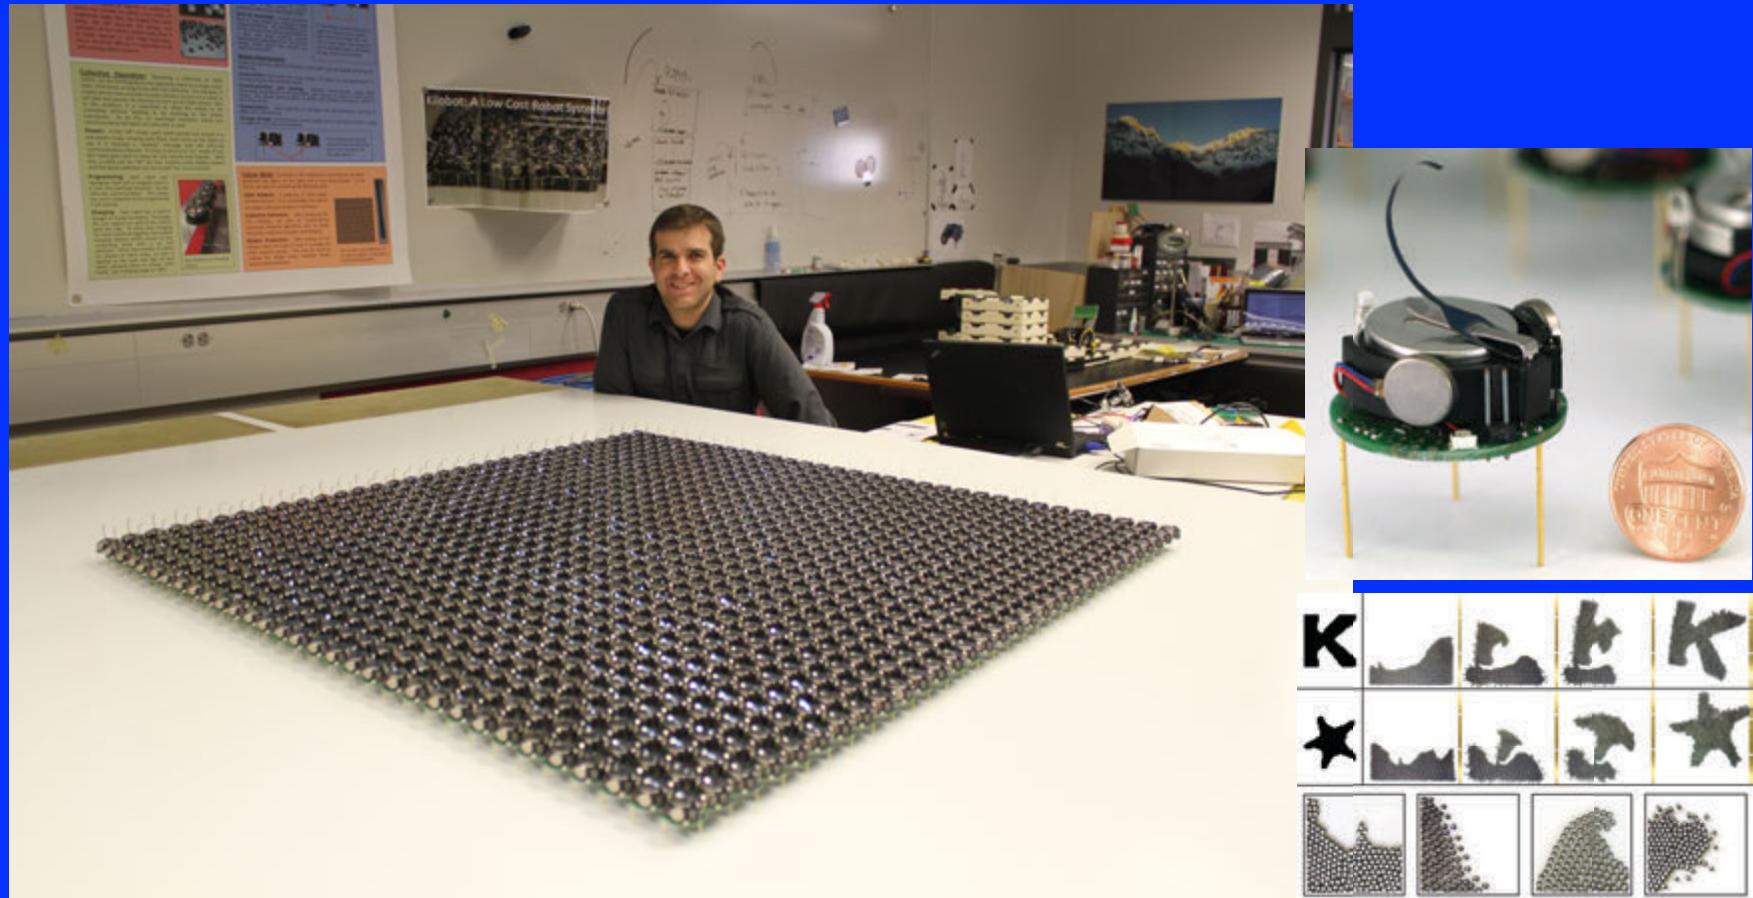
Hormones in “Feather Formation”



Scalable Self-Healing



Robotic Swarms



UAV Swarm and Applications

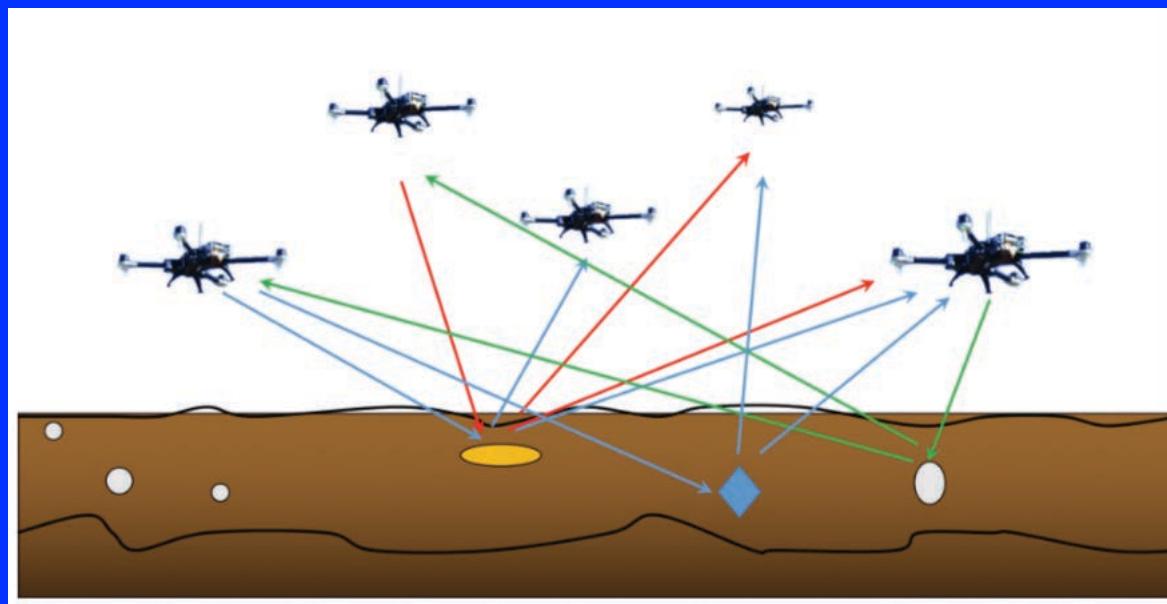
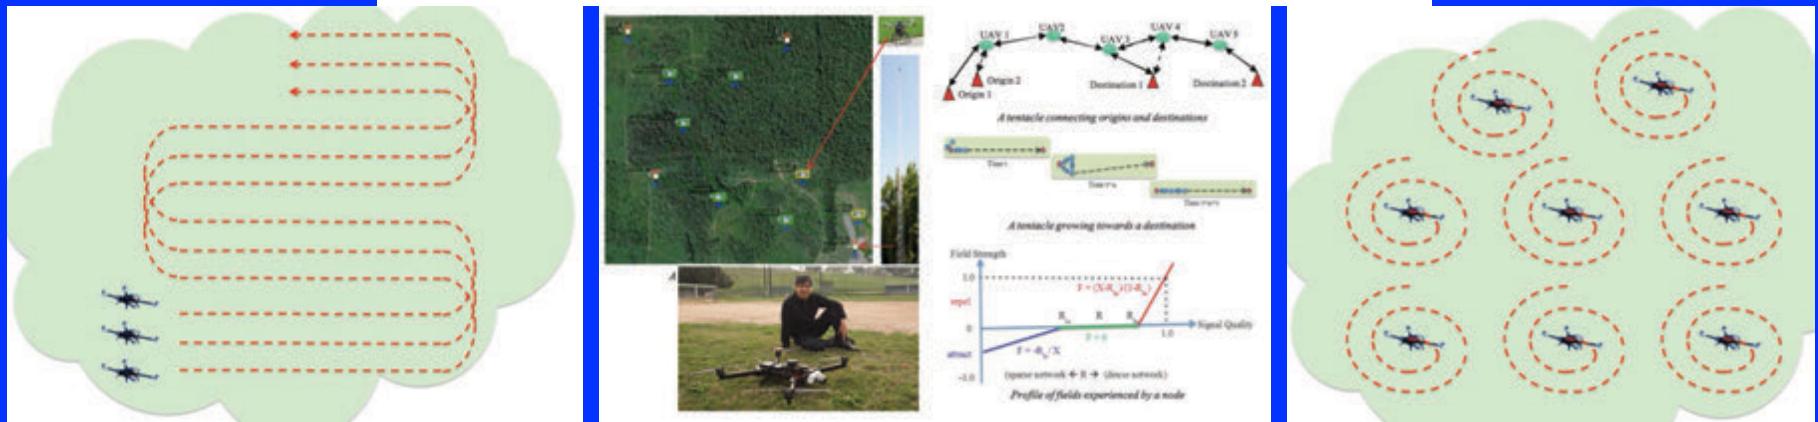


UAV Swarms

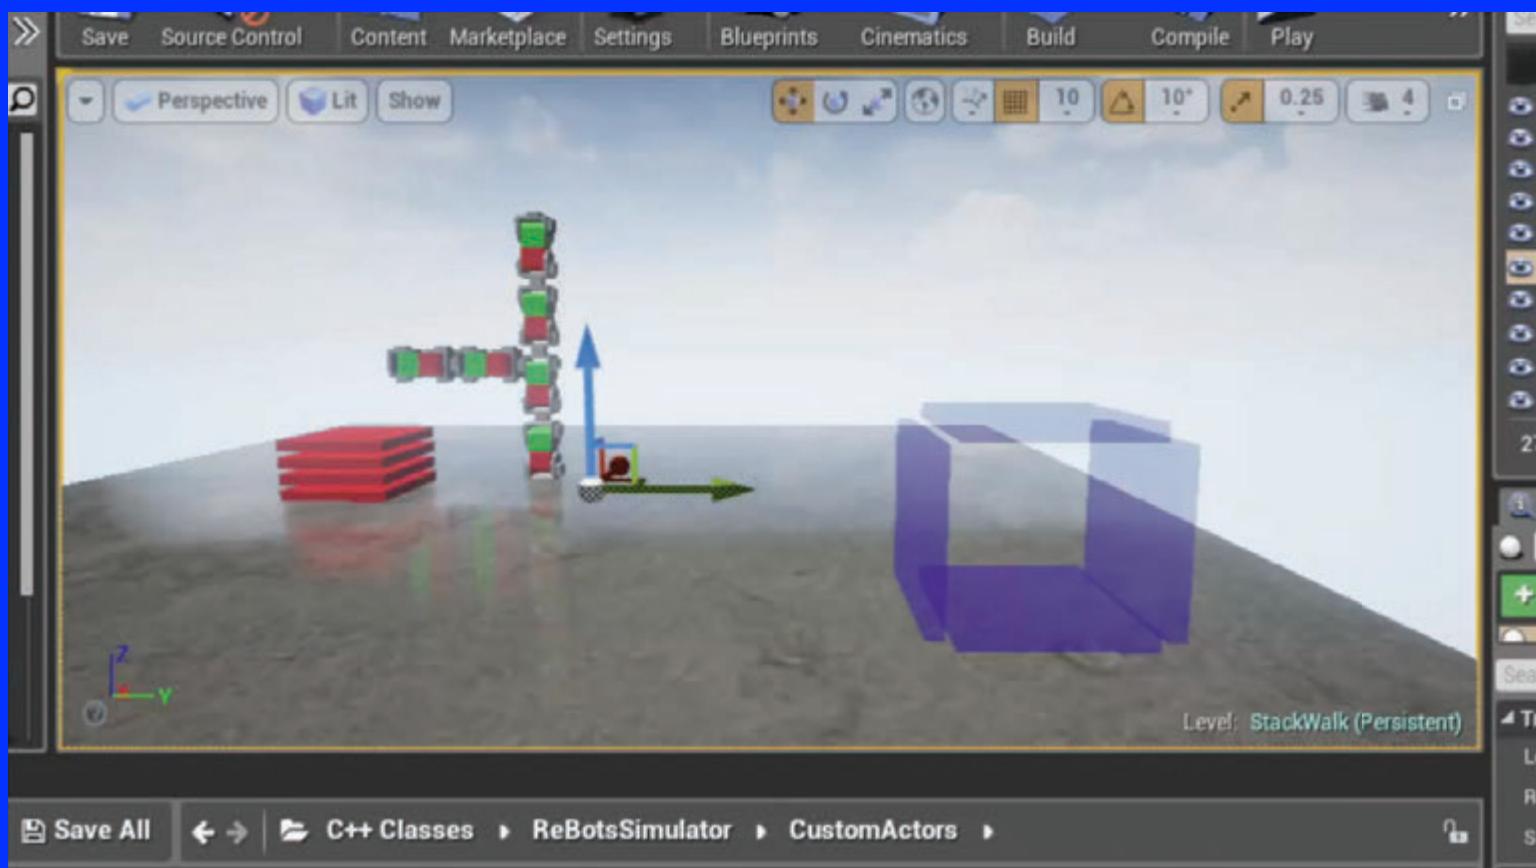


Mine Detection by UAV Swarms

(Masri, Moghaddam, Shen)



Self-Assembly

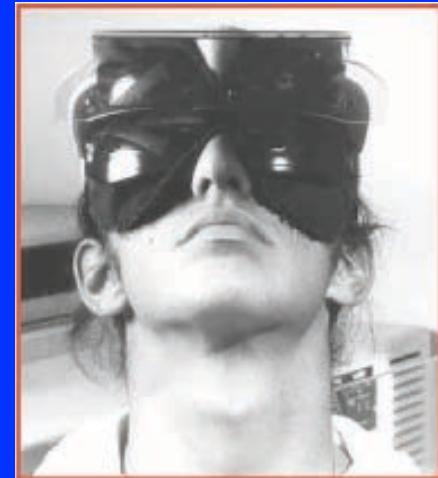


Very-Large Physics-Based Simulation

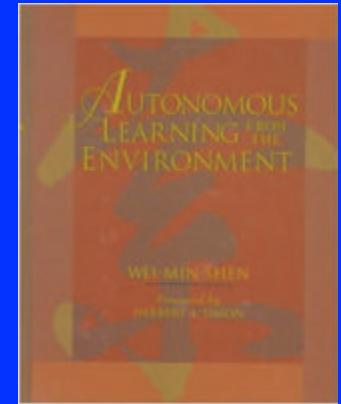
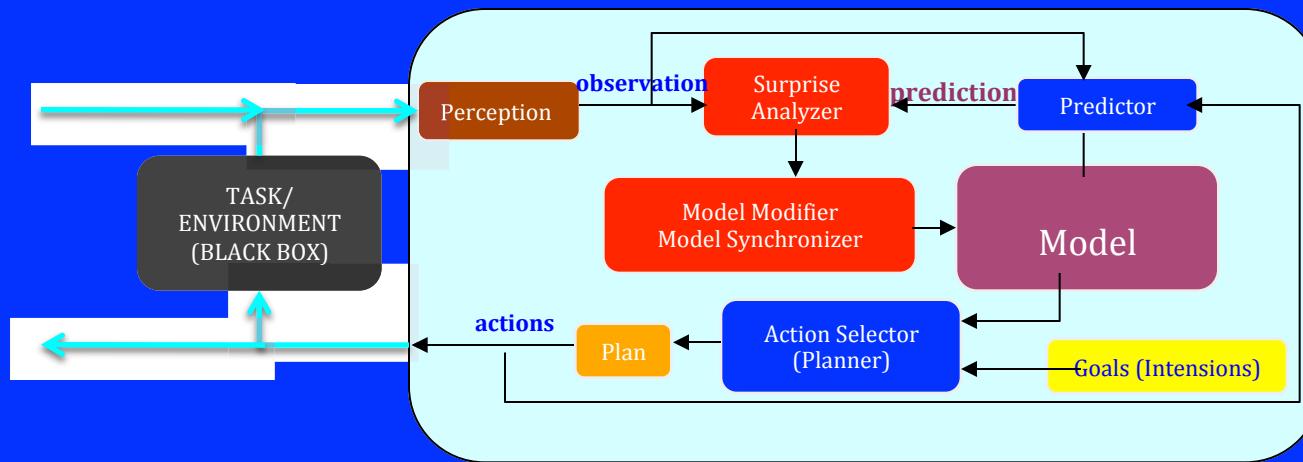


Surprise-Based Learning

- New Environment/Task → Learning → Knowledge/Skill
- Key Idea: To detect and learn from “surprises”
 - Adaptive to new environments
 - Learn to accomplish new tasks
 - Self-heal from unexpected failures or dynamics (e.g., inverted visions)
 - Know-how → Surprises → Learn → Recovery



Surprise-Based Learning



- The Learner continuously makes predictions, detects surprise, analyzes surprises, extracts critical information from surprises, and improves and uses its action models

Surprise ==> Model ==> Prediction



Related Research Topics

- Biological morphogenesis
- Self-organizing swarm robots
- Self-healing (scalable) systems
- Underwater robot swarm
- Self-assembly in space
- Oil and energy robotics
- Self-configureble network
- Evolution of brains

Conclusions

- Self-Reconfigurable Robots Are Coming!
 - Fundamental research challenges
 - Physics, chemical, biology, structures, robots, networks, agents, ...
 - Offer critical capabilities for many applications
 - Search/rescue, homeland security, self-assembly, smart structures, space, ocean, deep-water, oil production, ...
 - Has many future potentials
 - Autonomous, Resilient, Adaptive, and Learn
 - Intelligently selection of configurations (very challenging)

CSCI 561

Foundation of Artificial Intelligence

Professor Wei-Min Shen
(Dr. Nadeesha Ranasinghe)

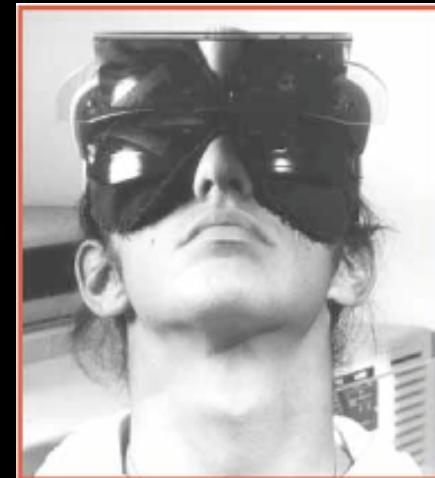
Session 25-1: Surprise-Based-Learning-1

Unsupervised Learning

- Type I: Clustering the data
 - Automatically group the data into clusters
- Type II: Parameter Learning (states are known)
 - Learn transitions, sensor models, & current state
- Type III: Structural non-parametric Learning
 - States have internal structures (not just “symbols”)
 - States are not known and must be learned

Why Surprise-Based Learning?

- New Environment/Task → Learning → Knowledge/Skill
- Key Idea: To detect and learn from “surprises”
 - Adaptive to new environments (may have no priori knowledge, “swim”)
 - Learn to accomplish new tasks (goals may change dynamically)
 - Self-heal from unexpected failures or dynamics (e.g., inverted visions)
 - Know-how → Surprises → Learn → Recovery

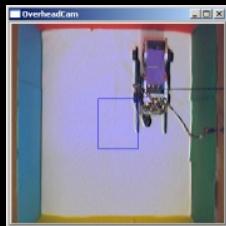


Technical Challenges (why hard?)

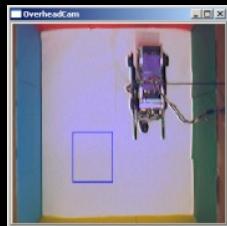
- Changes and mistakes are not easily detectable
 - Changes may not be predicted and occur simultaneously
 - Uncertainties in sensors and actions exist, not fault-free
 - Sufficient redundancy of components cannot be guaranteed
- Lack of complete and permanently correct knowledge
 - Sensor, action & environment models
 - Incomplete initial knowledge
 - Accidents happen
- Causes of interference may or may not be identifiable
 - Learn to reason with interference with/without adaptation
- May not have external supervision to detect changes
- Requires fast response times
 - Expensive solutions for memory/resource may not be feasible
- Must deal with both discrete and continuous data, uncertain and vast information space

Potential Unpredicted Changes

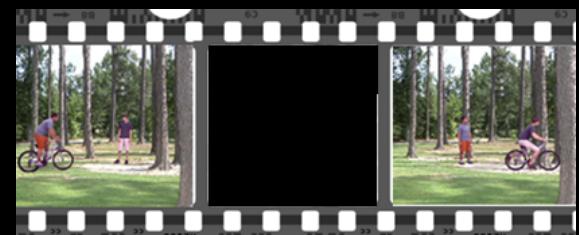
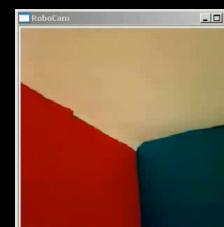
- Uninformed changes to the robot at runtime
 - Sensors: addition, deletion, definition-change, interference
 - Actions: addition, deletion, definition-change, interference
 - Goals: addition, deletion
 - Dynamic configurations of the environment
 - Objects: addition, deletion, relocation
- “Definition-change” alters mapping of input to output
 - Rotation, Translation, Scale
- Interference includes noise & missing data or gaps



Goal change



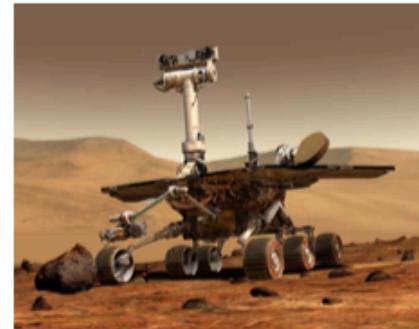
Rotated camera



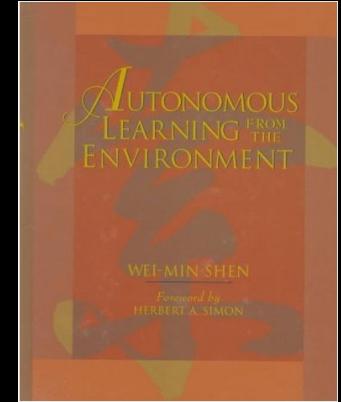
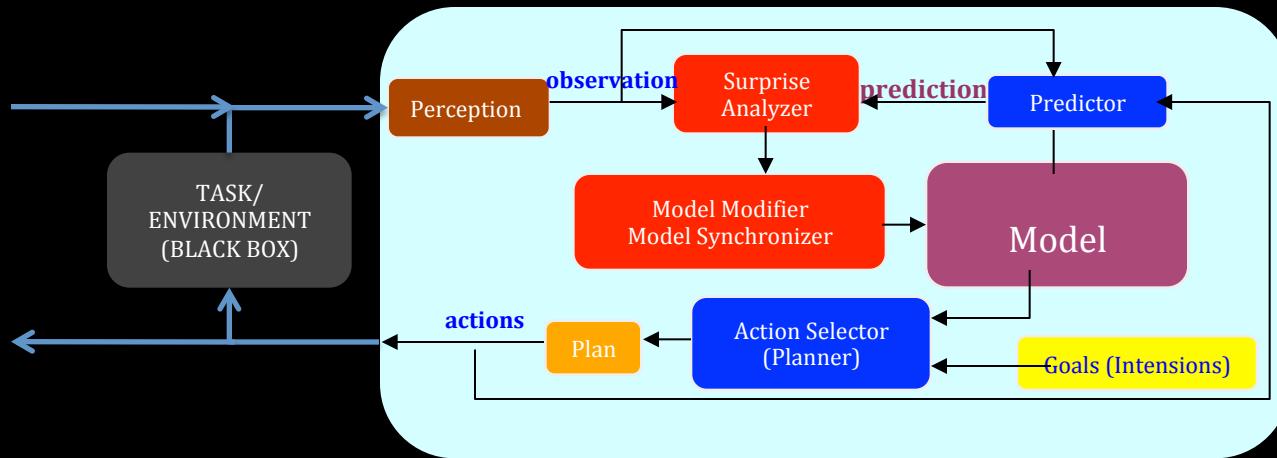
Gap

Surprises and Their Sources

- What is a Surprise?
 - A significant discrepancy of mental predictions and sensor measurements
- Where are its sources?
 - Incomplete initial knowledge
 - Designers cannot capture the richness of the real world
 - e.g. Spirit rover got stuck in soft soil
 - Dynamic nature of the environment and tasks
 - Environment changes
 - System or robot's components change
 - Accidents happen
 - e.g. Spirit's right-front wheel stopped working
 - Lack of critical knowledge in the current KB
 - e.g. Toddlers gradually learn to grab stuff
 - e.g. Things wear out with time



Surprise-Based Learning (CDL++)



- The Learner continuously makes predictions, detects surprise, analyzes surprises, extracts critical information from surprises, and improves and uses its action models

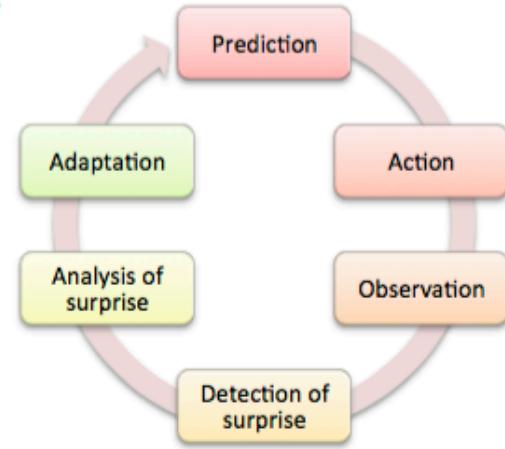


Surprise ==> Model ==> Prediction

Random exploration

The Cycle of SBL (CDL++)

- Autonomously learn **Predictive Rules** from experience
 1. Select **Actions** for the current goals
 2. Make **Predictions** before actions
 3. Perform actions
 4. Observe **events** and **entities** via sensors
 5. Detect **Surprises**
 6. Analyze and explain Surprises
 7. Adapt or create prediction rules
- **Detecting Surprises**
 - When a prediction does not match the actual outcome of an action
- **Explaining Surprises**
 - Finding the differences of a normal satiation and the unexpected situation
- **Revise Prediction Models**
 - Use the found differences to revise the relevant prediction rules
- **Adapt Decisions and Actions**
 - Use the adapted prediction rules to decide actions in future



Example Environments/Tasks

Simulated hunter-goal

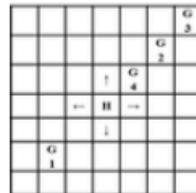
Sensor: location

Entities: H, G

Attributes: x, y, u

Actions: N, S, E, W

Goal: $H(x,y)=G(x,y) \parallel H(u)=G(u)$



Simulated hunter-prey

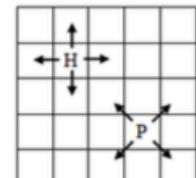
Sensor: location

Entities: H, P

Attributes: x, y, u

Actions: N, S, E, W

Goal: $H(x,y)=P(x,y) \parallel H(u)=P(u)$



Real office room

Sensors: camera

Entities: SURF objects

Attribute: s, x, y

Actions: L, R

Goal: current=shown camera



Visint.org video dataset

Sensors: tracks

Entities: actors, objects, rel. distance

Attribute: s, x, y, velocity

Actions: 12 verbs

Goal: seen=learned verb



Contributions

- ❖ Structure Learning
- ❖ Learning from Uninterpreted Sensors & Actions
- ❖ Detecting and Adapting to Unpredicted Changes
- ❖ Detecting and Reasoning with Interference

Prediction Rules and Models

- A learned Predictive Model is represented by Prediction Rules
 - Rule = Conditions → Action⁺ → Predictions
 - Conditions describe the state of observed entities and attributes prior to executing an action
 - Condition = (Entity, Attribute, Comparison Operator, Value)
 - e.g. c1 = (red, size, >, 10)
 - Predictions describe the forecasted state of entities and attributes as a result of executing an action (when conditions are satisfied)
 - Prediction = (Entity, Attribute, Expected Change, Value)
 - e.g. p1 = (blue, size, ↑, 5)
 - Rule example: [c1 → left → p1], or $P(p1|c1,\text{left})=0.62$
- Rules can be sequenced
 - i.e. $[o_0, a_1, p_1, \dots, a_n, p_n]$
 - Represent concepts such as plan, experiment, example and advice
 - Prediction sequences can capture temporal relations in observation sequences

Model Representation

- Prediction Rules
 - $[C,A] \rightarrow P$, e.g., [Raining, Step-out] → Get-Wet
 - In the form of probability $P(R|C,A)$, where R is the prediction
 - Different from $C \rightarrow A$ (production rules or action policy)
- Prediction Sequences
 - $PS = (S_0, A_0, P_1, A_1, P_2, A_2, \dots, A_n, P_n)$
- The multifunction of prediction sequences
 - Plan: if P_n is a goal
 - Exploration: if P_n guarantees a surprise
 - Experiment: if P_n may generate a surprise
 - Example: if the sequence will generate some surprises
 - Advice: if the sequence must be converted into prediction rules
- Layered networks for representing prediction rules from action-sensory data

Components of Prediction Rule

Goals
 $G = \{(s_1 e_1 b_1 = v_1) \wedge (s_1 e_1 b_2 = *) \wedge (s_2 e_1 b_x \neq \emptyset) \wedge \dots \wedge (s_i e_j b_k = v_l)\}$
 True when $O_t \models G$
 e.g. desired observation shown

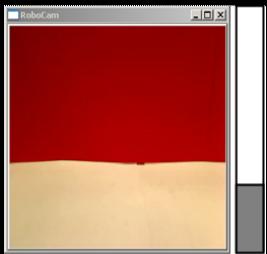
Observations
 $O_t = \{(s_1 e_1 b_1 = v_1) \wedge \dots \wedge (s_n e_i b_j = v_k)\}$

Sensors
 $\{s_1, s_2, \dots, s_n\}$
 e.g. {proximity, camera, goal}

Models

Actions
 $\{a_1, a_2, \dots, a_n\}$
 e.g. {F,B,L,R}

Comparison Operators
 $\{\%, \sim, \uparrow, \downarrow, \Theta_1, \Theta_2, \dots, \Theta_n\}$
 e.g. {%, ~, <, <=, !=, >, >=, >, <, >, <=}



Entities
 $\{e_1, \dots, e_f\}$
 e.g. {blob_white, blob_red etc.}

Attributes
 $\{b_1, \dots, b_g\}$
 e.g. {size, center-x, center-y}

Values

continuous [p, q] or discrete $\{v_1, v_2, \dots, v_h\}$

Predictions

Prediction for t is a forecasted observation in CNF prior to t
 $P_t = \{(s_1 e_1 b_1 \Theta v_1) \wedge \dots \wedge \neg(s_2 e_1 b_1 \Theta v_k)\}$
 True when $O_t \models P_t$

Surprises

A discrepancy between a prediction and its observation at t
 True when $O_t \not\models P_t$

Background and Related Work

- Shen-Simon89, Shen92, Shen94(book), Nadeesha-Shen2011, Nadeesha2012, Collins-Shen 2017, 2018
- Three different attitudes towards unpredicted changes [Saffiotti97]
 - Get rid of it: with precisely engineered robots, e.g. industrial robots
 - Tolerate it: using redundancy with carefully designed contingency strategies
 - Reason about it: with techniques for the representation and manipulation of uncertain information
- Tolerate it
 - Detect changes using models of sensors, actions and environments
 - Sensors: Model is given [Visinsky91] [Roumeliotis98] or learned from examples [Murphy96] [Stronger06]
 - Actions: [Ferrell94] [Kececi08]
 - Goals: [Wolpert98] [Doya02]
 - The set of **contingency strategies** may **not cover all unpredicted changes**
- Reason about it = data-driven adaptation
 - Sensors: Sensor fusion [Pierce97] [Soika02]
 - Actions: [Lipson04] [Pierce97] [Bongard06] [Mugan10]
 - Goals: Update model & re-plan [Yildirum99] [Busch07] [Stone08]
 - None focus on **simultaneous unpredicted changes to sensors, actions, goals & environment** in an **unsupervised** manner

Properties of Surprises

- Types of surprise
 - Unexpected failures
 - Unexpected successes
 - Null prediction surprise
 - When there is no a priori model
- Differentiate “new information” from noise
 - Focus attention on the relevant features
 - Seek differences that are statistically significant

Analyzing and Using Surprises

- Extract the critical information from a surprise
 - Create new model, if it is a null-prediction surprise
 - (\sim non-exists \wedge exists) = action \Rightarrow (appeared \wedge disappeared)
 - Improve the learned model
 - Compare the surprised situation with previous situations where the prediction was successful (i.e., “find the differences”)
 - Identify salient features that discriminate the situation
 - Which part of the memory to compare? (the effect of “first kiss”)
- Finding the true cause of a surprise could be difficult
 - E.g., Raining days? Mondays?
 - Too many differences
 - Too few differences (no differences)

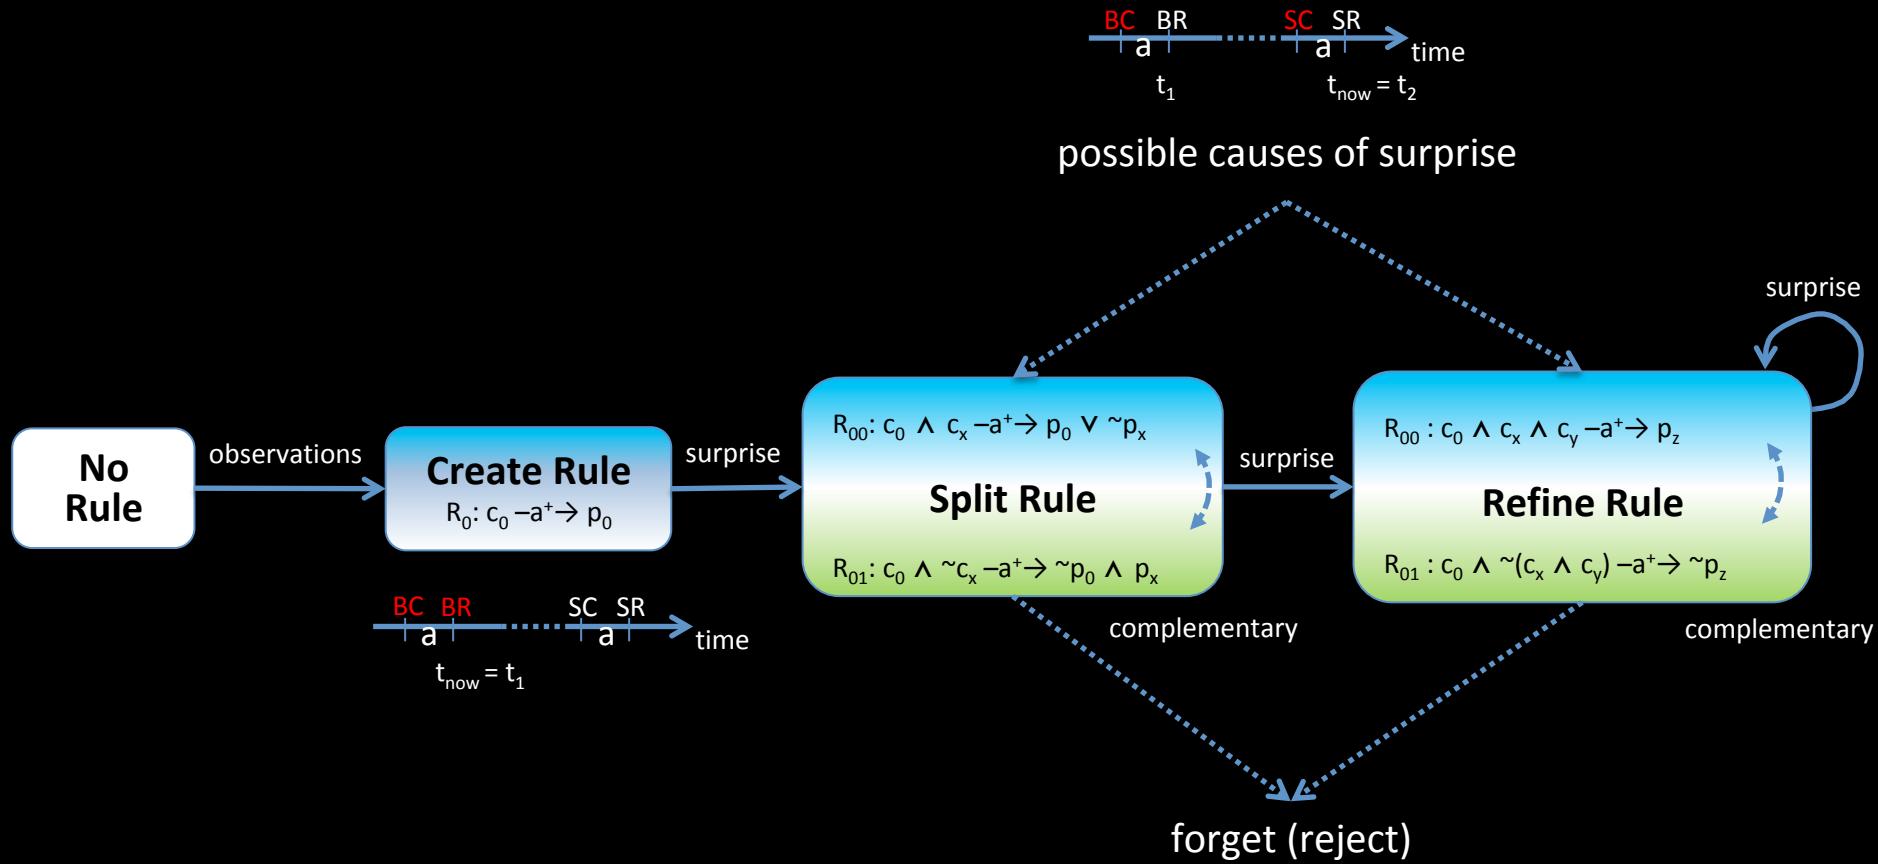
Desirable Features

- Autonomously extracts the states and corresponding state machines from the training dataset object tracks
- Number of states, observed features of states & transitions between states need not predetermined
- Top-down approach to meet low-level measurements and evidences
- Learn the structure of states for description and explanation of revisions
- “Surprise” connects to “anomaly” and “unexpected interference”

Previous Results

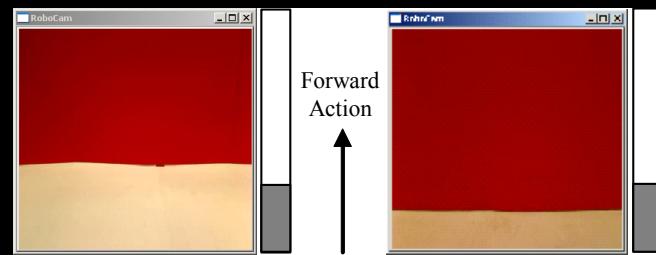
- Discrete or symbolic environments
 - Developmental learning to use novel tools
 - Scientific discovery of hidden features (genes)
 - Game playing
 - Learning from large knowledge bases
- Robotics environments
 - Sensor recovery in office environment
 - Recognize/predict visual acDons from video
 - Detect abnormality in robot/UAV swarms

Life Cycle of a Prediction Rule



Creation of New Prediction Rules

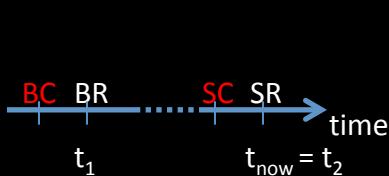
- When to create rules?
 - No predictions exist
 - No change indicated in all predicted rules (general rules)
 - All predicted rules have been rejected
- How to create rules?
 - Base-condition ‘BC’ = observation made at time t_1-1 before action was executed
 - Base-result ‘BR’ = observation made at time t_1 after action was executed
 - Compare BC and BR with operators $\{\%, \sim, \uparrow, \downarrow\}$ & create a rule for each difference



- e.g.
 - RuleR1: (White, All, %, 0) → FORWARD → (White, S, <, Value1) // size of white decreased
 - RuleR2: (White, All, %, 0) → FORWARD → (White, Y, >, Value2) // y-location of white increased
 - RuleR3: (Red, All, %, 0) → FORWARD → (Red, S, >, Value3) // size of red increased
 - RuleR4: (Red, All, %, 0) → FORWARD → (Red, Y, >, Value4) // y-location of red increased

Surprise Detection & Analysis

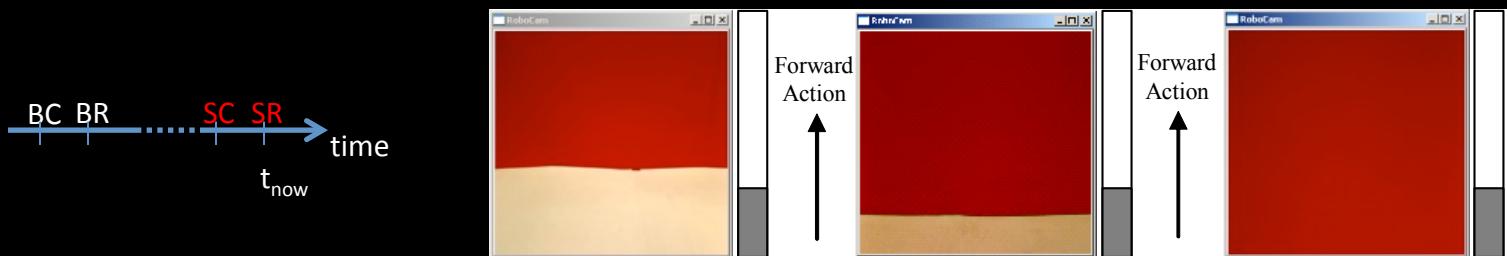
- Detect surprises by comparing the prediction against the observation
- Surprise analysis identifies possible cause(s) of the surprise
 - Surprised-condition ‘SC’ = observation made at time t_2-1 before action was executed
 - Surprised-result ‘SR’ = observation made at time t_2 after action was executed
 - Compare the current SC to a previous BC where the rule succeeded, with the given operators (typical set $\{\%, \sim, >, <, =, !=\}$)
 - Return a set of possible causes $[c_1, \dots, c_z]$
 - Each cause is an expression of an entity or attribute that was true in BC but false in SC
 - SBL assumes that a good cause can be found, no hidden states [Shen94]



- e.g.
 - Causes of surprise for RuleR1 = [(Red, S, $<$, Value5), (Red, Y, $<$, Value6), (White, S, $>$, Value7), (White, Y, $<$, Value8)]

Splitting Prediction Rules

- When to split rules?
 - A surprise occurs in a single rule for the first time
 - Single rule $R_0: c_0 \rightarrow p_0$
- How to split rules?
 - Get possible causes $C = [c_1, \dots, c_z]$ from surprise analysis
 - For each cause c_x identify the expected change p_x from SC and SR
 - Create a complementary pair of rules
 - Specialization $R_{00}: c_0 \wedge c_x \rightarrow p_0 \vee \neg p_x$
 - Generalization $R_{01}: c_0 \wedge \neg c_x \rightarrow \neg p_0 \wedge p_x$



- e.g. surprise analysis for RuleR1 returns 4 possible causes and 3 new consequences resulting in 12 pairs of rules:
 - RuleR1.1.1: (White,All,%,0) \wedge (Red,S,<,Value5) \rightarrow FORWARD \rightarrow (White,S,<,Value1) $\vee \neg$ (Red,S,>,Value9)
 - RuleR1.1.2: (White,All,%,0) $\wedge \neg$ (Red,S,<,Value5) \rightarrow FORWARD $\rightarrow \neg$ (White,S,<,Value1) \wedge (Red,S,>,Value9) ...

Refinement of Prediction Rules

- When to refine rules?
 - Any further surprises occur in a complementary rule
 - $R_{00}: c_0 \wedge c_x \neg a^+ \rightarrow p_0 \vee \neg p_x$
 - $R_{01}: c_0 \wedge \neg c_x \neg a^+ \rightarrow \neg p_0 \wedge p_x$
- How to refine rules?
 - Get possible causes $C = [c_1, \dots, c_z]$ from surprise analysis
 - When R_{00} is surprised, for each cause c_y that is not specified in c_0 or c_x
 - $R_{00}: c_0 \wedge c_x \wedge c_y \neg a^+ \rightarrow p_z$
 - $R_{01}: c_0 \wedge \neg(c_x \wedge c_y) \neg a^+ \rightarrow \neg p_z$
 - If no new cause, add the entire observation ($e_1, \dots, e_b.v_n$)
 - Similarly, when R_{01} is surprised, add c_y to its conditions



- e.g. by RuleR2.1.1 a forward action predicts that the y-location of white will increase, but it disappears, leaving 6 causes:
 - RuleR2.1.1.1: $(White, All, \%, 0) \wedge (Red, S, <, Value5) \wedge (Red, X, >, Value12) \rightarrow FORWARD \rightarrow (White, Y, >, Value2) \vee \neg(White, Y, >, Value9)$
 - RuleR2.1.2.6.2: $(White, All, \%, 0) \wedge \neg((Red, S, <, Value5) \wedge (White, S, >, Value17)) \rightarrow FORWARD \rightarrow \neg(White, Y, >, Value2) \vee (Red, S, >, Value9) \dots$

Dealing with Unpredicted Changes

- Goal change
 - Plan with the learned model to observable goals
 - Maintain a list of goal observations for hidden goals
- Dynamic configurations in the environment
 - Forget obsolete rules through rule rejection when contradictions or the success probability of a rule drops below a cutoff
- Action & sensor changes
 - Maintain a table with each row recording the relevance of a sensor, entity, attribute and action
 - Identify irrelevance by monitoring entities and attributes in the learned prediction rules
 - Force relevance when none of the active entities or attributes change

Some Evaluation Environments

Simulated hunter-goal

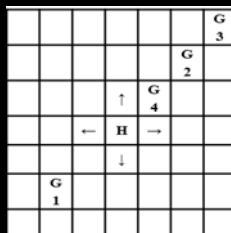
Sensor: location

Entities: H, G

Attributes: x, y, u

Actions: N, S, E, W

Goal: $H(x,y)=G(x,y) \mid\mid H(u)=G(u)$



Simulated hunter-prey

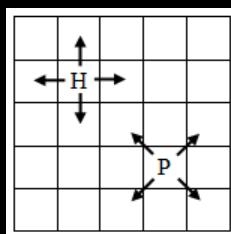
Sensor: location

Entities: H, P

Attributes: x, y, u

Actions: N, S, E, W

Goal: $H(x,y)=P(x,y) \mid\mid H(u)=P(u)$



Real office room

Sensors: camera

Entities: SURF objects

Attribute: s, x, y

Actions: L, R

Goal: current=shown camera



Visint.org video dataset

Sensors: tracks

Entities: actors, objects, rel. distance

Attribute: s, x, y, velocity

Actions: 12 verbs

Goal: seen=learned verb

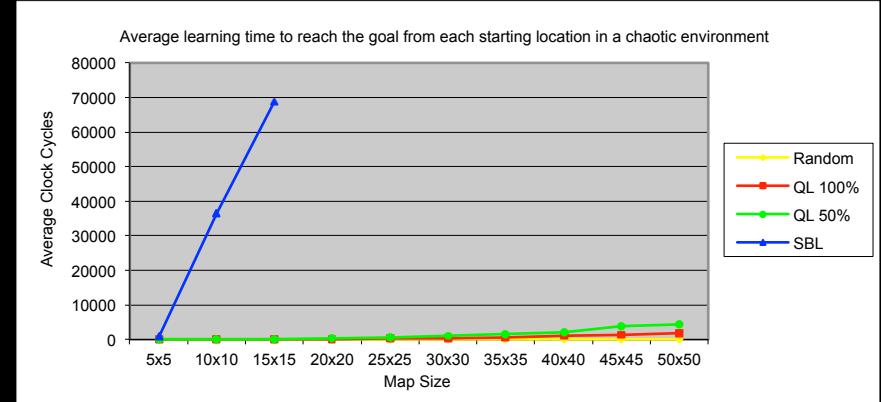
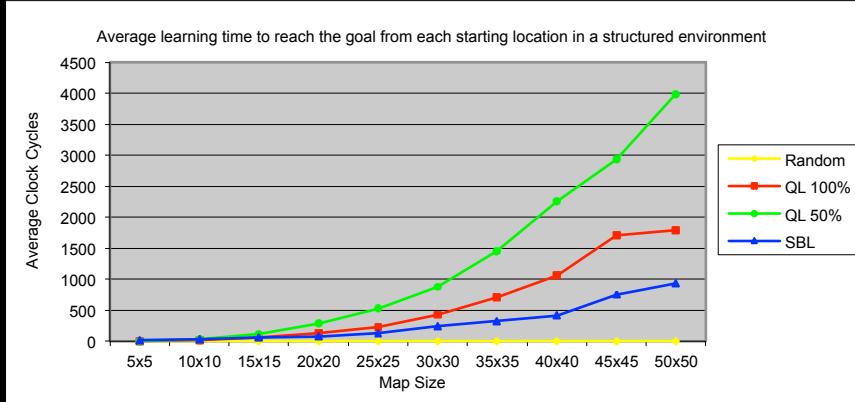
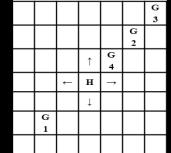


Contributions

- ❖ Structure Learning
- ❖ Learning from Uninterpreted Sensors & Actions
- ❖ Detecting and Adapting to Unpredicted Changes
- ❖ Detecting and Reasoning with Interference

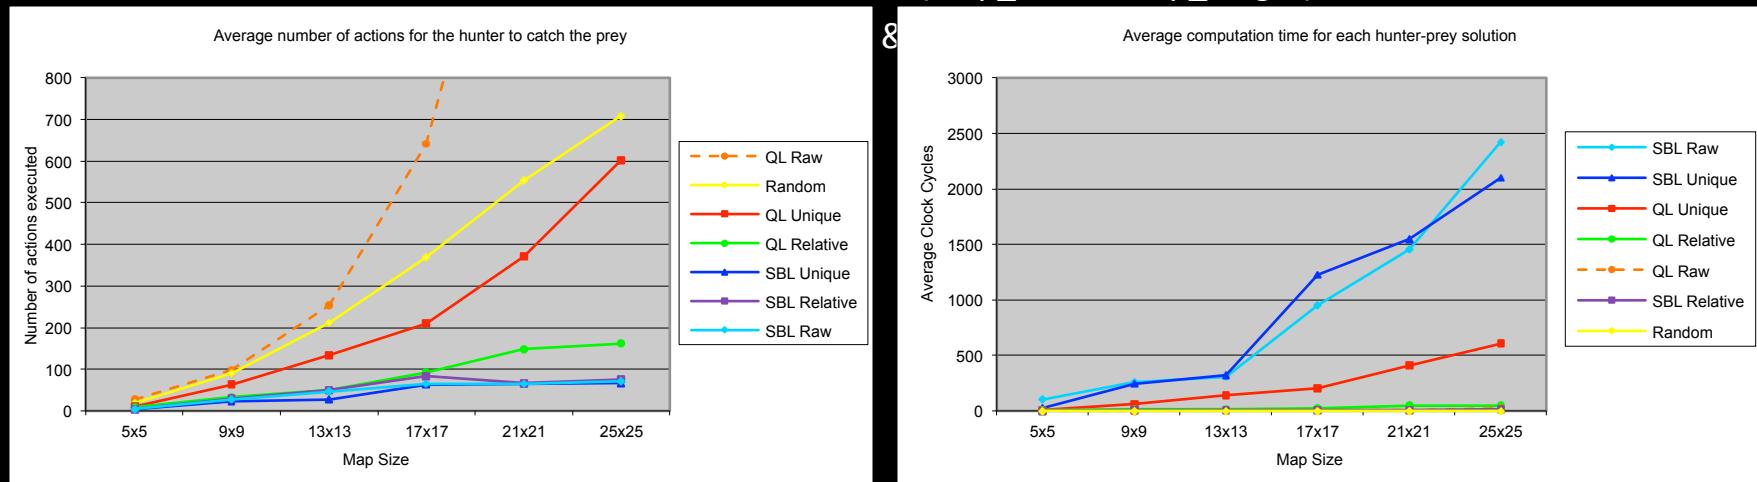
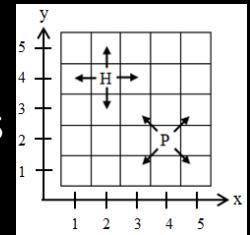
Structure Learning Results

- Structured environment
 - Properties monitored by sensors demonstrate patterns associated with each action
 - Scales by exploiting structure
 - Number of prediction rules proportional to structure, not size of environment
 - In chaotic environments SBL memorizes the data



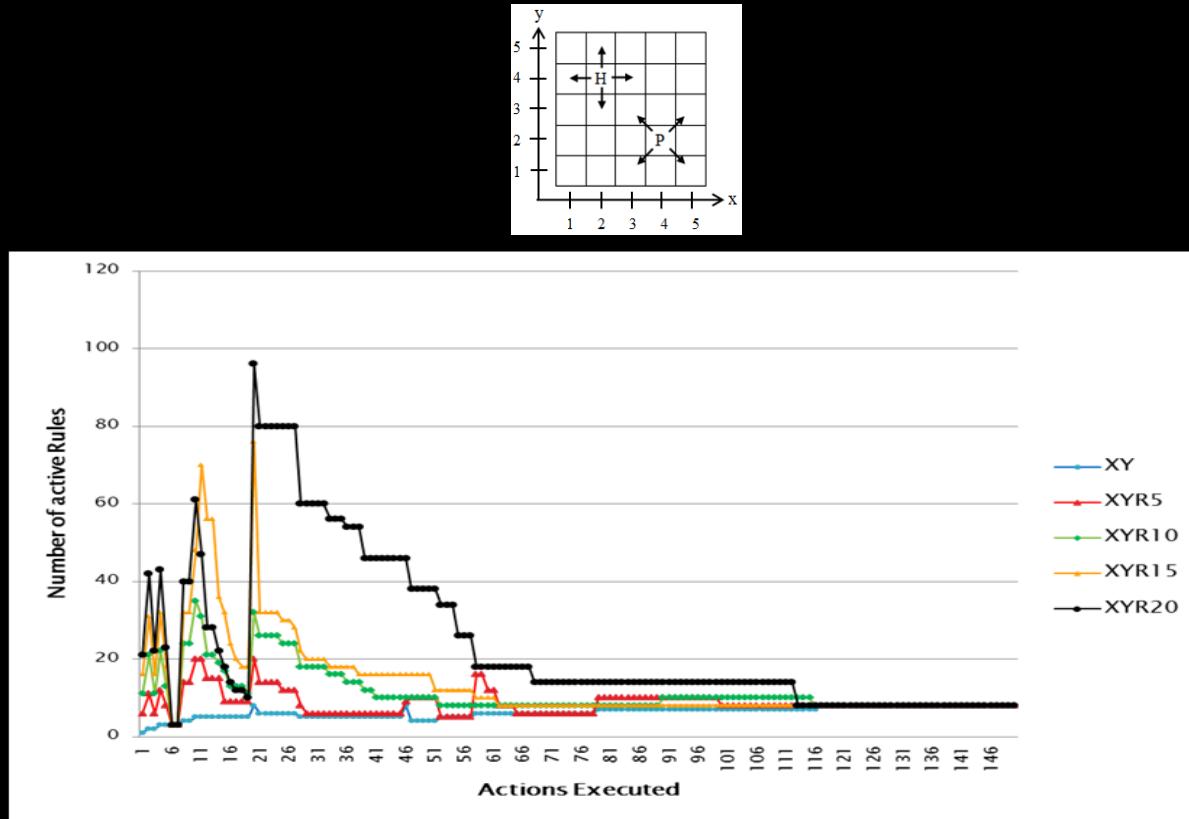
Learning from Uninterpreted Sensors & Actions

- Discretizes continuous data
 - Identifies data ranges and trends with comparison operators
 - e.g. Learned prediction rules forecast wrap around at boundaries
- Combines multiple uninterpreted sensors
 - Dimension reduction without hand-crafted approximation functions
 - Unique hcF: Locations of H & P merged into $(\text{map_width} * \text{map_height})^2$ states
 - Relative hcF: Relative location of P from H $(\text{map_width} * \text{map_height})$ states



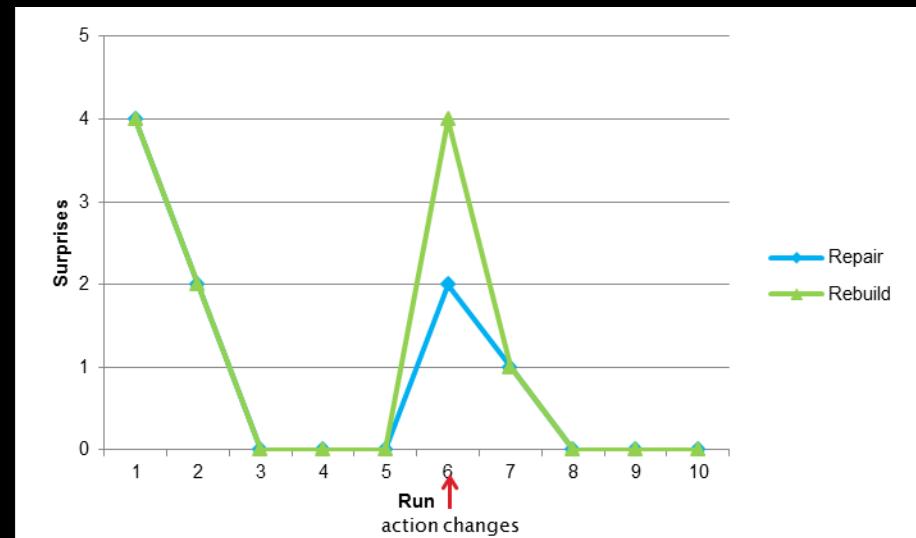
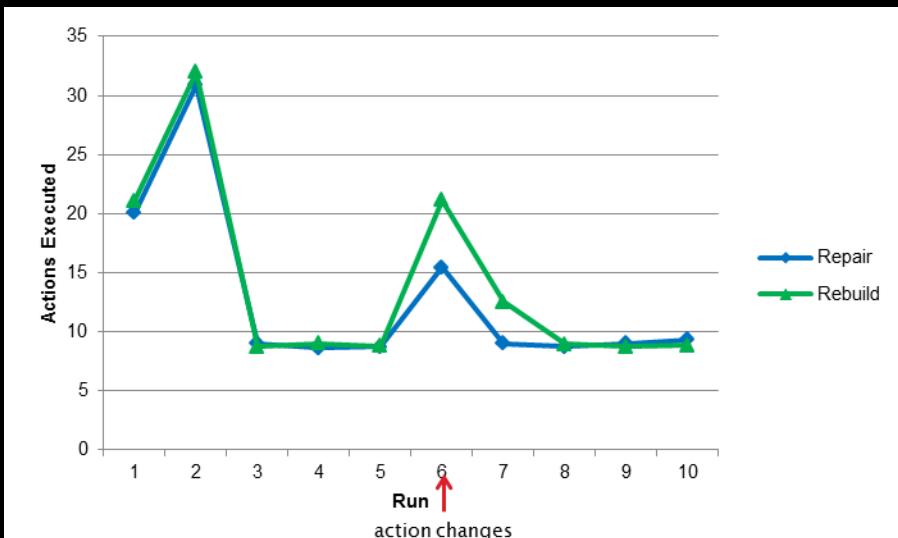
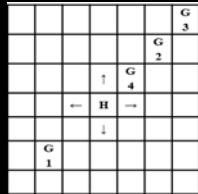
Scalability in Sensors

- Scales by exploiting structure, and sensor & action relevance
 - A different number of random number sensors was added to each test
 - The same sequence of 150 actions was executed in each test



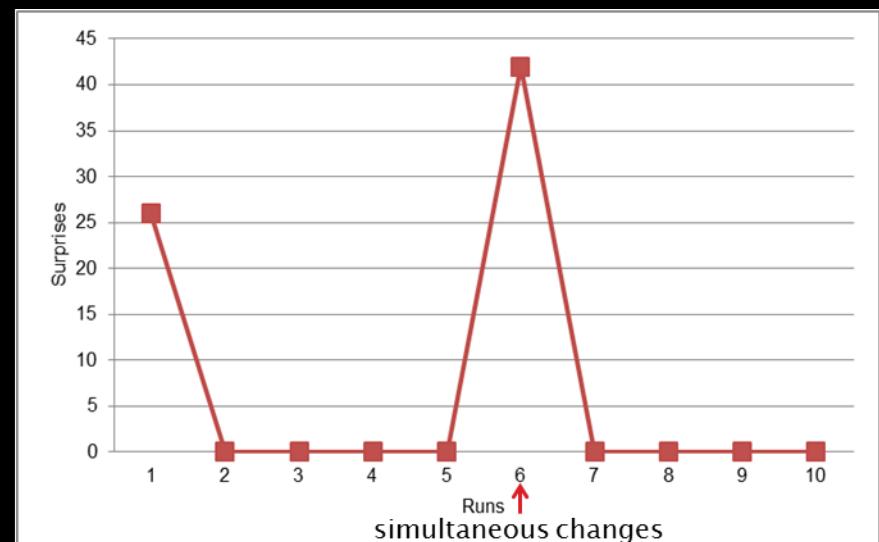
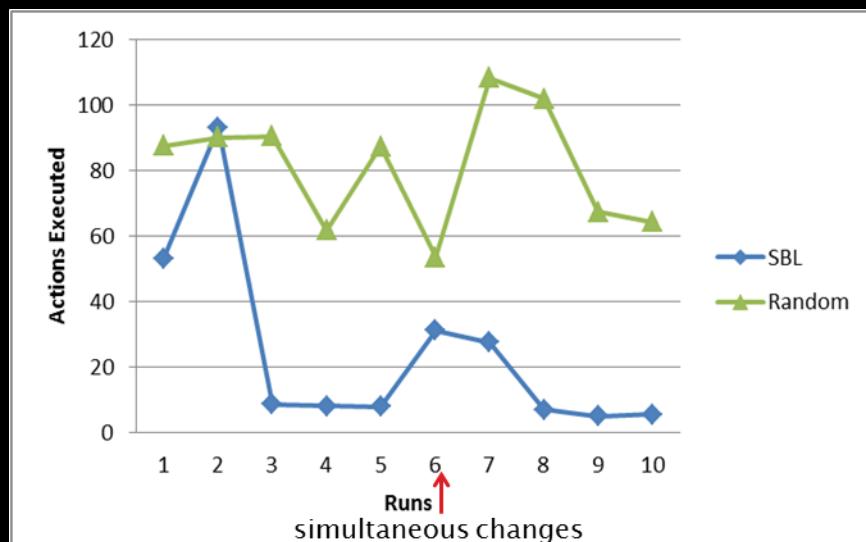
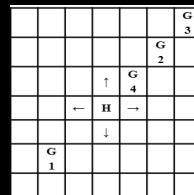
Repairing the Learned Model vs. Rebuilding from Scratch

- SBL detects unpredicted changes through surprises
 - Naturally repairs, does not rebuild from scratch
 - Forced to rebuild from scratch in this experiment
- Repairing preserves past experiences and prevents resource wastage
 - Run completed when goal reached, alternated between goals G_1 & G_2
 - In run 6, swapped N & E



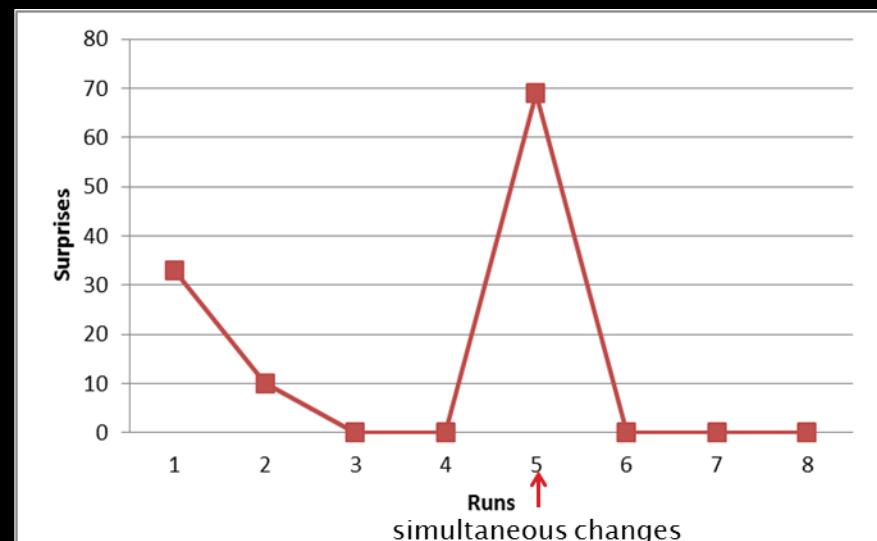
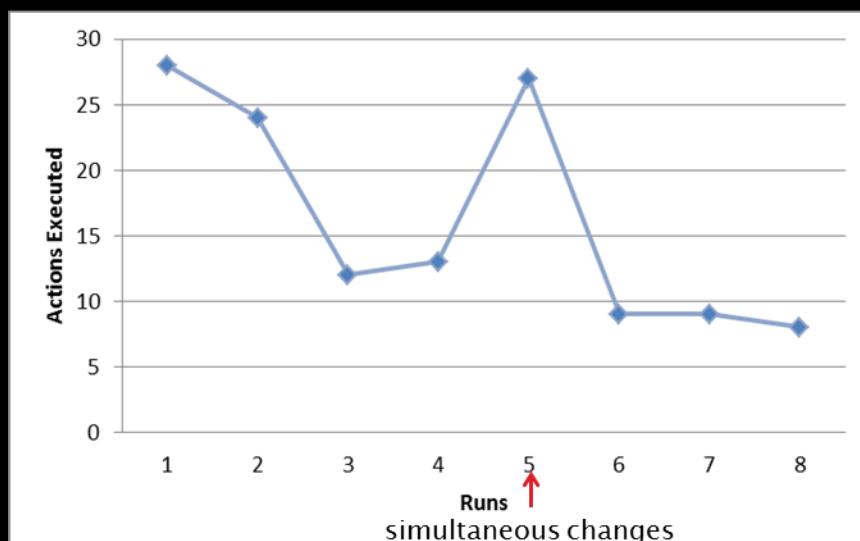
Detect and Adapt to Simultaneous Unpredicted Changes in Actions, Sensors & Goals

- SBL is open-minded for all possible causes, and eventually adapts
 - Run completed when goal reached, alternated between goals
 - In run 6, swapped N & W, removed u, moved G1 & G2 to G3 & G4



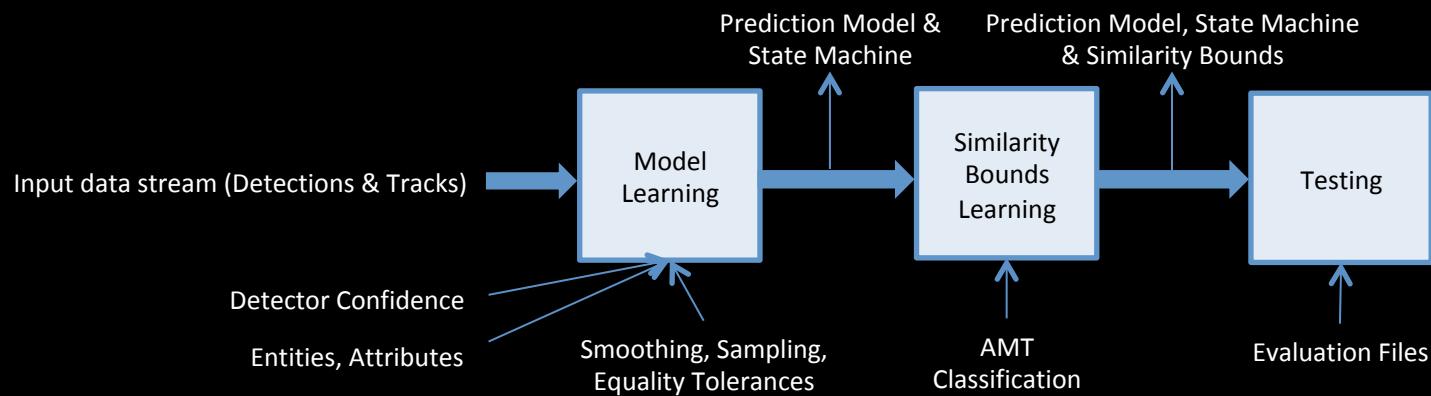
Detect and Adapt to Simultaneous Unpredicted Changes in Actions, Sensors, Goals & Environment

- SBL is feasible for real-world robots in unpredictable environments
 - Run completed when goal reached, alternated between goals
 - In run 5, swapped L & R, rotated 90° & translated camera, changed one goal book, moved the other goal book to a closer wall



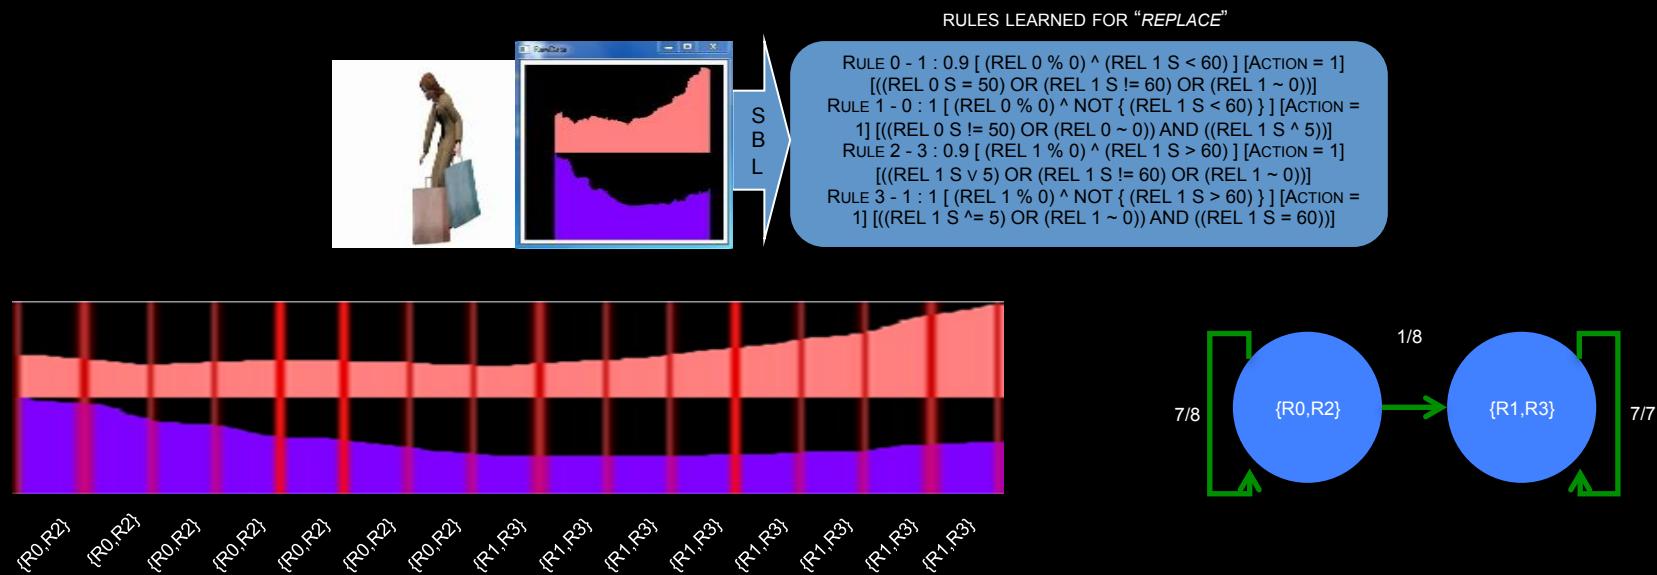
Detecting and Reasoning with Interference

- Interference is short term disturbance to sensors or actions
 - Only “Noise” & “Missing data” (Gaps) were considered
 - Detected as surprises
 - Do not adapt the learned model as the cause is unidentifiable
- SBL applied to action recognition & gap filling in video data streams
 - approach, arrive, bounce, carry, catch, chase, collide, drop, exchange, exit, flee, go, haul, leave, lift, move, pass, pick up, push, put down, raise, replace, throw, walk
 - Entities: Actors, Objects, Relative Distance Composite
 - Attributes: x-location, y-location, velocity, size



Action Recognition – Learning I

- Model Learning
 - Sample the data stream at fixed intervals, e.g. 10 frames
 - Learn a prediction model by observing data before and after each segment
 - Group prediction rules that fire together into a state
 - Extract the sequence of states (prediction sequence) for each example
 - Store state transition probabilities for positive examples in a Markov chain



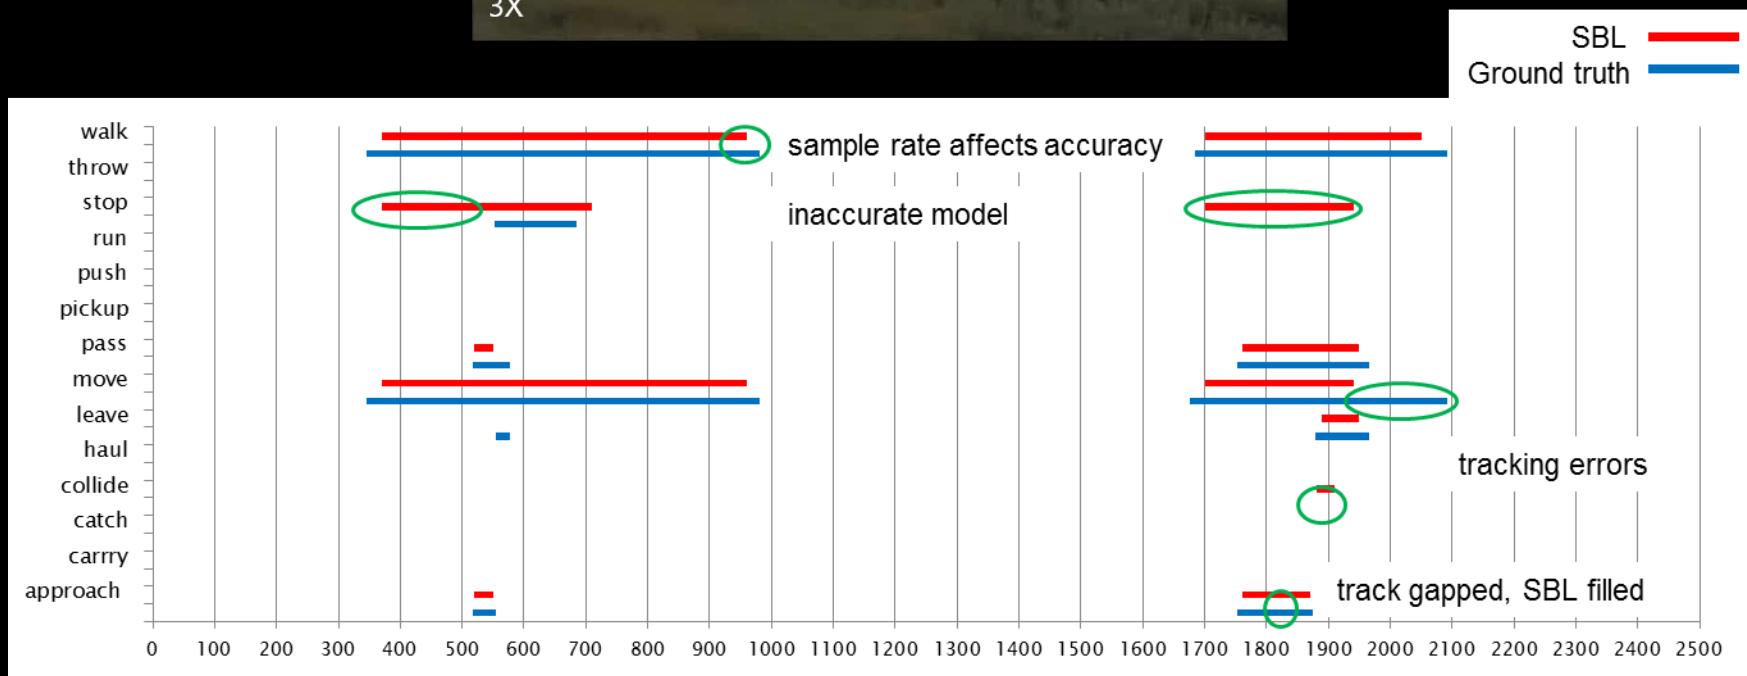
Action Recognition – Learning II

- Similarity Bounds Learning
 - Compute a prediction sequence by applying a model to the data stream
 - Calculate similarity metric to validate the observed state sequence
 - Valid state transitions score 0
 - Fixed negative penalties for invalid start state, end state and state transitions
 - $\text{similarity}(\text{sequence}, \text{MC}) = 0 - \text{noStartPenalty} - \text{noEndPenalty} - n * \text{transitionPenalty}$
 - Allows comparison of actions and videos with variable durations
 - Record lowest similarity value observed for all training examples
 - Highest similarity value = 0
 - Provides noise tolerance by accepting detections between these bounds

Action Recognition - Testing

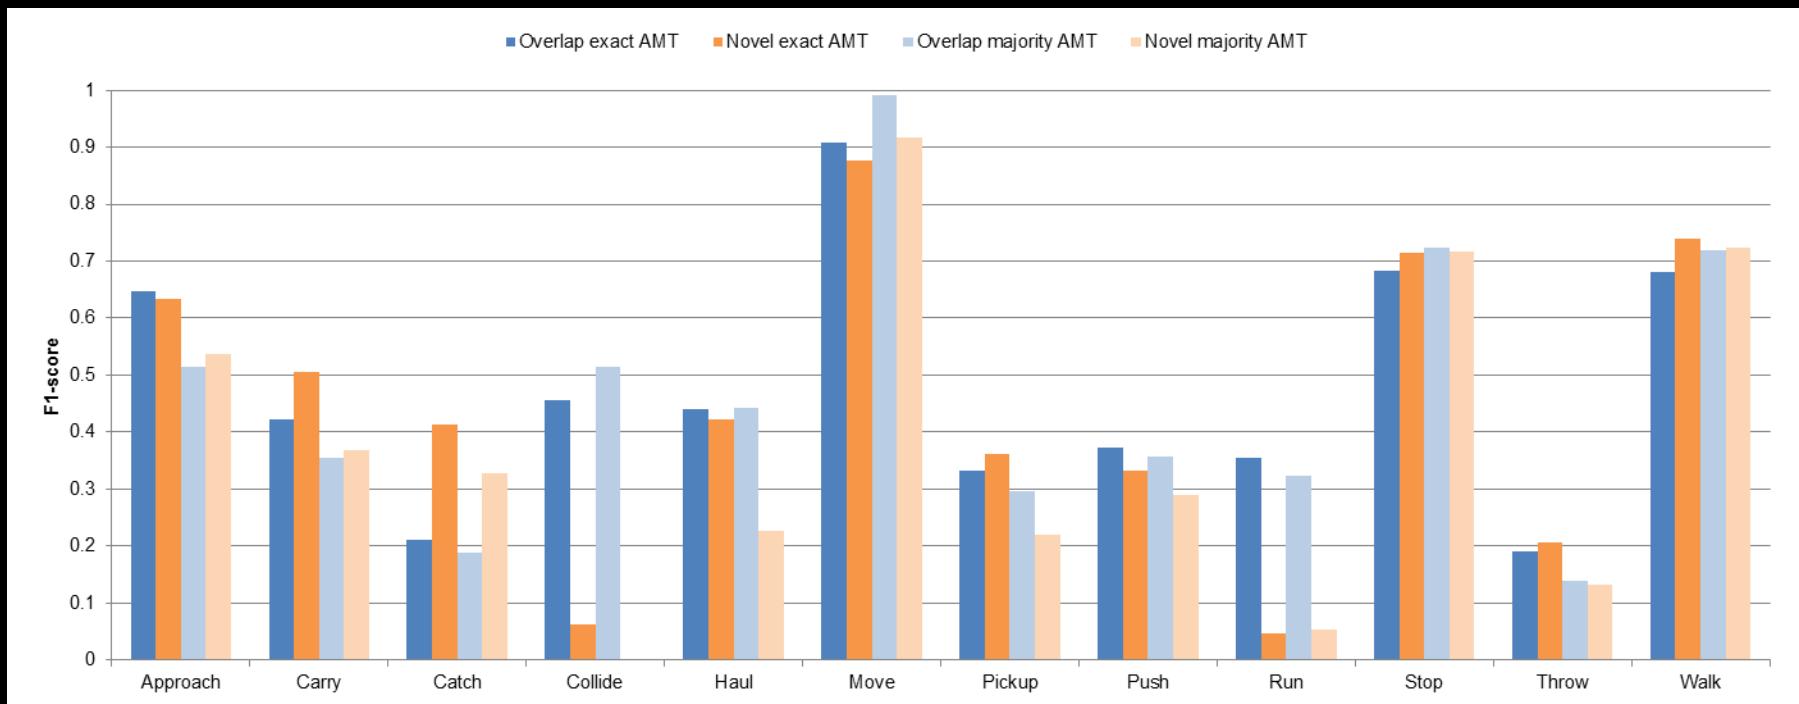
- Variable binding problem
 - Entities in a learned model may match several entities in the testing data
- For an unclassified video, given a model
 - Select a valid binding and extract a prediction sequence
 - Compute similarity metric for the prediction sequence
 - Test every valid binding and keep the highest value
 - If the highest value is within the learned similarity bounds the detection is true for that model
- When an action occurs multiple times in a video
 - Report matching start and end states as separate detections

Recognition Example



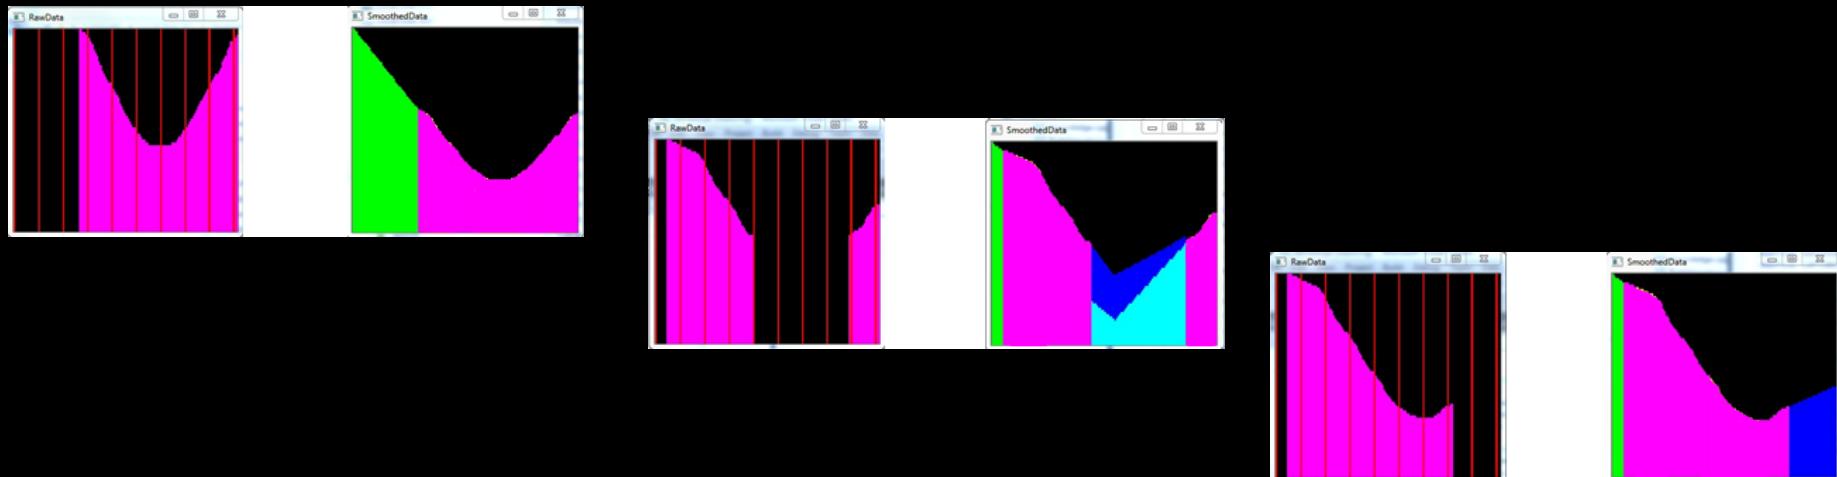
Recognition Results

- Visint.org Year 1 Dataset, USC tracks, 1200 videos
- $F1 = (2 * \text{precision} * \text{recall}) / (\text{precision} + \text{recall})$
- Outperformed hand-coded SAM by more than 10%
- Reasoning with detector accuracies below 20%



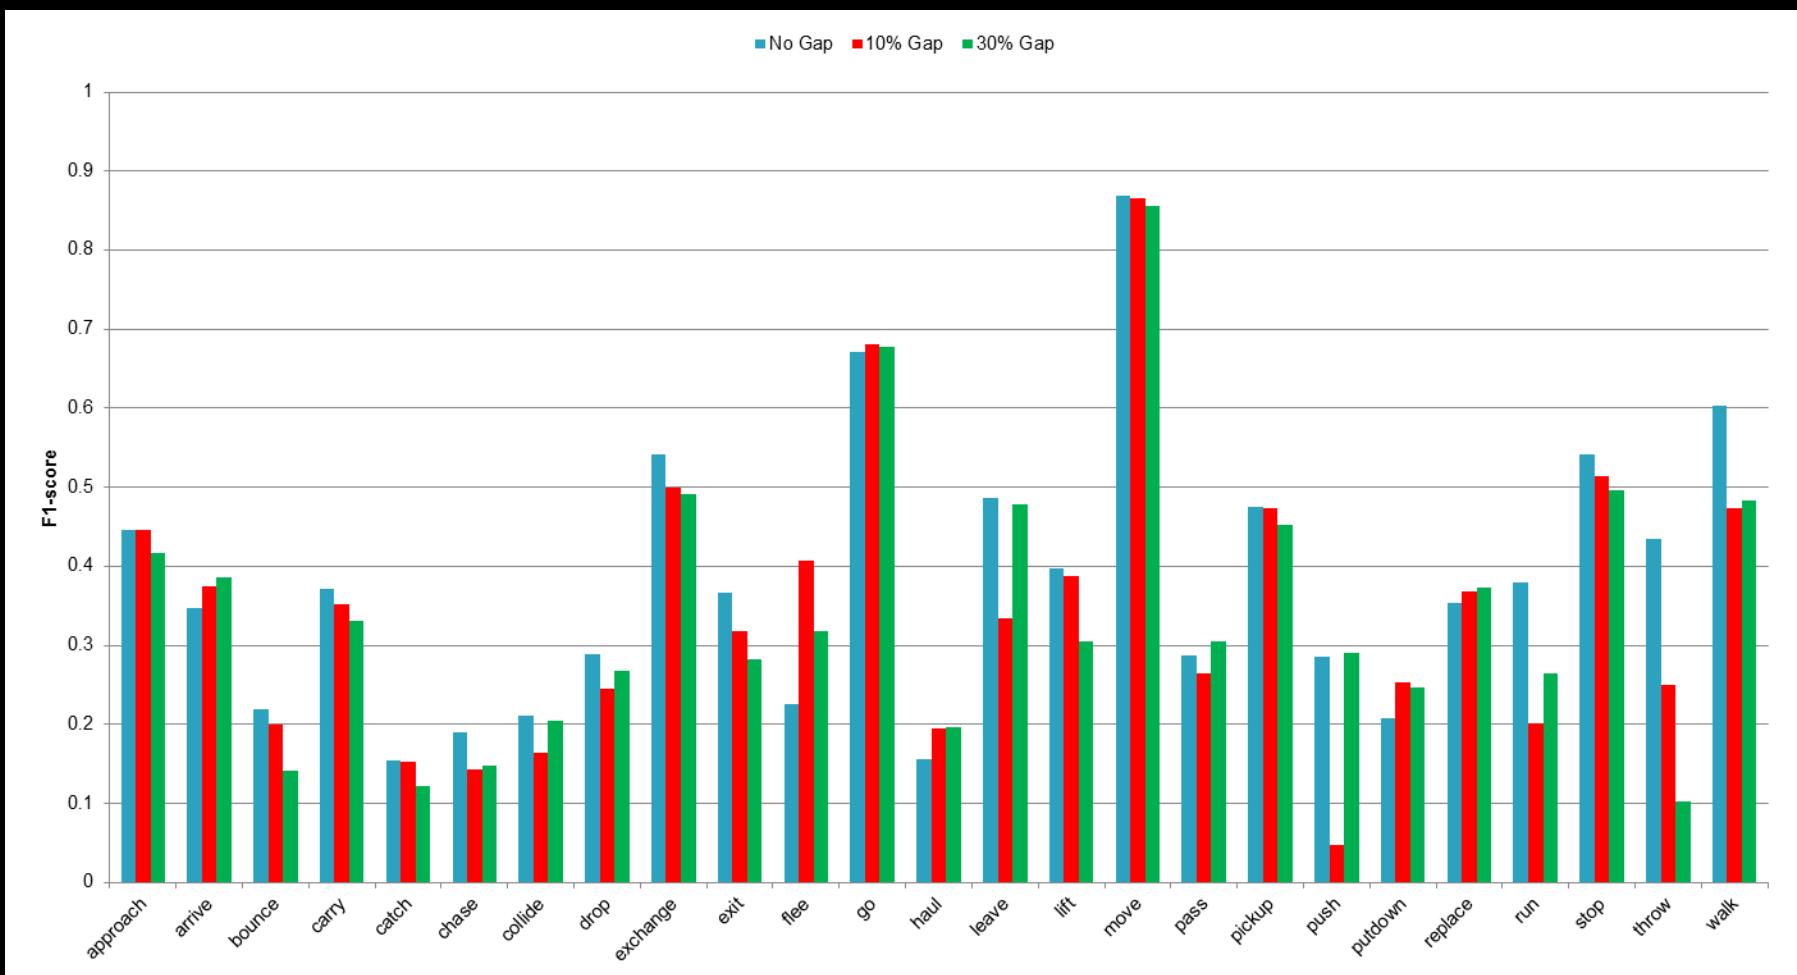
Missing Data - Gap Filling

- Perform BFS with the Markov Chain to fill gaps
 - **Beginning** – Post diction from first known state
 - **Middle** - Interpolation between known states
 - **End** - Prediction from the last known state
- Does not recover each trend, recovers state
- Slight penalties applied to decrease the likelihood action detection



Gap Filling Results

- Synthetic gaps introduced into 700 annotated videos



Summary of SBL Capabilities (so far)

Structure Learning

Surprise-Based Learning for Developmental Robotics, LAB-RS 2008

- +Discover the number of states
- +Scalable/not over fitting in structured environments
- +Few training examples produce a good model
- +Learns generic models that can solve specific goals
- +Temporal modeling with predictive capability

Learning from uninterpreted sensors and actions

The Surprise-Based Learning Algorithm, ISI Technical Report 2008

- +Preprocessing not required
- +Can use preprocessed data
- +Discretizes continuous data
- +Identifies useful sensors and ignores irrelevant sensors

Detecting and adapting to unpredicted changes

Autonomous Adaptation to Simultaneous Unexpected Changes in Modular Robots, IROS 2011 Workshop

Surprise-based developmental learning and experimental results on robots, ICDL 2009

- +Adapt to action changes
- +Adapt to sensor changes
- +Adapt to environmental changes
- +Adapt to goal changes
- +Adapt to simultaneous action, sensor, environment & goal changes
- +Repairs model faster than rebuilding it from scratch

Detecting and reasoning with interference

Autonomous Surveillance Tolerant to Interference, TAROS 2012

- +Detects noise
- +Detects and fills gaps
- +Quantify similarity between learned model and observed data
- +Reason with data of variable durations

Next Improvements

- Explicitly tolerate interference during model learning
 - Probabilistic rules can accommodate noise
 - Keep copies of the original rules during splitting and refinement
 - Planner must consider these probabilities to remove invalid plans
- Faster model convergence
 - Incorporate knowledge that some actions are reversible
 - Hypothesize and test
 - Create abstracted prediction rules during rule creation
 - Identify hidden states with Local Distinguishing Experiments [Shen90]
 - Utilize multiple SBL learners (distributed or nested hierarchy)
- Improve scalability in the number of sensors
 - Limit search space of surprise analysis through goal directed learning
- Reduce planning overheads
 - Generate policies with the prediction model for common goals

Conclusions

- SBL is a general learning method for structural learning from stream observations
- Useful for predictive learning, monitoring, and detecting abnormality
- Easy interface and comprehensible for human operators
- Has a great potential for the other projects

Stochastic Distinguishing Experiments

Thomas Joseph Collins and Wei-Min Shen



School of Engineering
Information Sciences

Overview

Motivation

Problem

Challenges

Related Work

Stochastic Distinguishing Experiments

Surprise-based POMDPs (sPOMDPs)

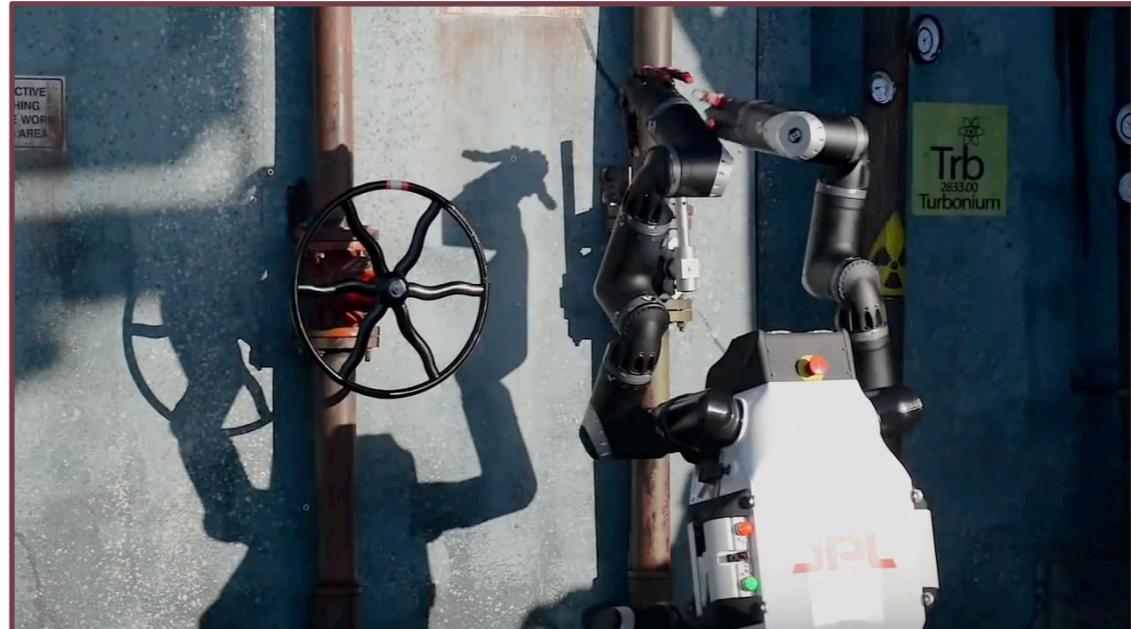
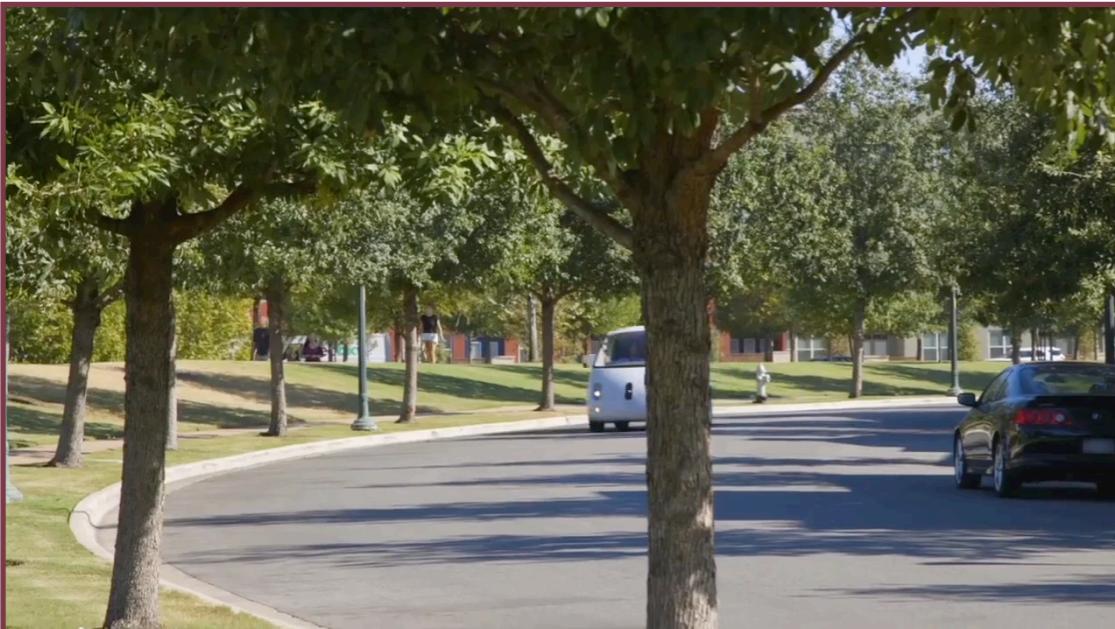
Learning sPOMDPs

Results

Future Work

Motivation

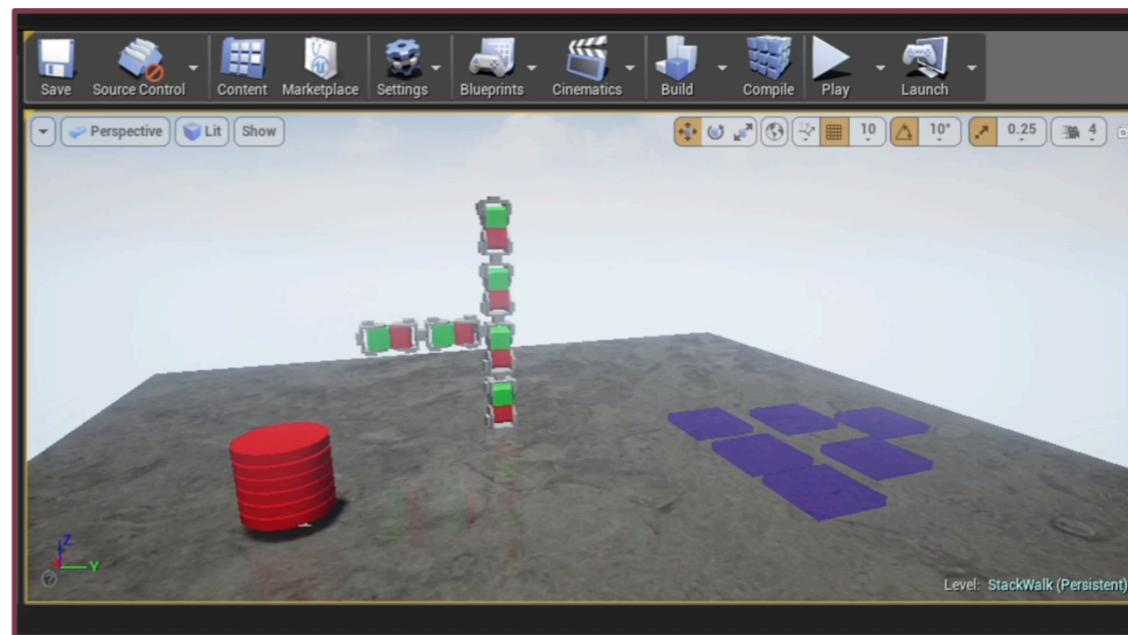
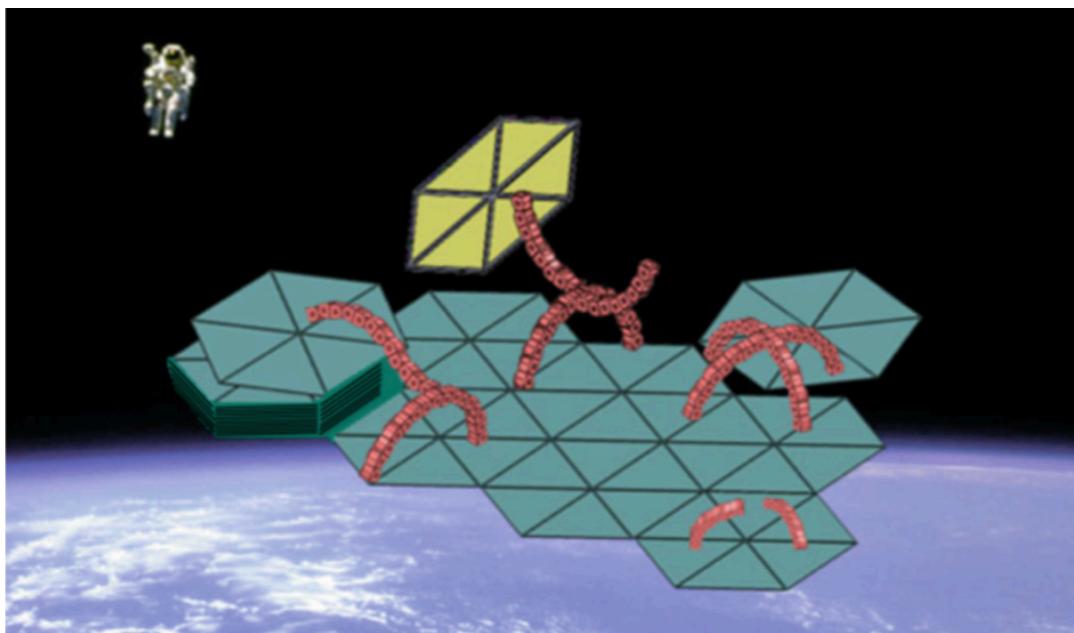
- There is an ever-increasing need for autonomous robotic systems that are able to adapt to and operate in a range of challenging environments.



- These environments vary substantially depending on the task, but virtually all environments of interest exhibit high-levels of **partial-observability** and **stochasticity**.

Motivation

- Human-engineered features and prior knowledge remain some of the most important inputs to algorithms solving fundamental problems in robotics.



- One of the most common forms of prior knowledge is an **underlying state space and/or dynamics representation** specified in terms of human-engineered features.

The Problem

Autonomous Learning from the Environment (ALFE, first formulated by Shen 1993) in which an autonomous agent is tasked with **actively** learning a **task-independent** model of the **state** and **dynamics** of an unknown discrete, partially-observable and stochastic environment from a stream of experience (**actions** and **observations**).

The Problem

Autonomous Learning from the Environment (ALFE), first formulated by Shen 1993) in which an autonomous agent is tasked with **actively** learning a **task-independent** model of the **state** and **dynamics** of an unknown discrete, partially-observable and stochastic environment from a stream of experience (**actions** and **observations**).

- This problem is highly unique in requiring that agent simultaneously decide:
 - **What actions to take** in the absence of external rewards.
 - **How much experience is needed** to build a useful environment model.

The Problem

- Solving ALFE is an important theoretical step toward **lifelong model learning and decision making** in unknown partially-observable and stochastic environments.
 - By **lifelong**, we mean a learning procedure that is meant to be invoked multiple times (or even continuously).

Key Challenges

- Partial-observability
 - Must consider (potentially infinite) history.
- Stochasticity
 - Noise vs. surprise.
- Minimal prior knowledge
 - Learned model must have nonparametric form that grows/shrinks with data.
- Action selection
 - How to choose actions with no reward/goal?
- State space formalization
 - Formalize the learned model to be used in place of human-designed state spaces.

Related Work

- Latent state (hidden variable) approaches:
 - **Representation/deep learning** (Bengio et al. 2013, Goodfellow et al. 2016)
 - **Probabilistic graphical models/temporal models** (Koller and Friedman, 2009)
 - **HMMs, DBNs** (Koller & Friedman 2009)
 - **Reinforcement learning** (Murphy 2000, Pineau et al. 2003, Mnih et al. 2015)
 - **Intrinsically-motivated** (Chentanez et al. 2005)
 - **Curiosity-based** (Pathak et al. 2017)
- Predictive (action/observation) approaches:
 - **Predictive state representations** (Littman & Sutton 2002, Boots et al. 2010)
 - **Stochastic distinguishing experiments** (Collins & Shen, 2017)
- **Stochastic Distinguishing Experiments** and **Surprise-based POMDPs** are a novel hybrid latent-predictive approach in which the results of predictive experiments are used to uniquely identify latent states emitting the same observation.

Stochastic Distinguishing Experiments (SDEs) [BICA 2018]

- **Stochastic Distinguishing Experiments** (Collins & Shen 2017) are sequences of actions – i.e., **experiments** the agent can perform – that statistically disambiguate identical-looking states.
- Let M be a set of latent (environment or model) states. We say that $s = (a_1, \dots, a_k)$ is a **Stochastic Distinguishing Experiment (SDE)** for m^i and m^j if the most likely sequences of observations generated by executing s from m^i and m^j differ ($g_{m^i}^* \neq g_{m^j}^*$).
- The higher the probability of observing $g_{m^i}^*$ from state m^i and $g_{m^j}^*$ from state m^j , the more useful this SDE is in distinguishing m^i and m^j because $g_{m^i}^*$ must have low probability when s is executed from m^j and vice versa.

Surprise-based partially-observable Markov decision processes (sPOMDPs)

- A **Surprise-based POMDP (sPOMDP)** is a POMDP augmented with a set of SDEs S . Each $s \in S$ statistically disambiguates at least two latent model states $m^i \neq m^j \in M$. Each $m \in M$ is represented by the unique most likely sequence of observations resulting from executing the associated SDE from that model state, g_m^* .
- Named for the **surprise-based** algorithm for actively learning sPOMDPs in unknown environments from agent experience.

Stochastic Distinguishing Experiments: Predictive Definition [BICA 2017]

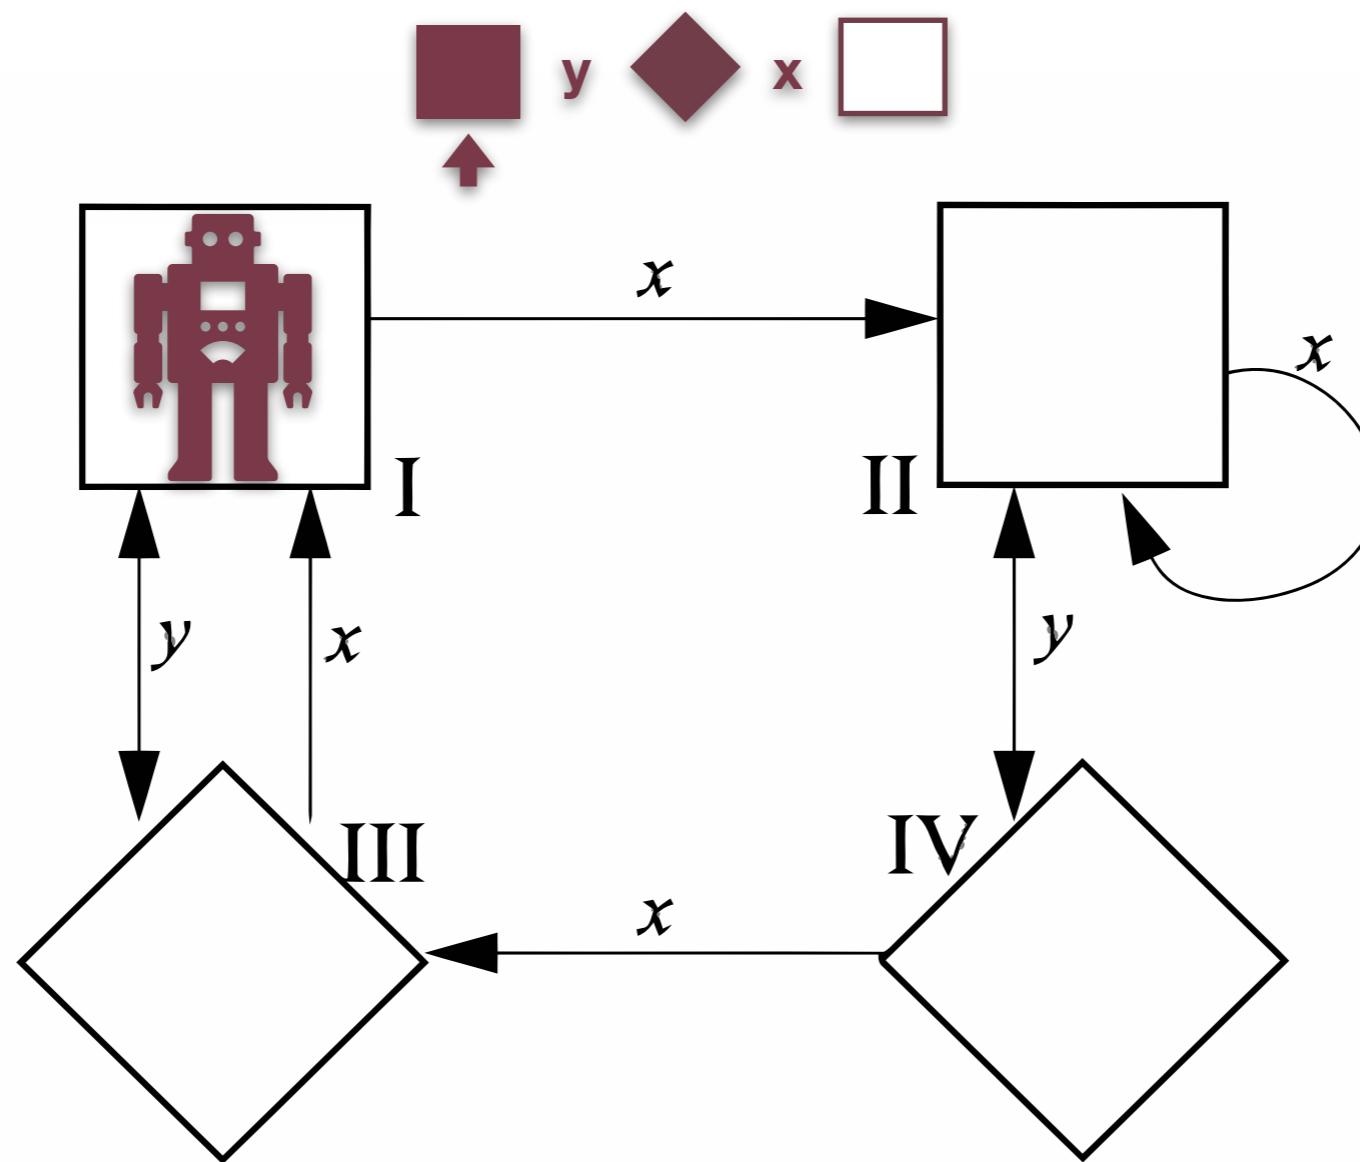
- SDEs were originally defined as **time-invariant conditional probability distributions** over the next observation given a finite, variable-length sequence of ordered actions and expected observations up to the present.
- The **predictive SDE learning algorithm** incrementally extends these experiments — beginning with a single SDE whose experiment is null — while ensuring the **conditions (experiments)** of SDEs are **mutually exclusive and exhaustive at each step**, with each possible agent history matching the experiment of **exactly one SDE**.

$$P(O_{t+1} | o_{t-k+1}, a_{t-k+1}, \dots, o_t, a_t)$$

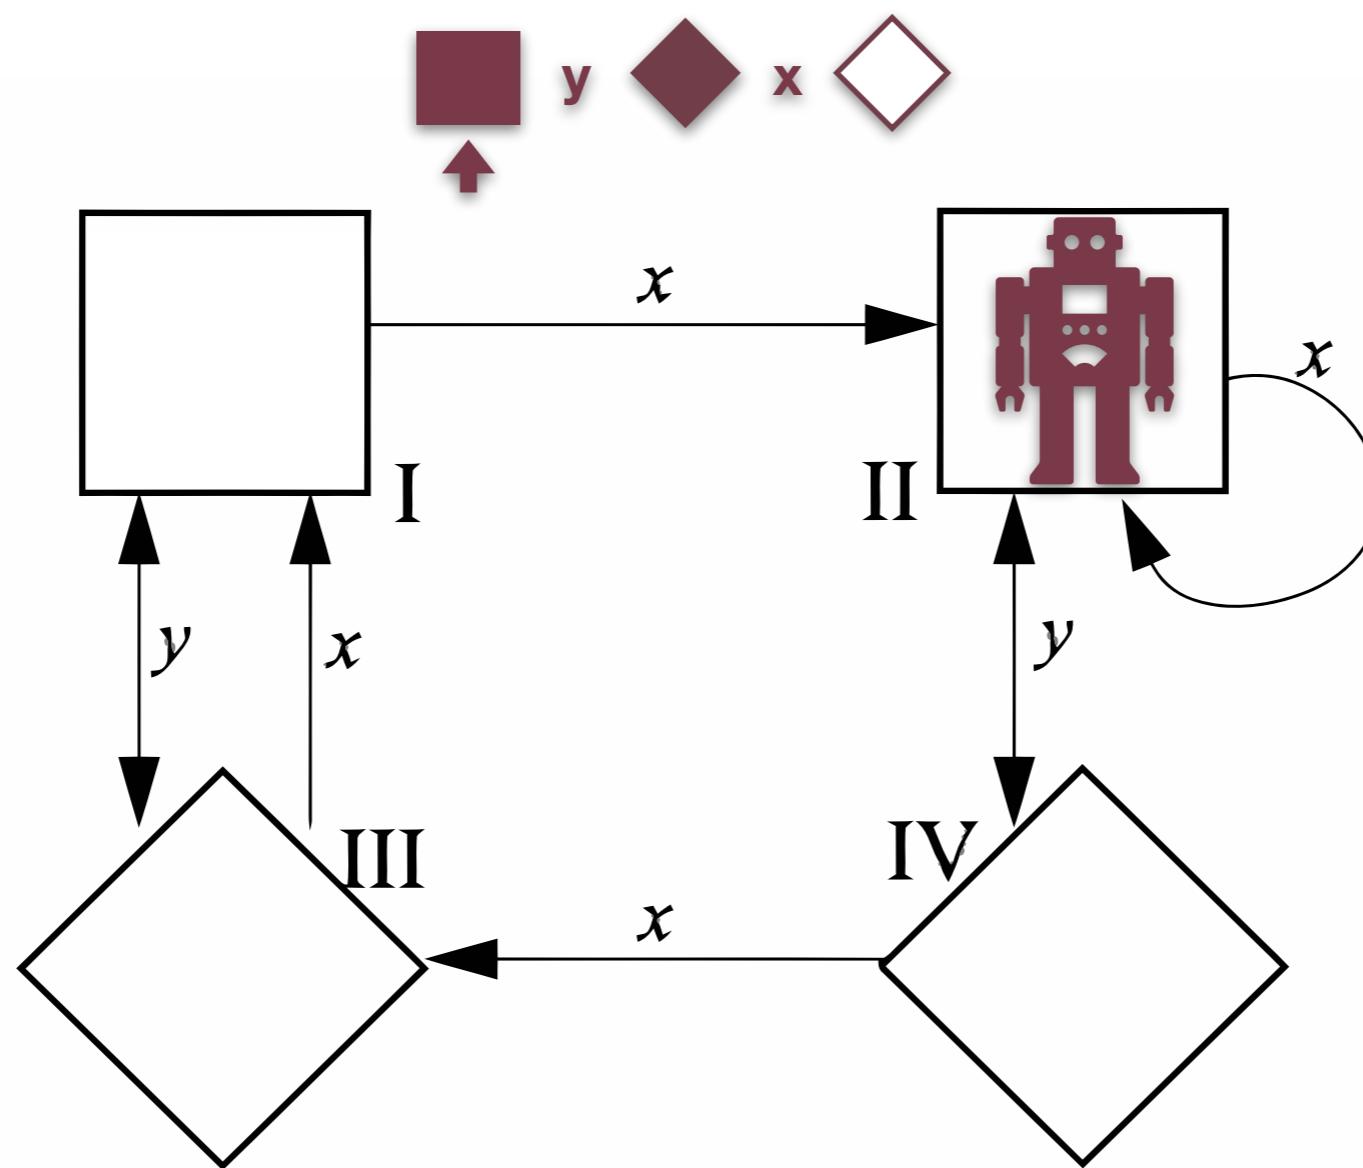
Experiment

Random variable representing the next observation

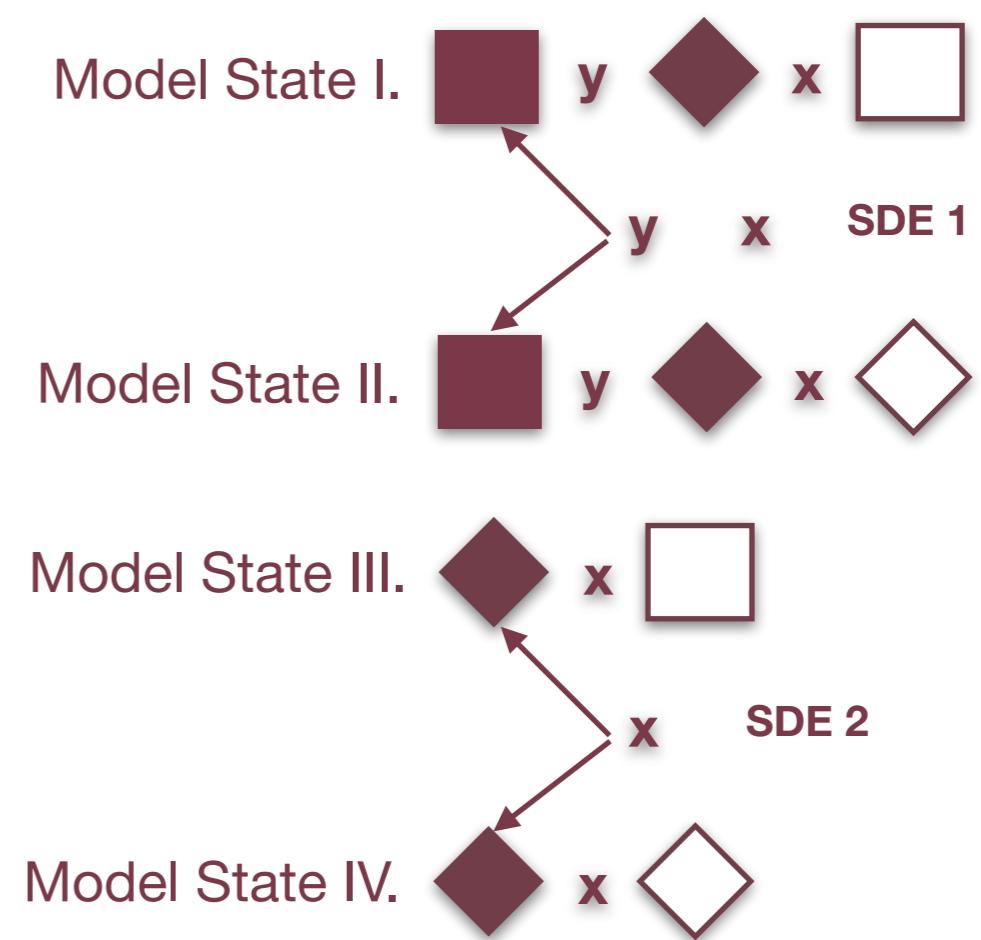
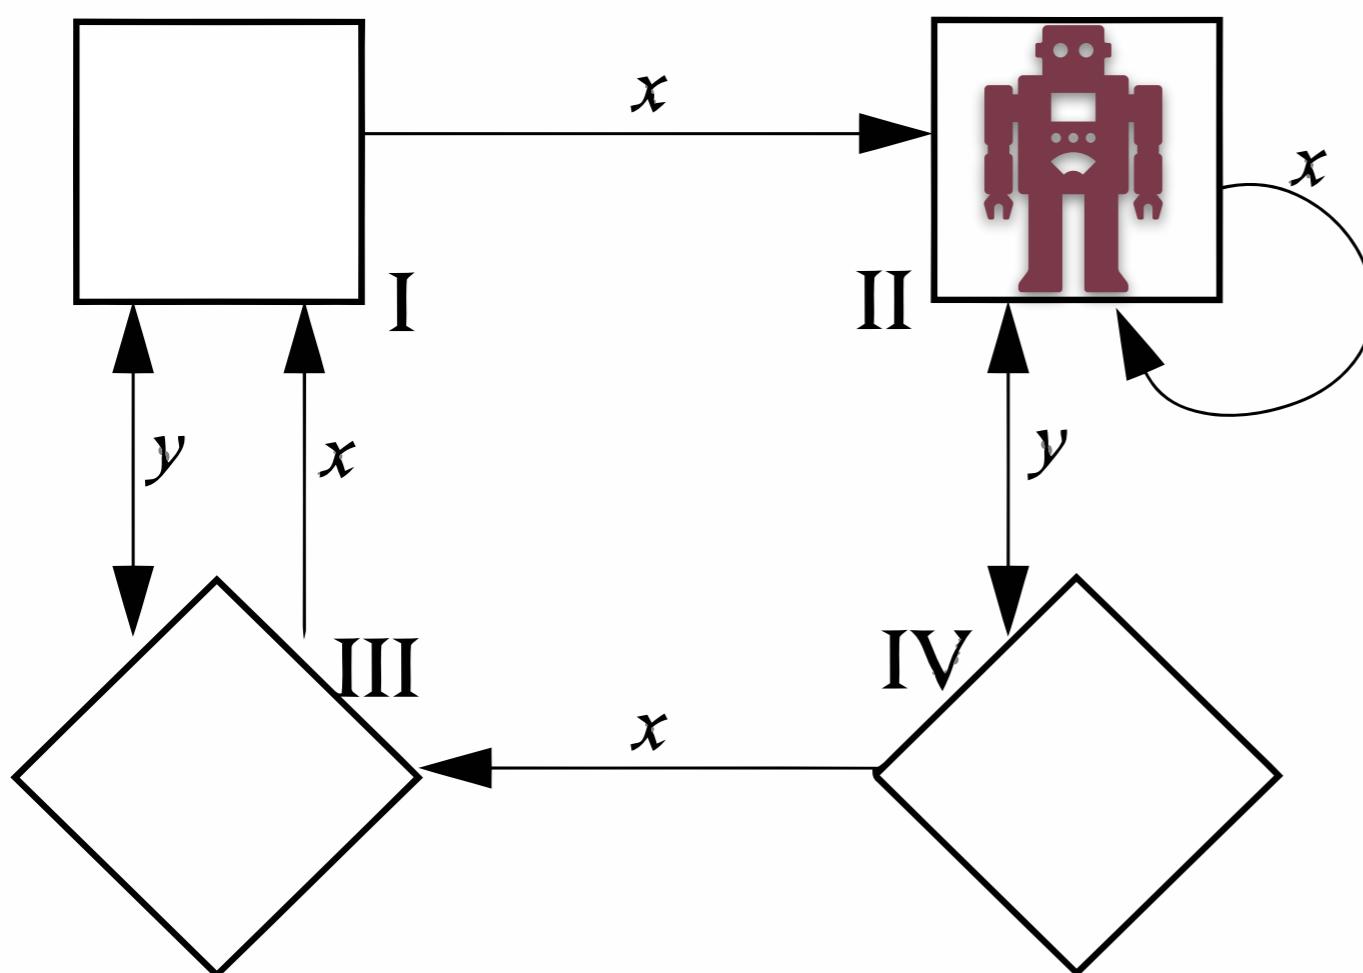
Surprise-based partially-observable Markov decision processes (sPOMDPs)



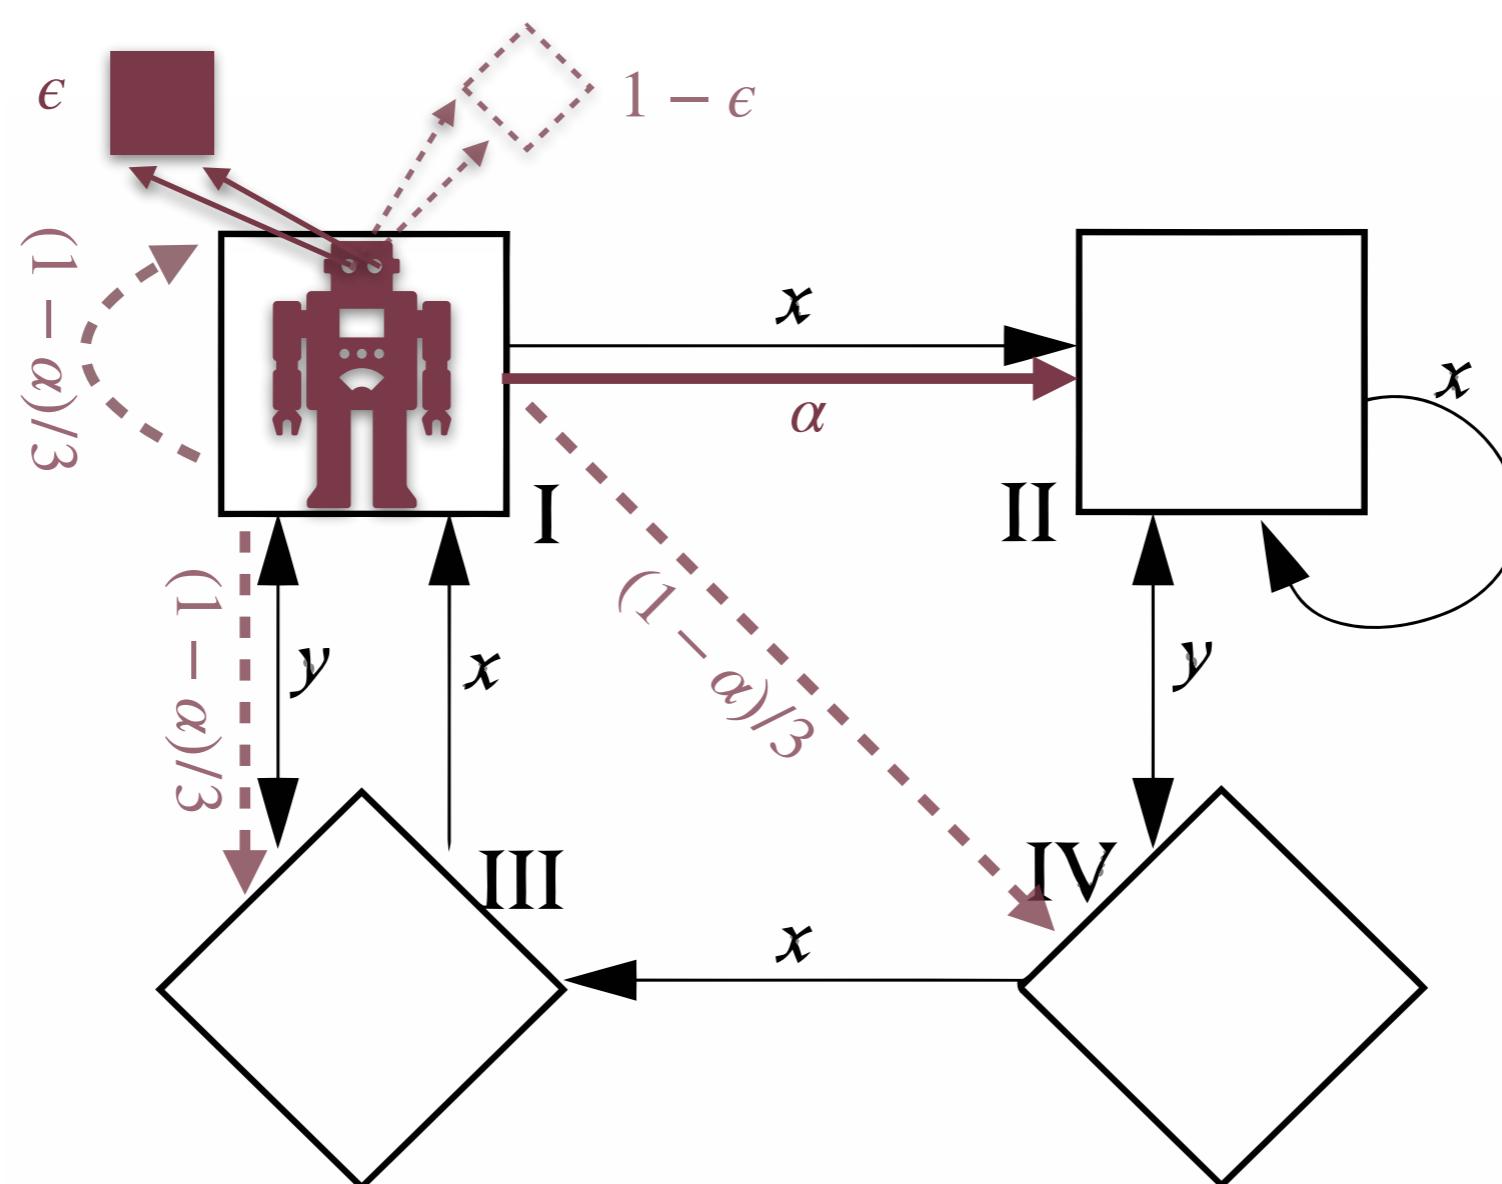
Surprise-based partially-observable Markov decision processes (sPOMDPs)



Surprise-based partially-observable Markov decision processes (sPOMDPs)

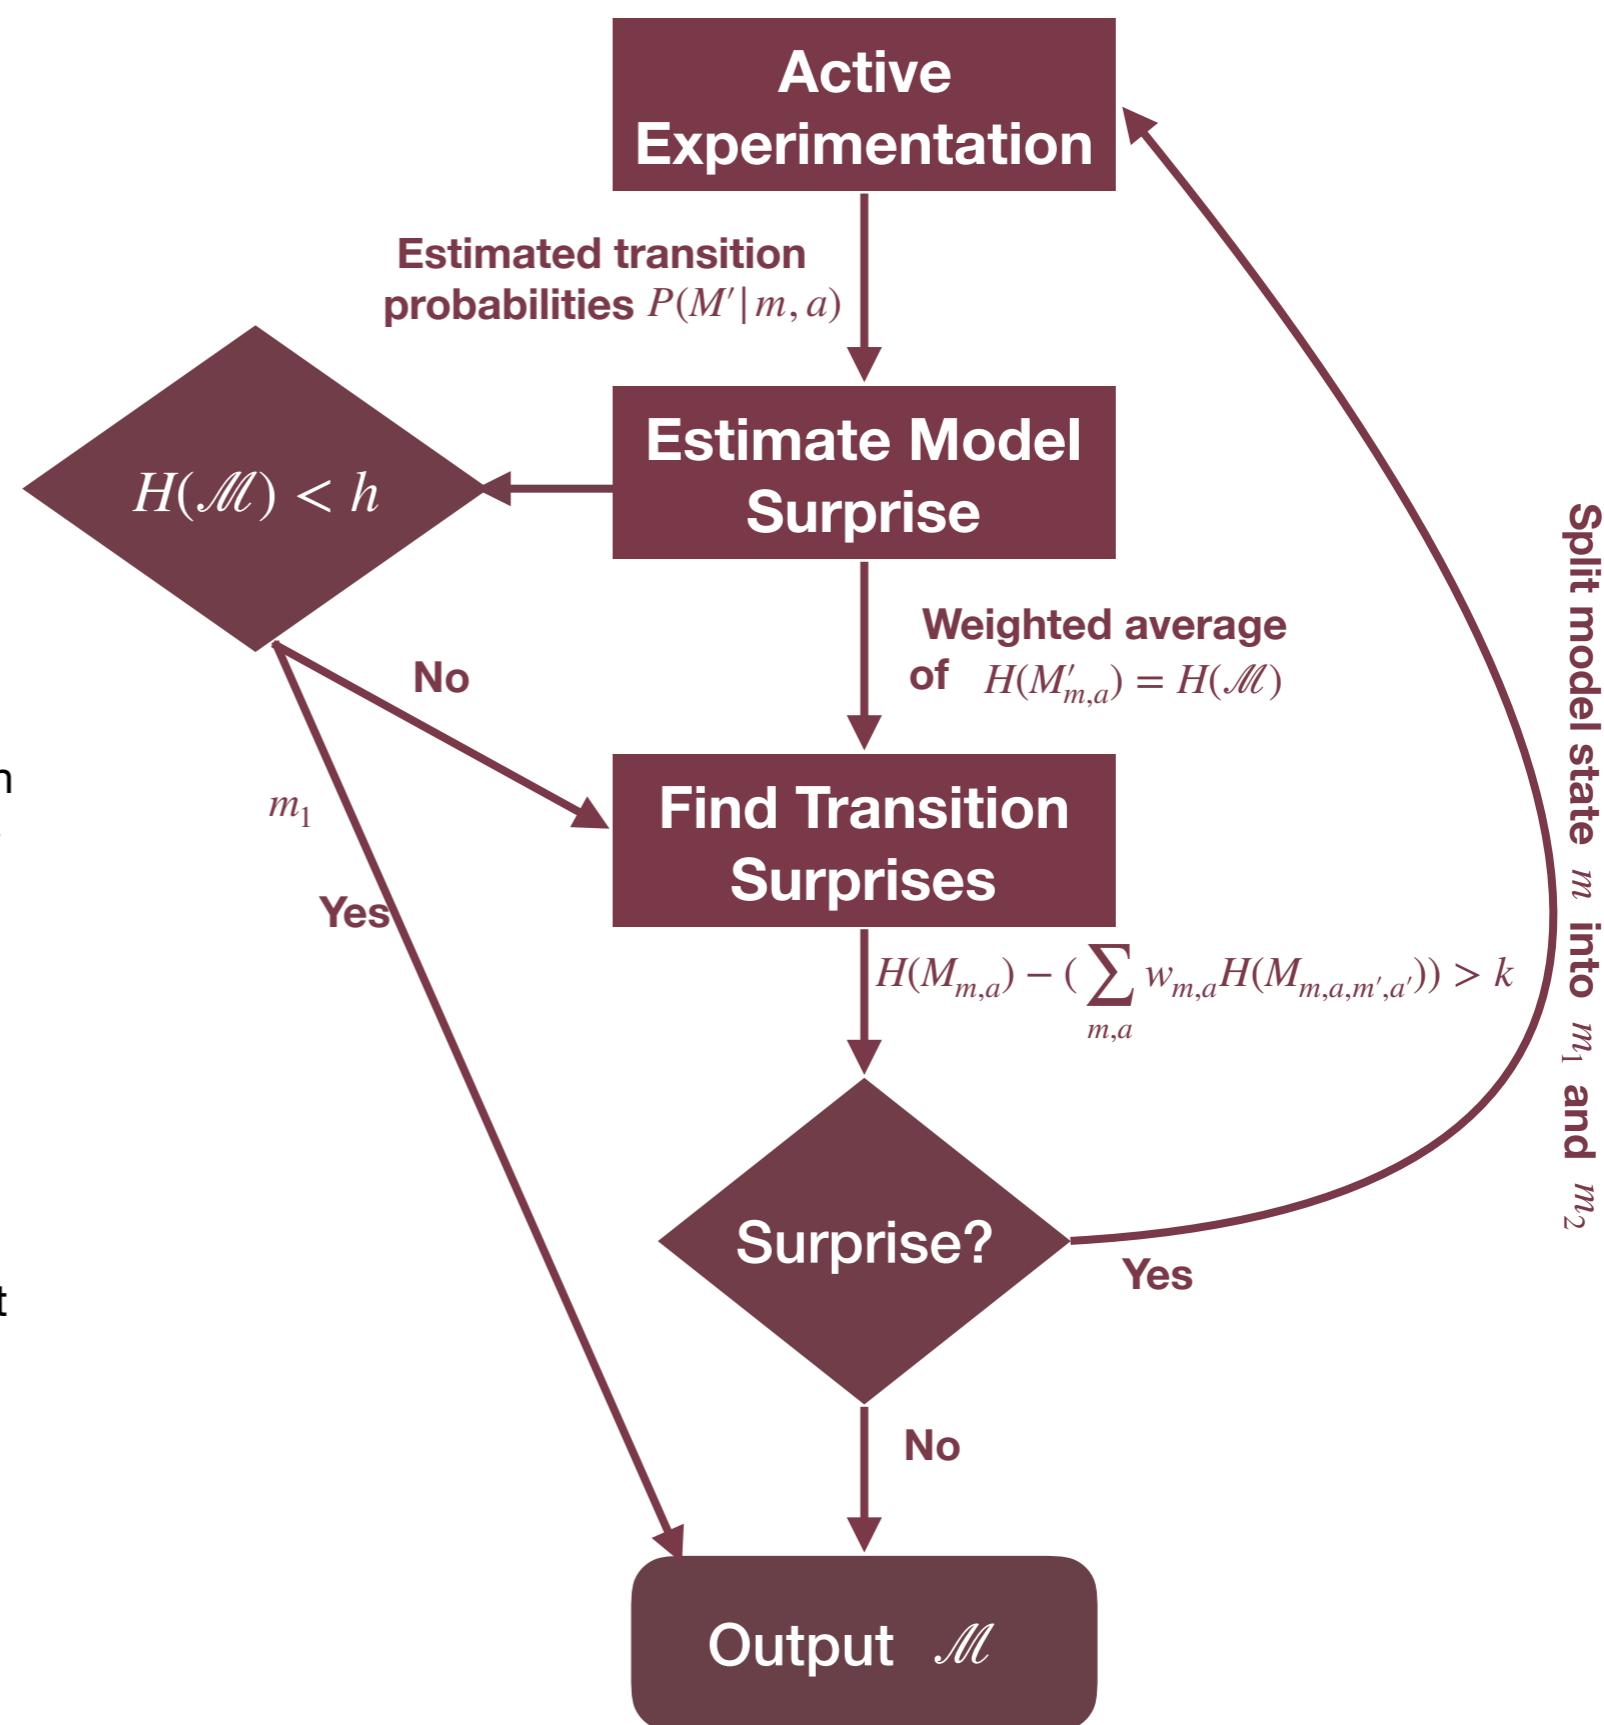


Alpha-epsilon POMDP environments

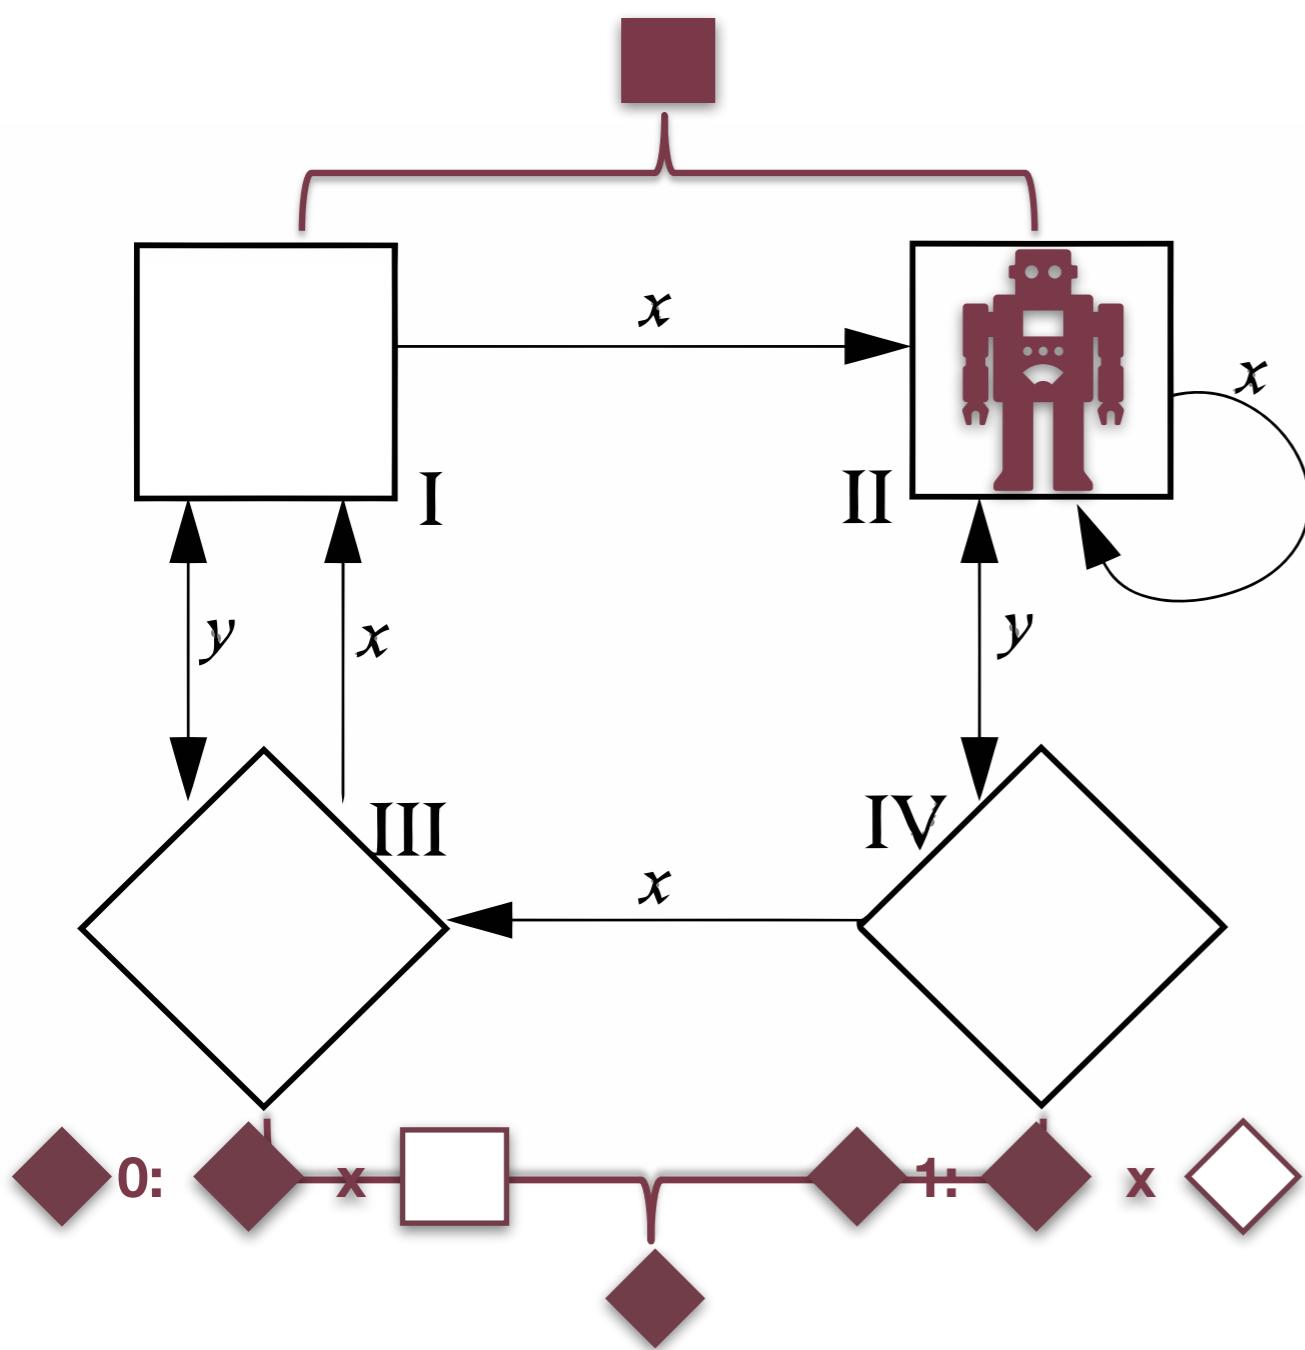


Learning sPOMDPs

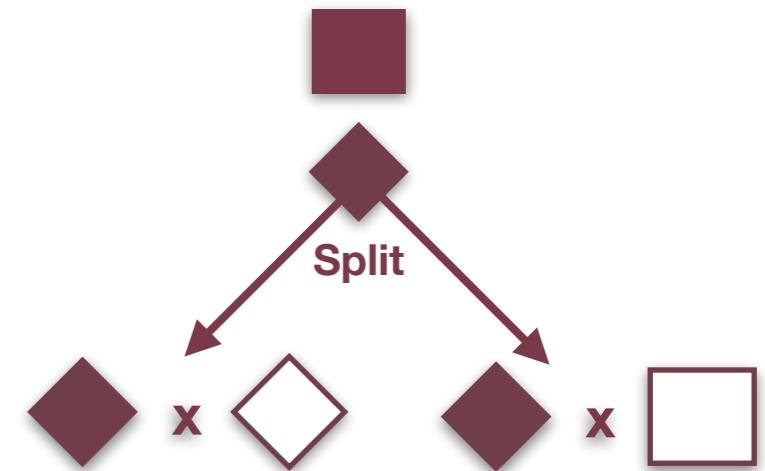
- Build a model whose **surprise value (weighted sum of transition distribution entropies)** is as low as possible.
- Begin by assuming observations are a suitable state space and look for experimental evidence that is **surprising** under this assumption.
- Large **reductions in transition entropy** resulting from considering an additional model state-action pair of history indicate a **violation of the Markov property**, and we **split** model state **m** into two model states:
 - **Intuition:** Model state **m** is representing underlying environment states with **differing dynamics**.
- By executing the **actions of SDEs** as well as random actions (with some probability), the agent can remain **localized in its model** while simultaneously **learning model state transition probabilities**.



Learning sPOMDPs



Current Model States:



Transition Probabilities:

$$P(\{\text{█}, \diamond\} | \text{█}, x) = \{0.99, 0.01\}$$

$$P(\{\text{█}, \diamond\} | \text{█}, y) = \{0.01, 0.99\}$$

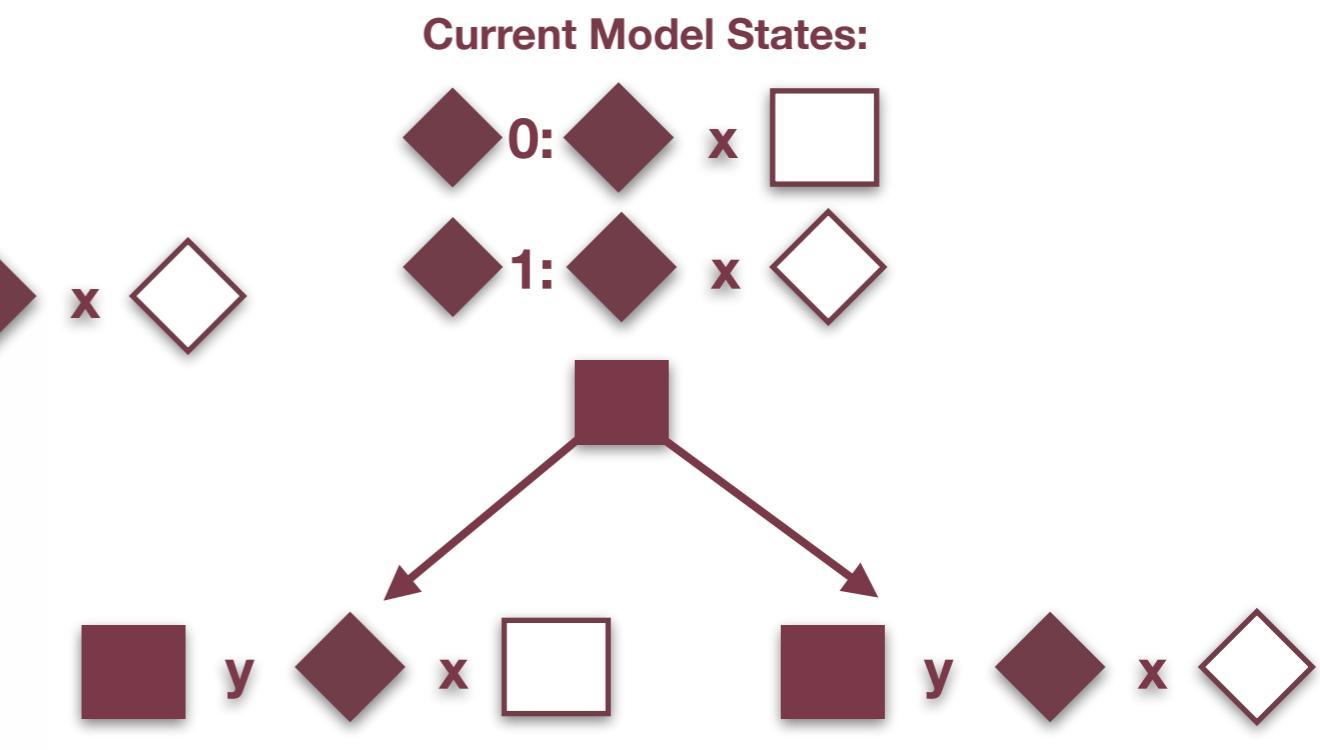
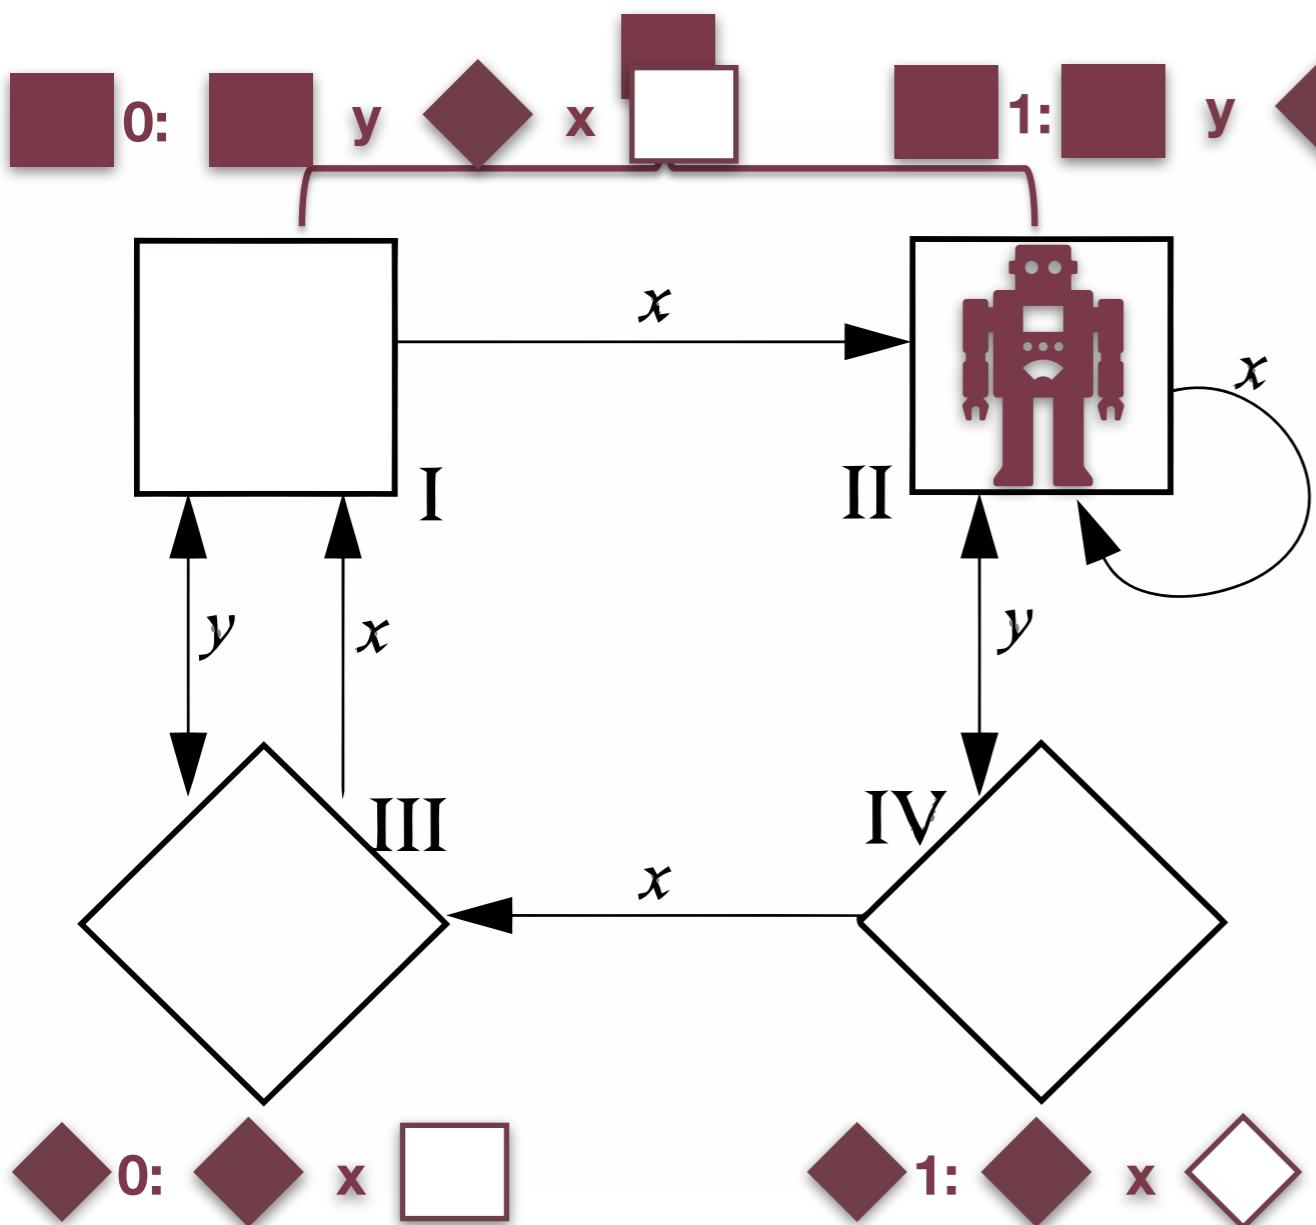
$$P(\{\text{█}, \diamond\} | \diamond, x) = \{0.51, 0.49\}$$

$$P(\{\text{█}, \text{◆}\} | \text{◆}, y) = \{0.99, 0.01\}$$

Surprise!

$$H(M_{m,a}) - \left(\sum_{m,a} w_{m,a} H(M_{m,a,m',a'}) \right) > k$$

Learning sPOMDPs



Transition Probabilities:

$P(\{\square, \diamond, 0, \diamond, 1\} | \square, x) = \{0.98, 0.01, 0.01\}$
 $P(\{\square, \diamond, 0, \diamond, 1\} | \square, y) = \{0.01, 0.33, 0.66\}$
 $P(\{\square, \diamond, 0, \diamond, 1\} | \diamond, 0, x) = \{0.98, 0.01, 0.01\}$
 $P(\{\square, \diamond, 0, \diamond, 1\} | \diamond, 0, y) = \{0.08, 0.01, 0.01\}$
 $P(\{\square, \diamond, 0, \diamond, 1\} | \diamond, 1, x) = \{0.01, 0.98, 0.01\}$
 $P(\{\square, \diamond, 0, \diamond, 1\} | \diamond, 1, y) = \{0.98, 0.01, 0.01\}$

Surprise!

$$H(M_{m,a}) - \left(\sum_{m,a} w_{m,a} H(M_{m,a,m',a'}) \right) > k$$

Theoretical Results

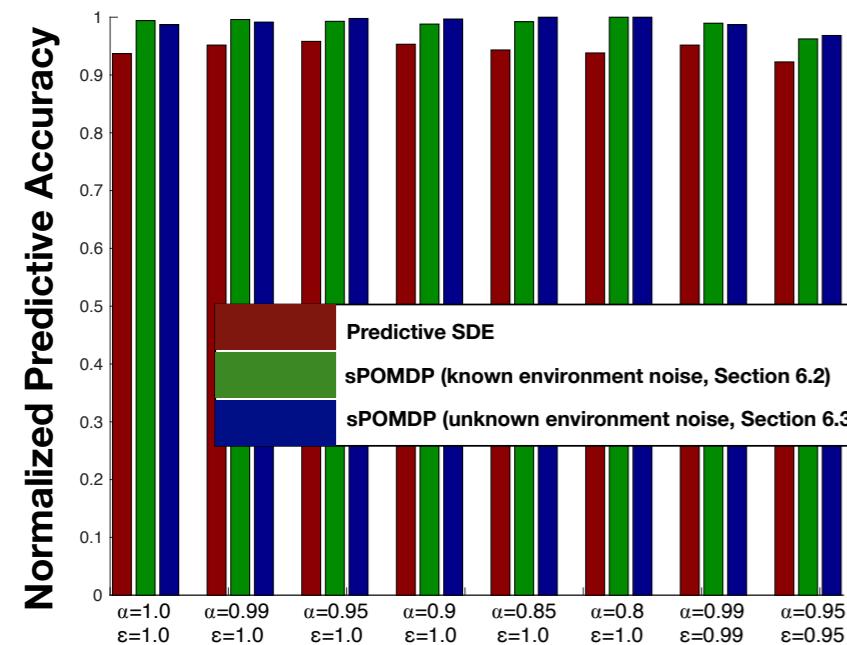
Theorem 1. *For any Moore machine environment \mathcal{E} , there exists a perfect sPOMDP model \mathcal{M} where $|M| \leq |Q|$, $|S| \leq |Q| - 1$, and each SDE $s \in S$ has, at most, $|Q| - 1$ actions, where $|Q|$ is the number of environment states in the minimal representation of \mathcal{E} .*

Theorem 2. *For any rewardless $\alpha\epsilon$ -POMDP environment \mathcal{E} in which $\alpha > 1/|Q|$ and $\epsilon > 1/|O|$, there exists a perfect sPOMDP model \mathcal{M} where $|M| \leq |Q|$, $|S| \leq |Q| - 1$, and each SDE $s \in S$ has, at most, $|Q| - 1$ actions, where $|Q|$ is the number of environment states in the minimal representation of \mathcal{E} .*

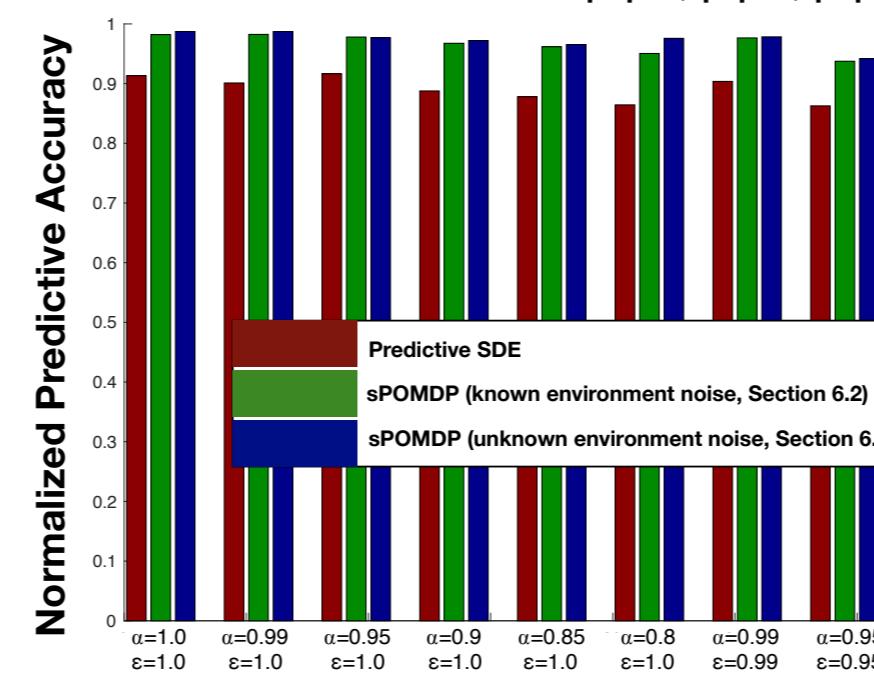
- The constructive proof of **Theorem 2** leads to an **optimal sPOMDP learning procedure**, where, given certain bounds on **alpha** and **epsilon** and an **oracle** that provides the true **model state to model state transition probabilities**, we can provably learn a perfect sPOMDP representation of the given environment.

Experimental Results

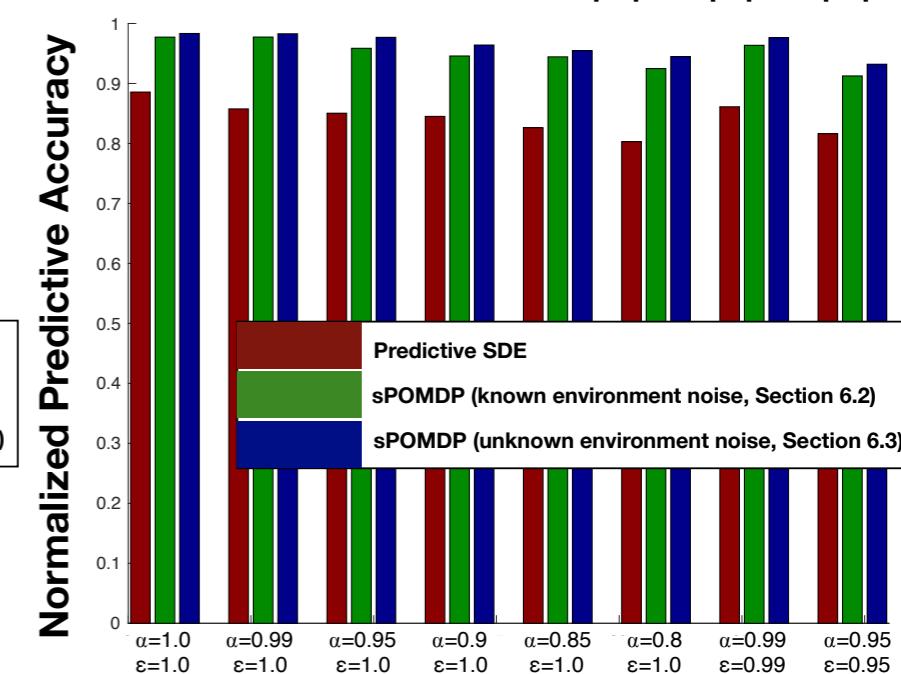
sPOMDP vs Predictive SDE $|Q|=4, |A|=2, |O|=2$



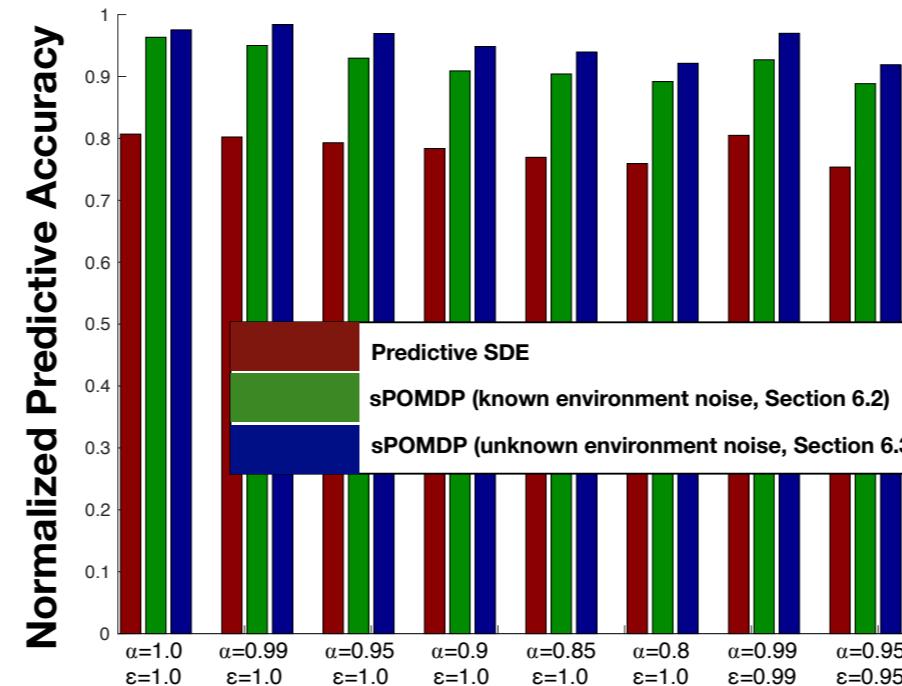
sPOMDP vs Predictive SDE $|Q|=6, |A|=2, |O|=3$



sPOMDP vs Predictive SDE $|Q|=8, |A|=2, |O|=4$

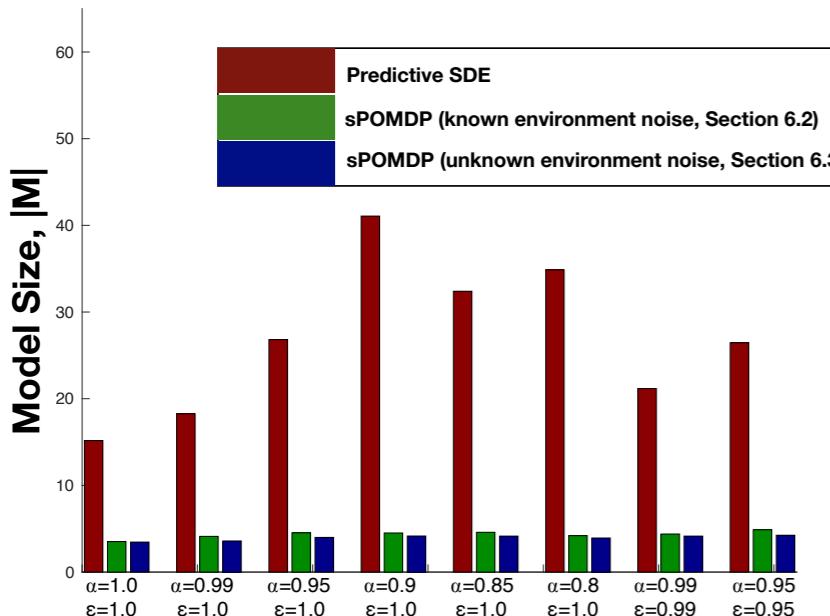


sPOMDP vs Predictive SDE $|Q|=12, |A|=2, |O|=6$

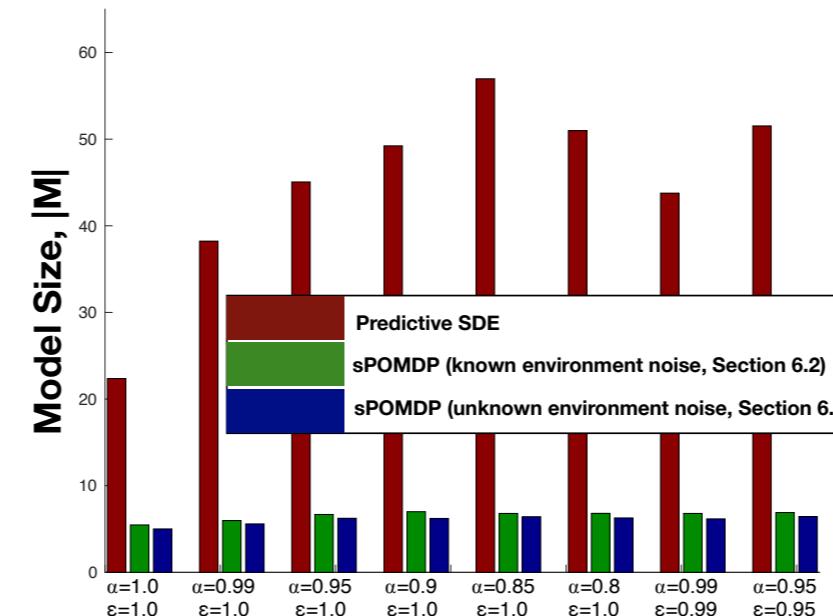


Experimental Results

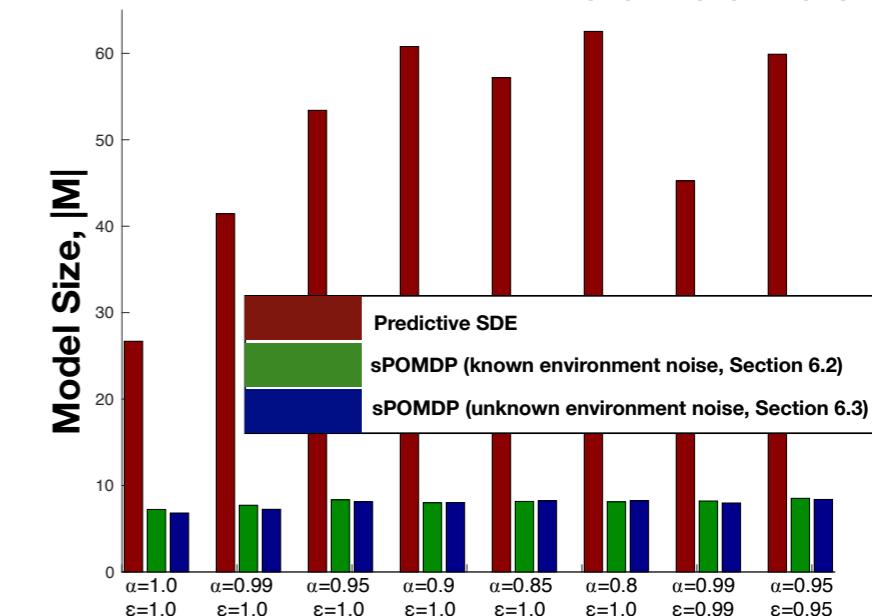
sPOMDP vs Predictive SDE $|Q|=4, |A|=2, |O|=2$



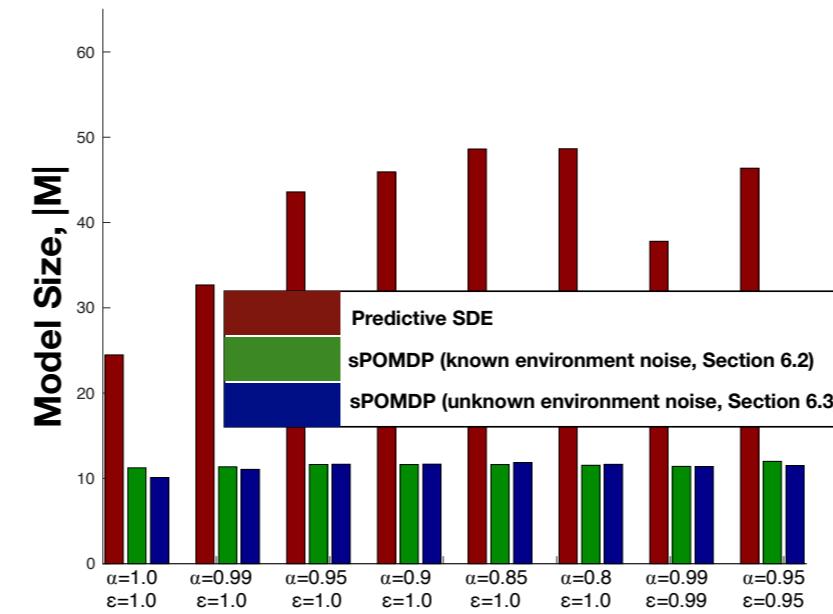
sPOMDP vs Predictive SDE $|Q|=6, |A|=2, |O|=3$



sPOMDP vs Predictive SDE $|Q|=8, |A|=2, |O|=4$



sPOMDP vs Predictive SDE $|Q|=12, |A|=2, |O|=6$



Future Work

- Continuous observation spaces.
- Function approximation techniques for scalability and generalization to unexplored areas of the action/observation space.
 - Balance between **parametric** and **nonparametric** approaches.

